

ORIOUS: A Professor Review Draft on Universal Runtime Safety for Physical AI under Degraded Observation

Author details omitted for professor review draft

Abstract—ORIOUS treats degraded observation as the fundamental hidden hazard for Physical AI: controllers act on telemetry, beliefs, and summaries of the world, while safety is judged on the physically true state. This manuscript presents ORIOUS, Observation–Reality Integrity for Universal Safety, as a reusable runtime layer that detects reliability collapse, inflates the observation-consistent state set, tightens the admissible action set, repairs or replaces unsafe actions, and emits auditable certificates before release. The paper is organized around the universal runtime claim rather than a witness-row-specific forecasting claim. Its scientific object is a shared runtime grammar grounded in one adapter contract, one benchmark schema, one CertOS governance surface, and one theorem bridge. Battery remains the witness row with the deepest theorem-to-code-to-artifact closure; AV and healthcare are defended bounded rows. The central claim is therefore strong but disciplined: ORIOUS establishes a universal runtime safety layer for Physical AI under degraded observation across the active three-domain program, without claiming broader empirical closure beyond that repo-truth boundary.

Index Terms—Physical AI safety, degraded observation, observation–action safety gap, runtime assurance, conformal uncertainty, safety filters, cyber-physical systems, certificates

I. INTRODUCTION AND CONTRIBUTION STATEMENT

Physical AI systems do not act on truth. They act on telemetry, inferred state, delayed measurements, missing packets, stale buffers, and learned surrogates of the world. The resulting engineering problem is deeper than forecast error or planner suboptimality: the software stack can remain internally coherent while the physical system becomes externally unsafe. Batteries may violate hidden state-of-charge limits, vehicles may enter unsafe headway regimes on stale perception, and clinical monitoring pipelines may suppress timely intervention.

ORIOUS is built around the claim that these failures share one universal object: degraded observation changes the semantics of what may still be safely done. The field therefore needs a runtime layer whose job is not to outperform every planner, but to mediate the boundary between degraded observation and physical actuation. ORIOUS, Observation–Reality Integrity for Universal Safety, is proposed as that layer. It detects reliability loss, calibrates the observation-consistent state set, constrains the admissible action set, repairs or replaces unsafe candidate actions, and certifies release through an auditable

lifecycle. The framework is written around the universal runtime claim rather than a witness-row-specific framing. Battery is retained as the deepest witness row, not as the conceptual center of the claim.

This manuscript positions ORIOUS against a precise gap in the literature. Distribution-free inference can quantify uncertainty [1–6]. Runtime assurance and shields can intercept unsafe behavior [7–12]. Safety filters and constrained control can map uncertainty into admissible action sets [13–18]. Fault, anomaly, and drift detection can identify when observation should no longer be trusted [19–24]. What none of these families alone provides is a reusable runtime grammar that links degraded observation to repaired actuation and auditable certificate release across multiple physical domains.

The paper therefore makes four contributions.

- 1) It defines the observation–action safety gap (OASG) as the right universal hazard object for Physical AI under degraded observation.
- 2) It presents the Detect–Calibrate–Constrain–Shield–Certify kernel as a typed runtime safety layer with explicit formulas, pseudocode, and certificate semantics.
- 3) It unifies domain evaluation through one adapter contract, one benchmark schema, and one CertOS governance surface.
- 4) It argues for a bounded-universal evidence posture: battery is the witness row; AV and healthcare are defended bounded rows.

A. What ORIOUS is and is not

ORIOUS is not a new nominal controller, a new world model, or a universal planner. It does not ask every domain to share the same action space, dynamics model, or learning stack. Instead it introduces a narrow but rigorous intervention layer at the moment where candidate actuation is about to be released. That distinction matters. Many physical AI projects already contain estimators, predictors, planners, and local guards. What they usually lack is a reusable semantics for what happens when the controller-visible state is no longer trustworthy. ORIOUS supplies that semantics.

Equally important, ORIOUS is not an excuse to hide ad hoc emergency logic behind a new name. Every intervention mode in the framework must be typed and auditable:

keep, repair, fallback, degrade, or reject. The release decision is therefore promoted from implementation detail to scientific object. That is the conceptual step that lets the paper talk about batteries, vehicles, and healthcare monitoring within one formal vocabulary.

B. Design principles

The framework is organized around four design principles.

- 1) **Post-nominality:** upstream intelligence remains replaceable; ORIUS is evaluated at the release boundary rather than by forecasting or planning quality.
- 2) **Typed degradation:** reliability loss, widened uncertainty, repaired action, fallback mode, and certificate status are explicit runtime objects.
- 3) **Domain-bounded portability:** each adapter owns telemetry semantics and safety geometry, but the kernel and benchmark semantics do not change from one row to the next.
- 4) **Auditable release:** safety is judged on what actuation is actually released, under what certificate, and with what lifecycle trace.

These principles explain why ORIUS should be read as infrastructure for Physical AI safety rather than as one more domain benchmark result. A battery controller can be good or bad; an AV planner can be good or bad; a healthcare monitor can be good or bad. ORIUS asks a different question: when those nominal systems become observation-degraded, what runtime object decides whether the next action remains legally releasable?

The draft is intentionally ambitious but disciplined. It uses a universal writing posture, but it keeps one explicit evidence-boundary section in the main body. It includes formulas, theorem statements, proof sketches, runtime pseudocode, decisive results tables, and cross-domain synthesis, while pushing proof-depth and artifact-depth material into two separate appendices.

II. PROBLEM DEFINITION: DEGRADED OBSERVATION AND OASG

Consider a discrete-time physical system observed through a degraded telemetry channel:

$$x_{t+1} = f(x_t, a_t, \omega_t), \quad y_t = h(x_t, \nu_t), \quad \hat{x}_t = \eta(y_{0:t}), \quad (1)$$

where $x_t \in \mathcal{X}$ is the true state, y_t is the observation packet, $\hat{x}_t$ is the controller-visible state, and $a_t^{\text{cand}} = \pi_t(\hat{x}_t)$ is the candidate action produced by the nominal controller. Let $\mathcal{S} = \{x \in \mathcal{X} : \phi(x) \geq 0\}$ denote the safe set, and let $\mathcal{A}(x)$ denote the set of actions that preserve safety when the true state is x .

The observation-action safety gap is the event

$$\text{OASG}_t \equiv \mathbf{1}[a_t^{\text{cand}} \in \mathcal{A}(\hat{x}_t) \wedge a_t^{\text{cand}} \notin \mathcal{A}(x_t)]. \quad (2)$$

The key point is that OASG is an action-semantics failure, not only a state-estimation failure. A better estimator can reduce its frequency, but the safety event itself is defined




Fig. 1: Degraded observation turns controller legality into a latent safety hazard: the candidate action can be valid on the observed state while unsafe on the true state.

by the legality of released actuation under degraded observation.

ORIUS turns this hazard into a runtime object by making observation quality explicit. Let $w_t \in [0, 1]$ be a reliability score computed from freshness, missingness, sensor disagreement, residual behavior, or domain-specific degradation signatures. The observation-consistent set is then inflated as

$$U_t = \{x \in \mathcal{X} : d_t(x, \hat{x}_t) \leq \beta_0 + \lambda(1 - w_t) + \delta_t\}, \quad (3)$$

where β_0 is the nominal uncertainty radius, $\lambda(1 - w_t)$ is the reliability-conditioned inflation, and δ_t captures drift- or history-based correction. The universal contract does not insist on a single metric d_t ; it insists only that the set widens monotonically when reliability falls.

The corresponding tightened safe-action set is

$$\mathcal{A}_t^{\text{safe}}(U_t) = \{a \in \mathcal{A} : g(x, a) \leq 0, \forall x \in U_t\}, \quad (4)$$

where $g(x, a)$ is the domain safety constraint. The released action is then

$$a_t^{\text{safe}} = \begin{cases} a_t^{\text{cand}}, & a_t^{\text{cand}} \in \mathcal{A}_t^{\text{safe}}(U_t), \\ \mathcal{R}_t(a_t^{\text{cand}}, \mathcal{A}_t^{\text{safe}}(U_t)), & \mathcal{A}_t^{\text{safe}}(U_t) \neq \emptyset, \\ \mathcal{F}_t(c_t), & \mathcal{A}_t^{\text{safe}}(U_t) = \emptyset, \end{cases} \quad (5)$$

with repair map $\mathcal{R}_t$, fallback policy $\mathcal{F}_t$, and certificate context c_t .

This yields the central runtime question for the paper: once degraded observation is acknowledged, what set of actions remains certifiably safe and auditable? ORIUS answers that question with one runtime kernel, one benchmark schema, and one governance discipline. The theorem bridge later shows that the resulting safety layer is not a heuristic wrapper but a structured bound-preserving construction under explicit assumptions.

The same object also governs evaluation. ORIUS measures safety on the true state rather than on the observed state alone. At the episode level this yields the true-state violation rate

$$\text{TSVR} = \frac{1}{T} \sum_{t=0}^{T-1} \mathbf{1}[x_t \notin \mathcal{S}], \quad (6)$$

and the runtime intervention rate

$$\text{IR} = \frac{1}{T} \sum_{t=0}^{T-1} \mathbf{1}[a_t^{\text{safe}} \neq a_t^{\text{cand}}]. \quad (7)$$

TABLE I: Canonical runtime objects for OASG-aware release.

Object	Role
w_t	Reliability score derived from freshness, missingness, and disagreement.
U_t	Observation-consistent state set widened when trust falls.
$\mathcal{A}_t^{\text{safe}}(U_t)$	Tightened action set that remains legal for every state in U_t .
$\mathcal{R}_t$	Repair map used when the candidate action is unsafe but the safe set is non-empty.
$\mathcal{F}_t$	Explicit fallback used when no repaired safe action exists.
c_t	Certificate payload recording margin, reliability, release mode, and audit continuity.

These metrics make the claim falsifiable. A system cannot claim degraded-observation safety if it only looks legal on the observed state, or if it relies on undocumented emergency overrides rather than explicit runtime interventions.

III. RELATED WORK AND THE INSUFFICIENCY OF PRIOR FAMILIES

ORIOUS does not claim that prior work ignored uncertainty or safety. The claim is narrower and stronger: the major adjacent families solve pieces of the problem, but not the full runtime bridge from degraded observation to repaired actuation and auditable release.

A. Conformal prediction and calibration are necessary, but not enough

Distribution-free and conformal methods provide finite-sample uncertainty surfaces for learned predictors [1–6, 25]. These tools are indispensable for ORIOUS because the runtime must widen the observation-consistent set when reliability drops. But they do not, by themselves, specify which action remains legal after that widening. A prediction interval is not a certificate of actuation legality. ORIOUS therefore treats calibration as the input to a safety layer rather than as the safety layer itself.

B. Runtime assurance and Simplex-style supervision do not close the observation gap

Runtime assurance, supervisory control, and Simplex-style architectures [7–12, 26] show that nominal intelligence should not be the final authority over actuation. ORIOUS inherits that insight directly. The gap is that many supervisory designs begin after a candidate action already exists and assume that the state estimate is coherent enough to be monitored. ORIOUS starts earlier: it treats the observation surface itself as unsafe, which forces the runtime to reason jointly about reliability, uncertainty inflation, repair geometry, fallback, and certificate validity.

C. Safety filters and constrained control are mathematically strong but plant-local

Barrier functions, robust MPC, tube MPC, predictive safety filters, and reachability-based controllers can turn model uncertainty into admissible action sets [13–18, 27, 28]. ORIOUS does not compete with that mathematics. It packages those ideas behind a typed runtime contract that can host multiple geometries across multiple plants. The scientific move is not a new universal controller; it is one reusable way to make degraded observation alter actuation semantics.

D. Anomaly, drift, and fault detection stop too early

Change detection, drift detection, anomaly scoring, and CPS fault analysis provide the language for recognizing when telemetry should no longer be trusted [19–24, 29, 30]. Yet many systems stop at an alarm or a reliability score. ORIOUS treats that as an incomplete runtime response. Once trust has fallen, the system still needs to decide which actions remain legal, which actions must be repaired, when fallback is required, and how that decision is certified.

E. Domain-specific safety work rarely becomes a universal runtime grammar

Energy dispatch, autonomous driving, clinical monitoring, and other physical-AI safety settings each contain strong plant-level ideas, but most remain tied to one domain’s state geometry, one benchmark, and one intervention semantics. The result is local excellence without a shared runtime object. ORIOUS positions itself as a cross-domain safety-layer grammar that can sit above those domain-specific methods while keeping the data and proof posture explicitly tiered.

The ORIOUS thesis is therefore integrative but not vague. The universal object is not “safe AI” in the abstract. It is the typed runtime layer that binds reliability, uncertainty, admissible actions, repaired release, and certificate governance under one action-semantics problem.

IV. ORIOUS RUNTIME ARCHITECTURE

ORIOUS is a post-nominal runtime layer. It does not replace the upstream predictor, planner, optimizer, or policy. It wraps them with a narrow contract that begins at raw telemetry and ends at certificate-gated release. This is the architectural reason the framework can remain domain-agnostic even when the domains do not share nominal models or action geometry.

The kernel consists of five stages.

- 1) **Detect:** estimate reliability w_t and drift flags from freshness, missingness, cross-sensor consistency, residual behavior, and domain-specific diagnostics.
- 2) **Calibrate:** inflate the observation-consistent set U_t rather than trusting a nominal point estimate.
- 3) **Constrain:** recompute the safe action set under the widened state set.

TABLE II: Why prior families remain necessary but insufficient for the ORIUS object.

Method family	What it contributes	Why it is insufficient for OASG
Conformal and adaptive calibration	Distribution-free uncertainty and shift-aware coverage control	Does not decide repaired actuation, fallback legality, or certificate release after trust degrades.
Runtime assurance and shielding	Supervisory interception and safety-controller fallback	Often assumes the observation surface is coherent enough to supervise directly.
Safety filters and robust constrained control	Strong admissible-set mathematics under uncertainty	Usually remains plant-local and does not define a reusable benchmark/governance contract.
Anomaly, drift, and fault detection	Detects that telemetry has become unreliable	Stops at alarms or scores unless paired with repair, fallback, and audit semantics.
Domain-specific CPS safety stacks	High-quality local solutions in one plant family	Rarely expose a universal runtime grammar that transfers across domains without changing the scientific object.

- 4) **Shield:** keep, repair, or replace the candidate action.
- 5) **Certify:** emit a causal certificate and update the governance state.

A. Execution path through the kernel

The practical execution path is simple enough to be implemented in a few typed calls, but rich enough to support formal analysis. A telemetry packet enters through the domain adapter with time stamps, missingness indicators, residual channels, and any domain-specific diagnostics. The detect stage converts those observations into a reliability score and degradation flags. The calibrate stage then refuses the usual engineering shortcut of keeping a point state estimate when trust has fallen; instead it materializes the observation-consistent state set that the controller must now respect. The constrain stage computes what actuation remains safe for every state in that set. Shielding then either preserves the candidate action, repairs it into the safe set, or invokes an explicit fallback. Certify binds the release decision to a certificate payload and governance state transition.

That same execution path has different physical interpretations across domains. In battery systems it becomes charge or discharge limitation under hidden state-of-charge uncertainty. In AV it becomes governed braking or maneuver restriction under degraded headway confidence. In healthcare it becomes intervention legality under uncertain vital-sign trust. The point of ORIUS is not that these domains share dynamics. The point is that they share a release-boundary problem once degraded observation is made explicit.

The kernel becomes scientific rather than architectural once its objects are written explicitly. ORIUS uses a reliability-conditioned uncertainty envelope

$$U_t = \{x : m_t(x, \hat{x}_t) \leq q_{1-\alpha_t}(w_t, \Delta_t)\}, \quad (8)$$

where m_t is a domain-specific mismatch measure and $q_{1-\alpha_t}$ is a reliability-aware quantile or inflation operator. The tightened action set is

$$\mathcal{A}_t^{\text{safe}}(U_t) = \{a \in \mathcal{A} : g(x, a) \leq 0, \forall x \in U_t\}, \quad (9)$$

and the certificate-relevant margin is

$$\gamma_t(a) = \inf_{x \in U_t} \phi(f(x, a, 0)), \quad (10)$$

Algorithm 1 Detect–Calibrate–Constrain–Shield–Certify Runtime Step

Require: observation packet o_t , history $\mathcal{H}_t$, candidate action a_t^{cand}
compute reliability score w_t and degradation flags from o_t
build observation-consistent set U_t using reliability-aware inflation
construct tightened safe-action set $\mathcal{A}_t^{\text{safe}}(U_t)$
if $a_t^{\text{cand}} \in \mathcal{A}_t^{\text{safe}}(U_t)$ **then**
 $a_t^{\text{safe}} \leftarrow a_t^{\text{cand}}$
else if $\mathcal{A}_t^{\text{safe}}(U_t) \neq \emptyset$ **then**
 $a_t^{\text{safe}} \leftarrow \mathcal{R}_t(a_t^{\text{cand}}, \mathcal{A}_t^{\text{safe}}(U_t))$
else
 $a_t^{\text{safe}} \leftarrow \mathcal{F}_t(c_t)$
end if
emit certificate c_t with $(w_t, U_t, \gamma_t, \text{mode}, \text{audit_hash})$
release a_t^{safe} only if certificate is valid or fallback is explicit

so that $\gamma_t(a) \geq 0$ indicates one-step safety for every state in the observation-consistent set.

The three equations above are the core of the framework. The first says that trust in observation is itself a runtime variable. The second says that admissible actuation is defined against a set, not only against a point estimate. The third says that the certificate is not a paperwork artifact; it is tied to a concrete safety margin over that set. Together they turn degraded observation from a monitoring nuisance into a first-class control-theoretic and systems object.

This algorithm is the real reusable object in ORIUS. It explains why the framework can support batteries, vehicles, healthcare monitoring, and future additional domains without pretending those settings share the same nominal control law. What is shared is the runtime grammar that mediates degraded observation before actuation.

Two design choices keep this object scientifically clean. First, ORIUS is intentionally narrow: it sits between nominal intelligence and release rather than trying to subsume forecasting, planning, or identification. Second, the runtime layer is typed all the way through. Reliability,




Fig. 2: ORIUS as a universal runtime layer between nominal intelligence and physical actuation. Domain-specific adapters feed a shared Detect–Calibrate–Constrain–Shield–Certify kernel whose outputs are benchmarked and governed under one contract.

uncertainty, repair mode, fallback mode, certificate state, and audit continuity are explicit objects. That is why the framework can be benchmarked and governed as a reusable safety layer instead of described only as a software pattern.

B. Why the architecture is universal

ORIUS achieves universality by freezing the meaning of a few objects and refusing to freeze everything else. The meaning of reliability, uncertainty widening, tightened action legality, repaired release, fallback, and certificate validity is invariant across domains. Telemetry parsers, nominal controllers, local geometry, and domain constraints are not invariant and are deliberately delegated to adapters. This split is what lets the same runtime layer sit above an optimizer in energy systems, a planner in autonomy, a control stack in process industries, or a monitored intervention pipeline in healthcare without becoming mathematically empty.

V. UNIVERSAL CONTRACTS: ADAPTER, BENCHMARK, AND CERTOS GOVERNANCE

ORIUS is organized around three contracts: the adapter contract, the benchmark contract, and the CertOS governance contract. Together they explain how degraded observation is converted into a typed runtime safety layer rather than a battery-only control recipe.

A. Adapter contract

The adapter is the only place where domain-specific telemetry and constraints are allowed to enter the universal runtime. The repository and its supporting design documents make this boundary explicit: `src/orius/certos/runtime.py` provides a domain-neutral governance surface, `docs/UNIVERSAL_BENCHMARK_SPEC.md` defines the common step schema, and `docs/UNIVERSAL_GOVERNANCE_SPEC.md` requires a domain governance policy instead of battery-shaped runtime state.

In practical terms, each domain adapter must expose a true-state violation predicate, an observed-state satisfiability predicate, an optional margin computation, a fallback action, and certificate-compatible action repair. That is the minimum needed to make ORIUS portable across battery energy storage, autonomous vehicles, healthcare monitoring, and future additional domains without changing the meaning of safety at the kernel level.

The adapter can therefore be read as a typed semantic projection

$$\mathcal{D}_t : (o_t, \mathcal{H}_t, a_t^{\text{cand}}) \mapsto (w_t, U_t, \mathcal{A}_t^{\text{safe}}, \mathcal{R}_t, \mathcal{F}_t, c_t),$$

where telemetry and nominal actions are translated into the universal runtime objects that ORIUS knows how to govern. This is one of the most important clarifications in the paper: universality does not mean a universal plant model. It means one typed safety projection from domain observation to governed actuation.

B. Benchmark contract

ORIUS-Bench is the manuscript-facing replay contract. The universal step record keeps the same fields across domains: `true_constraint_violated`, `observed_constraint_satisfied`, `true_margin`, `observed_margin`, `intervened`, `fallback_used`, `certificate_valid`, `latency_us`, and `domain_metrics`. The metrics family is equally fixed: TSVR, OASG, certificate validity accuracy, intervention rate, audit completeness, latency percentiles, and coverage diagnostics when conformal uncertainty is active.

The benchmark contract matters because it decides what evidence counts. A step is not judged by whether the nominal controller felt confident. It is judged by whether the true constraint was violated, whether the observed state looked safe, whether ORIUS intervened, whether fallback was used, whether the certificate remained valid, and what latency budget was consumed. This prevents the evaluation from collapsing into prediction accuracy or task reward while still remaining agnostic to the nominal model.

This contract prevents the argument from collapsing into a single battery result table. It also keeps the evidence boundary honest: the same schema can be reused across the defended three-domain lane and future additional domains, but the rows do not carry the same proof depth.

C. CertOS governance contract

CertOS is the runtime governance layer. Its role is deliberately narrow: do not release an action unless the runtime can either validate a current certificate or replace the action with an explicit fallback. The policy surface is domain-neutral and consists of five hooks: `is_actuation(action)`, `fallback_action(constraints, state)`, `horizon_update(certificate, reliability, drift, constraints, validity_horizon)`,

Algorithm 2 CertOS Certificate Release and Lifecycle Update

Require: repaired action a_t^{safe} , certificate candidate c_t , reliability w_t , constraints κ_t
project a_t^{safe} through `integrity_projection`
if `required_certificate_fields()` missing **then**
mark certificate invalid and request fallback
end if
validate lifecycle state and audit hash continuity
if certificate valid **then**
compute horizon update using `horizon_update`
append validated certificate to audit ledger
release projected action
else
replace with `fallback_action` and append degraded certificate
release fallback only if lifecycle invariants still hold
end if

TABLE III: The three public ORIUS contracts used throughout the professor draft.

Layer	Contract obligation
Adapter	Convert degraded telemetry into reliability, uncertainty, safe-action, repair, fallback, and certificate objects.
Benchmark	Keep one true-state safety schema and one metric family across all domains.
CertOS	Release only validated or explicit fallback actions with audit continuity.

`required_certificate_fields()`, and `integrity_projection(action)`.

The invariants are correspondingly small and auditable: no actuation without a valid or fallback certificate path, an unbroken certificate hash chain, and fallback only when validity collapses. That is the right level of abstraction for a professor-facing draft that wants to argue for a fundamental safety layer.

This is also where ORIUS separates itself from many safety wrappers that are difficult to compare scientifically. If a runtime layer cannot explain what is released, under what certificate, for how long that certificate remains valid, and how fallback is recorded, then it is hard to distinguish principled safety governance from an implementation-specific patch. CertOS is deliberately narrow so that those questions are answerable.

The claim boundary is therefore straightforward: ORIUS is a universal runtime safety contract for Physical AI under degraded observation, but the evidence remains tiered by domain and is not presented as equal-depth proof across all rows.

That methodological discipline matters as much as the API itself. It prevents the paper from reading like a battery controller with extra documentation. By freezing the adapter boundary, the benchmark record, and the governance hooks, ORIUS defines one reproducible scientific object whose meaning does not change from one domain




Fig. 3: The theory bridge separates three levels of the claim: formal runtime objects, domain-specific evidence obligations, and promotion status under the defended 3-domain discipline.

chapter to the next.

VI. THEORY BRIDGE AND PROMOTION DISCIPLINE

This section distills the formal spine of ORIUS for a professor-facing IEEE draft. It connects the battery proof ladder, the universal kernel, and the cross-domain promotion logic without reintroducing the full monograph narrative. The key point is that ORIUS is not presented as a battery theorem generalized by analogy. It is presented as a runtime safety grammar whose transfer obligations are explicit, checkable, and separate from the evidence hierarchy.

A. Formal setting

We consider a partially observed safety-constrained control system on a discrete horizon $t = 0, 1, \dots, T - 1$:

$$x_{t+1} = f(x_t, a_t, \omega_t), \quad y_t = h(x_t, \nu_t), \quad w_t = \psi(y_{0:t}),$$

with true state $x_t \in \mathcal{X}$, action $a_t \in \mathcal{A}$, observation history $\mathcal{H}_t = \sigma(y_{0:t})$, and reliability score $w_t \in [0, 1]$. The safe set is

$$\mathcal{S} = \{x \in \mathcal{X} : \phi(x) \geq 0\},$$

and the episode violation count is

$$V_T = \sum_{t=0}^{T-1} \mathbf{1}[x_t \notin \mathcal{S}].$$

The ORIUS runtime layer takes a nominal candidate action $a_t^{\text{cand}} = \pi_t(\mathcal{H}_t)$, constructs an observation-consistent state set U_t , tightens the admissible action set $\mathcal{A}_t^{\text{safe}}(U_t)$, repairs or replaces the candidate action, and emits a causal certificate.

B. Promotion discipline

The active ORIUS branch uses a promotion discipline rather than a flat universality gate. A domain instantiation is not promoted by rhetorical similarity; it is promoted when it supplies:

- 1) a declared observation model and a reproducible replay or experimental protocol;

- 2) a sound observation-consistent set U_t and a domain-specific safe-action set;
- 3) a repair or fallback operator that is measurable, causal, and certificate-backed;
- 4) a benchmark compatible with the universal schema and a certificate that can be audited;
- 5) a domain-specific safety reduction result that can be checked against the shared contract.

This discipline rule is the manuscript's main safeguard: architecture may be universal before evidence is broad, but the defended empirical claim remains limited to the promoted three-domain lane until additional domains discharge the same obligations. For theorem promotion specifically, status is governed by `reports/publication/theorem_promotion_matrix.csv` and `reports/publication/theorem_result_cards/*.json`. The dashboard and manuscript may call a row flagship only when the statement, assumptions, proof, code anchor, tests, artifacts, hashes, manuscript anchor, and claim-boundary note are all present.

Proposition VI.1 (Promotion discipline). *If a domain instantiation satisfies the universal adapter contract and discharges the promotion obligations above, then the ORIUS kernel can be interpreted on that domain as a runtime safety layer with the same formal objects used in the battery reference proof: observation reliability, uncertainty inflation, admissible repair, fallback, and certificate semantics. If those obligations are not discharged, the manuscript may retain only a bounded or gated claim.*

Proof sketch. The proposition is a packaging statement, not a new control theorem. It follows directly from the typed kernel and adapter surface described in the universal-completeness chapter: the runtime only requires an observation packet, a reliability score, an observation-consistent state set, a tightened safe-action set, a repair operator, a fallback, and a causal certificate. Appendix A lists these objects and their assumption dependencies. Once a domain supplies the same objects, the formal language of the battery ladder applies without changing the runtime contract; what changes is the evidence status of the domain. $\square$

C. T5–T7_Battery: Temporal certificate and fallback support

Certificates are temporal objects. T5 (Finite-Horizon Certificate Validity) is a flagship finite-horizon theorem: if the reachable tube from the observation-consistent state set stays inside the safe set through horizon h , the certificate remains valid through that horizon unless invalidating evidence arrives. T6 (Certificate Expiration Bound) adds a confidence-aware closed-form first-passage lower-bound proposition on the remaining certificate life, using the current boundary gap and the drift proxy σ_d . T7_Battery (Feasible Fallback Existence) then closes the operational loop as a supporting proposition with a piecewise battery fallback: safe hold on the interior, boundary-aware safe

landing near the boundary, and fail-closed behavior when neither case is certifiable. These surfaces are runtime-enforced: when the T6 lower bound drops below one step, the guarantee check fails and the shield must discharge the C7_BatteryFallback contract.

D. T9: Impossibility of quality-ignorant mandatory release

T9 is promoted only in its scoped empty-safe-core form. It is not a universal claim about controllers allowed to abstain, expand uncertainty, fallback, or deny release.

The first lower-bound claim is that ignoring reliability is not a benign simplification when the observation class hides latent states with disjoint safe actions.

Theorem VI.2 (T9: Impossibility of quality-ignorant mandatory release). *For an observation o , let $B(o) = \{x : h(x) = o\}$ and $K(o) = \cap_{x \in B(o)} C(x)$. If $K(o) = \emptyset$, no observation-only mandatory-release controller can guarantee true-state safety for every $x \in B(o)$.*

Proof sketch. The controller must choose one action $a = \pi(o)$. Empty common safe core means no single action belongs to all $C(x)$ for $x \in B(o)$, so some latent state in the ambiguity class makes that mandatory release unsafe. $\square$

E. T10: Boundary-indistinguishability lower bound

T10 is a two-state lower-bound theorem. It does not claim a global minimax frontier or matching optimality result.

Theorem VI.3 (T10: Boundary-indistinguishability lower bound). *Let x_0, x_1 be boundary states with observation laws P_0, P_1 . If $d_{TV}(P_0, P_1) \leq \epsilon$ and $C(x_0) \cap C(x_1) = \emptyset$, then every observation-only mandatory-release controller satisfies*

$$\sup_{i \in \{0,1\}} \Pr_{O \sim P_i} [\pi(O) \notin C(x_i)] \geq \frac{1 - \epsilon}{2}.$$

Proof sketch. This is a Le Cam style reduction. The two disjoint safe-action sets mean one released action must be wrong for at least one latent state. The smaller the total variation gap between the two observation laws, the less information the controller has about which state generated the observation. The two-point testing inequality gives the displayed worst-case risk lower bound. $\square$

F. Active theorem package after promotion

The promoted theorem package is deliberately stronger than a relabeling exercise. Each surface now has a narrow mathematical statement, an explicit assumption footprint, a proof source, a code anchor, tests, artifacts, and a claim-boundary note. The result is an active register with no defended draft rows. T5 is a finite-horizon certificate validity theorem rather than a definition-only row: the certificate is valid exactly while the reachable tube remains inside the safe set and no invalidating evidence has arrived. T8 is a graceful-degradation dominance theorem only after adding useful work: shutdown may be safe, but it does

not satisfy the same safety-work partial order unless it preserves the required operation fraction. T9 and T10 are promoted as scoped lower-bound results over observation-only mandatory release, not as unrestricted statements about every possible runtime architecture.

The observation-robustness rows follow the same rule. The Byzantine reliability result is stated for a trimmed aggregation model with fewer than half corrupted channels; the corollary to the ORIUS risk envelope is allowed only under that robust-score error bound. The stale-observation result is stated as bounded-drift uncertainty growth, not as a physical law of geometric decay. The runtime law suite is likewise narrow: reliability-monotone inflation, safe-set antitonicity, intervention thresholding, and the observation-ambiguity sandwich are small lemmas and propositions that support the larger proof spine. Tminimax is explicitly finite and scoped to two-state boundary ambiguity classes, while the sensor-necessity row reduces missing critical coordinates to the empty-safe-core condition. These constraints are the reason the theorem package can be ambitious without becoming overbroad.

G. T11: Structural transfer through the universal adapter contract

The third claim is the one that matters most for a universal paper draft. It says that the battery proof pattern can transfer if, and only if, the domain exposes the right runtime objects.

Theorem VI.4 (T11: Structural transfer). *Let $\mathcal{D}$ be any domain instantiation that satisfies the universal adapter contract: it provides an observation packet, a reliability score w_t , an observation-consistent state set U_t , a domain-safe action set $\mathcal{A}_t^{\text{safe}}(U_t)$, a repair map, and a fallback action. If $\Pr[x_t \in U_t \mid \mathcal{H}_t] \geq 1 - \alpha$ and every repaired action lies in the tightened safe-action set, then*

$$\Pr[x_{t+1} \in \mathcal{S} \mid \mathcal{H}_t] \geq 1 - \alpha.$$

Proof sketch. Condition on the observation history. On the event $x_t \in U_t$, soundness of the tightened safe-action set implies that every admissible repaired action preserves the safe set at the next step. If the tightened set is empty, the fallback action preserves safety by contract. Therefore the only remaining source of failure is the probability that the true state lies outside U_t , which is at most α . Appendix A and the theorem-surface register record the battery proof instances that instantiate this logic in the reference domain. $\square$

Corollary VI.5 (Episode aggregation under a domain risk budget). *If a domain supplies a predictable per-step risk budget r_t such that $\Pr[x_{t+1} \notin \mathcal{S} \mid \mathcal{H}_t] \leq r_t$, then*

$$\mathbb{E}[V_T] \leq \sum_{t=1}^T r_t.$$

When the budget has the battery-style envelope $r_t \leq \alpha(1 - w_t)$, the familiar upper bound $\mathbb{E}[V_T] \leq \alpha(1 - \bar{w})T$ follows immediately.

H. Interpretation for the IEEE draft

The practical message for the IEEE manuscript is that the theory bridge does not need to be a second thesis chapter. It needs to do three things well: state the universal objects, state the universality gate, and state the proof skeleton that connects the reference battery ladder to any new domain that clears the adapter contract. That is the strongest possible version of the claim without inventing unsupported empirical parity.

This is why the theory bridge belongs in the main paper instead of disappearing into a supplement. The ORIUS claim is not merely that a benchmark improved. It is that degraded observation changes the structure of what safe release means, and that the runtime layer makes that change explicit in a way that can be proved, benchmarked, and audited.

VII. BATTERY WITNESS DOMAIN

The battery row is the strongest witness in ORIUS because it closes the entire chain from assumptions and theorems to controller code, replay artifacts, latency budgets, and certificate traces. It is therefore the right place to show the runtime grammar in its deepest form without allowing the paper to collapse back into a battery-only manuscript.

A. Problem, state, action, and safety object

The witness task is storage dispatch under uncertain load, renewable injection, price signals, and telemetry faults. The true state includes at least state of charge, operational bounds, and dispatch context; the candidate action is charge/discharge power. The governing safety object is not cost. It is hidden state violation under degraded observation, measured on the true state rather than on controller-visible legality.

B. Data and replay surface

The battery evidence is built on the locked CPSBench-derived replay surface with 168-hour episodes, fixed seeds, and fault families that include dropout, stale sensing, delay, drift, and mixed degradation. This is what makes the battery row useful as a witness: it is a repeatable degraded-observation environment rather than a one-off case study.

C. Witness result

The decisive battery result is that ORIUS repair-aware controllers eliminate true-state violations under the locked protocol while non-repair baselines do not. The deep OQE variant is not presented as a miracle predictor; it is a quality signal routed into a repair-and-certify layer that prevents hidden violations.

The deeper trace, ablation, and raw-sequence evidence remains important, but it belongs in Appendix B. In the main paper the battery row serves one purpose: it proves that the runtime safety layer is not merely architectural.

TABLE IV: Battery witness summary from the tracked main battery result surface.

Controller	TSVR	IR (%)	Cost Δ (%)
ORIUS FTIT	0.0	1.9	0.3
ORIUS wrapped	0.0	0.0	0.1
Adaptive conformal + repair	0.0	0.0	-0.0
Deterministic LP	0.0	0.0	0.0
Robust fixed interval	0.0	0.0	-0.0




Fig. 4: Battery witness comparison for reliability-conditioned ORIUS variants. The important signal is structural: repair-aware variants preserve zero true-state violations under degraded telemetry.

It changes the released-action semantics and removes true-state failures on the strongest witness surface currently available in the repo.

VIII. RUNTIME AND EMPIRICAL RESULTS

The decisive empirical question for ORIUS is not whether one domain achieves the lowest cost or the smallest interval width. It is whether the runtime layer materially reduces true-state violations under degraded observation while remaining operationally cheap and auditable. The current repository truth supports that claim most strongly in the battery witness row and then at bounded depth across AV and healthcare.

The battery witness establishes the strongest result surface: ORIUS repair-aware controllers eliminate true-state violations under the locked degraded-telemetry protocol while maintaining bounded intervention rates. The cross-domain runtime evidence then shows that the same runtime grammar remains cheap enough to be plausible across non-energy domains. The important result is therefore joint: safety repair is doing real work, and the added runtime layer is not computationally extravagant.

Three results matter most.

- 1) **Safety reduction:** the battery witness row is the clearest proof that reliability-conditioned repair prevents hidden violations that remain visible under deterministic and static-interval baselines.
- 2) **Runtime plausibility:** the defended rows remain in the sub-millisecond regime at the P95 level, which means ORIUS behaves like a lightweight safety layer rather than a second controller.
- 3) **Governed release:** certificates, fallback paths, and audit continuity are observed rather than assumed, so

TABLE V: Runtime and governance summary for the professor-facing ORIUS draft.

Domain	Row	P95 latency (ms)	Fallback (%)	Audit / lifecycle (%)	Boundary
Battery Energy Storage	Witness row	0.812	100.0	100.0	deepest theorem-to-artifact closure
Autonomous Vehicles	Defended bounded row	0.744	100.0	100.0	bounded brake-hold runtime contract
Medical and Healthcare Monitoring	Defended bounded row	0.701	100.0	100.0	bounded monitoring and alert-release semantics

runtime safety is measured as a released-action discipline instead of an undocumented filter.

The combined runtime-governance table also makes the evidence boundary easy to read. Battery is the witness row. AV and healthcare are defended bounded rows with full certificate rates and bounded runtime semantics. That is exactly the disciplined claim shape a professor-facing draft should preserve.

What this means scientifically is that ORIUS is being evaluated as a release layer, not as a task policy. The battery witness demonstrates the deepest theorem-to-runtime closure, but the runtime-governance table shows the more general point: intervention, fallback, latency, and certificate validity can be measured under one contract across qualitatively different physical systems. That is why the paper uses bounded rows and a witness row rather than forcing every domain into one artificial leaderboard.

The deployment-scope table prevents the results section from sliding into unbounded deployment rhetoric. Even the strongest battery row stops short of unrestricted field deployment. The defended bounded rows stop at replay-backed and runtime-governed validation. Future additional domains remain outside the defended lane precisely so that the paper’s universal architecture claim stays larger than a single benchmark number but smaller than a universal deployment claim.

One further interpretive point matters for the professor review draft. ORIUS is not claiming that every domain should optimize the same surface metric. Some rows care about dispatch legality, others about headway or envelope preservation, others about intervention legality. The universal result is instead that each domain can expose a true-state safety object, and that ORIUS can reduce risk on that object by governing release under degraded observation. That is the empirical thesis the paper can defend without overstating parity.

IX. CROSS-DOMAIN SYNTHESIS

ORIUS is best read as a universal runtime safety layer for Physical AI, not as a single-domain controller that happens to generalize. The common object across the monograph is the degraded-observation hazard: the controller sees an observation surface, but safety is judged on the true state or the observation-consistent state set. That is why the same runtime grammar recurs across the active battery, autonomous-vehicle, and healthcare rows.

The synthesis that follows from this status is narrower and stronger than a generic transfer-learning claim.

ORIUS standardizes the runtime objects that matter for safety under degraded observation: reliability assessment, uncertainty widening, admissible-set tightening, repair or fallback, certificate emission, and auditable lifecycle logging. What varies across domains is the safety geometry, the telemetry parser, and the domain-specific repair policy. What stays invariant is the safety-layer grammar.

That distinction is the key cross-domain result. Battery remains the witness row, while autonomous vehicles and healthcare are defended bounded rows with replay-backed runtime evidence under explicit contract limits.

The professor-facing takeaway is therefore precise: ORIUS already supports a universal safety architecture across the active three-domain lane, but the evidence ladder is tiered. The architecture is universal in contract; the closure depth is not equal across Battery, AV, and Healthcare. That is the right level of honesty for an R1-grade monograph.

A. Why ORIUS is more than a multi-domain wrapper

Many systems can be described across multiple domains if the description remains qualitative enough. ORIUS aims at a stronger standard. It keeps the same runtime question, the same typed release objects, and the same benchmark semantics across all rows. That is why the framework should be understood as universal in *runtime logic* rather than merely broad in *application count*. The universal claim does not come from listing many domains. It comes from showing that the same degraded- observation semantics survive contact with distinct action geometries.

B. Autonomous vehicles

The AV row treats degraded observation as a time-to-collision and entry-barrier hazard under bounded vehicle semantics. Its runtime object is not the full autonomy stack. It is the release decision under uncertain headway and environment state. The defended bounded-row result matters because it shows the ORIUS contract moving into a fast-reacting closed-loop domain without importing battery-specific actuation logic. The row is bounded intentionally: the claim is about governed braking or maneuver legality, not full autonomous-driving field closure.

C. Healthcare monitoring

The healthcare row is the clearest reminder that ORIUS is not only about actuation- heavy control. Here the

TABLE VI: Deployment validation scope for the active three-domain ORIUS draft.

Deployment surface	Governing artifact	Scope type	Current status	Bounded manuscript claim	Exact non-claim or gap
Battery witness runtime	reports/publication/dc3witness_replay.csv	witness_replay	bounded_reference	Battery supports the deepest runtime and theorem witness surface.	This is still not unrestricted deployment.
Autonomous-vehicle defended replay	reports/orius_av/nuplab/autonomous-vehicle_defended_replay	bounded_replay	defended_bounded	AV supports full runtime all-zip grouped nuPlan replay under the narrow brake-hold runtime contract.	It does not claim full autonomous-driving field closure.
Healthcare defended replay	reports/healthcare/runtime_governance_summary.csv	bounded_replay	defended_bounded	Healthcare supports bounded monitoring and alert-release claims.	It does not claim regulated clinical deployment or a stronger clinical-release theorem beyond bounded replay.
OOD and adversarial completeness	chapters/ch34_outside_explicit_evidence_text; reports/publication/adversarial_probing_robustness_table.csv	explicit_evidence_text; bounded_non_claim	bounded_non_claim	The monograph can discuss bounded active probing and non-claim discipline.	It does not claim universal adversarial completeness or unrestricted OOD safety.

TABLE VII: Cross-domain status for the professor-facing ORIUS draft.

Domain	Current row	What is currently closed	Current boundary
Battery Energy Storage	Witness row	adapter, replay, theorem witness, governance	none
Autonomous Vehicles	Defended bounded row	adapter, replay, governance, bounded runtime contract	brake-hold runtime only
Medical and Healthcare Monitoring	Defended bounded row	adapter, replay, governance, bounded monitoring contract	runtime-emitted certificate validity metrics on bounded replay

danger is intervention suppression or alert illegality under degraded vital-sign observation. The defended bounded-row result shows that the same runtime grammar still applies when the safety object is intervention legality rather than dispatch power. It also marks a clean scientific boundary: the paper is not claiming bedside autonomy certification, only monitored intervention legality under a bounded runtime contract.

X. EXPLICIT EVIDENCE BOUNDARY

The manuscript does not overrun repository truth. The current repo supports a universal runtime-safety architecture, but it does not support a flat claim of equal empirical maturity across all possible Physical-AI domains. The evidence boundary is therefore part of the scientific contribution, not a defensive footnote.

The tracked matrices make the boundary explicit. The domain-closure matrix, runtime-budget matrix, and universal validation summary all now say the same thing in the active branch: battery is the reference witness, while autonomous vehicles and healthcare are defended bounded rows under explicit contracts.

The paper therefore uses the following claim discipline:

- It may claim a universal safety layer for Physical AI under degraded observation.
- It may claim bounded validation for AV and healthcare under their current contracts.

It does not claim equal-depth universality across all possible domains, universal hardware closure, or unrestricted field deployment readiness. Those would outrun

the present repository evidence and the locked validation outputs.

The practical implication is straightforward. The manuscript can be ambitious about architecture, but it must stay conservative about closure. That is not a weakness in the ORIUS program. It is the reason the program remains defensible under serious review.

XI. CONCLUSION

ORIUS argues that the right universal object for Physical AI safety is not a single controller, a single benchmark, or a single proof family. It is a runtime safety layer that mediates degraded observation before actuation. Once telemetry reliability becomes part of the safety problem, the system must expose reliability, uncertainty inflation, tightened action legality, repair, fallback, and certificate semantics as first-class objects. That is the contribution of the Detect-Calibrate-Constrain-Shield-Certify kernel.

The manuscript makes that contribution at three levels. First, it defines OASG as the dominant hidden hazard: actions can remain legal on the observed state while becoming unsafe on the true state. Second, it shows how a universal runtime contract can bind adapter semantics, benchmark semantics, and CertOS governance under one typed layer. Third, it demonstrates that the architecture is more than a diagram: the battery row provides witness-depth theorem-to-code-to-artifact closure, and the defended bounded rows show that the same runtime grammar travels beyond energy.

TABLE VIII: Allowed claims versus non-claims at the current evidence boundary.

Claim family	Supported by current repo truth	Not supported yet
Universal runtime grammar	One adapter contract, one benchmark schema, one CertOS lifecycle, and one degraded-observation hazard model across all rows.	A claim that all rows have identical proof depth or equal closure.
Battery witness depth	Deep theorem-to-code-to-artifact lineage with locked reference artifacts.	Treating battery as the conceptual center of the paper.
AV / healthcare	Defended bounded rows with locked replay and bounded runtime semantics.	Unqualified claims of unrestricted deployment readiness.

The scientific boundary is intentionally explicit. The paper does not claim equal-depth closure beyond the defended three-domain lane. Battery is the witness row, and AV plus Healthcare are bounded defended rows. Future additional domains remain outside the current defended evidence package. That boundary does not weaken the architectural thesis. It shows that ORIUS is disciplined enough to separate universality of contract from the present evidence depth.

What changes at the field level is the object of research. ORIUS suggests that Physical AI safety should be organized around reusable runtime layers that sit between nominal intelligence and physical release. Such a layer can host conformal uncertainty, safety filters, runtime assurance, and anomaly detection without collapsing into any one of them. If future additional domains clear the same obligations, ORIUS supports a broader empirical claim. Even before that point, it establishes a credible safety-layer grammar for Physical AI under degraded observation.

In that sense, ORIUS is also a writing and evaluation proposal for the field. It argues that future physical-AI systems should report not only predictive accuracy or planner reward, but also true-state violation behavior, intervention frequency, fallback usage, certificate validity, and release latency under degraded observation. That reporting discipline matters because many failures in real systems are not failures of nominal optimization; they are failures of release semantics once observation ceases to be trustworthy. A framework like ORIUS makes those failures measurable and governable.

REFERENCES

- [1] V. Vovk, A. Gammerman, and G. Shafer, “Algorithmic learning in a random world,” 2005.
- [2] J. Lei, M. G’Sell, A. Rinaldo, R. J. Tibshirani, and L. Wasserman, “Distribution-free predictive inference for regression,” 2018.
- [3] Y. Romano, E. Patterson, and R. J. Candès, “Conformalized quantile regression,” 2019.
- [4] I. Gibbs and E. J. Candès, “Adaptive conformal inference under distribution shift,” 2021.
- [5] A. N. Angelopoulos and S. Bates, “Conformal prediction: A gentle introduction,” 2023.
- [6] R. F. Barber, E. J. Candès, A. Ramdas, and R. J. Tibshirani, “Conformal prediction beyond exchangeability,” 2023.
- [7] Seto, Daisuke and Krogh, Bruce H. and Sha, Lui and Chutinan, Alongkritt, “The simplex architecture for safe online control system upgrades,” 1998, proceedings of ACC.
- [8] Sha, Lui, “Using simplicity to control complexity,” 2001, iEEE Software.
- [9] Desai, Ankush and Ghosal, Akash and Saha, Indranil and Yang, Yuhong and Dutle, Aaron and Ivanov, Radoslav and Lee, Insup and Seshia, Sanjit A. and Tiwari, Ashish, “Soter: A runtime assurance framework for programming safe robotics systems,” 2019, proceedings of IROS.
- [10] R. Bloem, B. Könighofer, R. Könighofer, and C. Wang, “Shield synthesis: Runtime enforcement for reactive systems,” 2015.
- [11] Bartocci, Ezio and Falcone, Ylies and Francalanza, Adrian and Reger, Giles, “Introduction to runtime verification,” 2018, lecture Notes in Computer Science.
- [12] M. Leucker and C. Schallhart, “A brief account of runtime verification,” 2009.
- [13] A. D. Ames, S. Coogan, M. Egerstedt, G. Notomista, K. Sreenath, and P. Tabuada, “Control barrier functions: Theory and applications,” 2019.
- [14] A. D. Ames, J. W. Grizzle, and P. Tabuada, “Control barrier function based quadratic programs with application to adaptive cruise control,” 2014.
- [15] K. P. Wabersich and M. N. Zeilinger, “A predictive safety filter for learning-based control of constrained nonlinear dynamical systems,” 2021.
- [16] Fisac, Jaime F. and Akametalu, Anayo and Zeilinger,

- Melanie N. and Kaynama, Sahar and Gillula, Jeremy and Tomlin, Claire J., “A general safety framework for learning-based control in uncertain robotic systems,” 2019, IEEE Transactions on Automatic Control.
- [17] D. Q. Mayne, M. M. Seron, and S. V. Raković, “Robust model predictive control of constrained linear systems with bounded disturbances,” 2005.
- [18] J. B. Rawlings, D. Q. Mayne, and M. M. Diehl, “Model predictive control: Theory, computation, and design,” 2017.
- [19] Basseville, Michele and Nikiforov, Igor V., “Detection of abrupt changes: Theory and application,” 1993, prentice Hall.
- [20] Page, E. S., “Continuous inspection schemes,” 1954, biometrika.
- [21] —, “A test for a change in a parameter occurring at an unknown point,” 1955, biometrika.
- [22] Hinkley, David V., “Inference about the change-point in a sequence of random variables,” 1971, biometrika.
- [23] Pasqualetti, Fabio and Dorfler, Florian and Bullo, Francesco, “Attack detection and identification in cyber-physical systems,” 2013, IEEE Transactions on Automatic Control.
- [24] A. Teixeira, I. Shames, H. Sandberg, and K. H. Johansson, “A secure control framework for resource-limited adversaries,” 2015.
- [25] M. Zaffran, O. Féron, Y. Goude, J. Jolivet, and A. Dieuleveut, “Adaptive conformal predictions for time series,” 2022.
- [26] Havelund, Klaus and Rosu, Grigore, “An overview of the runtime verification tool java pathexplorer,” 2004, formal Methods in System Design.
- [27] W. Langson, I. Chrysoschoos, S. V. Rakočić, and D. Q. Mayne, “Robust model predictive control using tubes,” 2004.
- [28] Camacho, Eduardo F. and Bordons, Carlos, “Model predictive control,” 2013, springer.
- [29] Breunig, Markus M. and Kriegel, Hans-Peter and Ng, Raymond T. and Sander, Jorg, “Lof: Identifying density-based local outliers,” 2000, proceedings of SIGMOD.
- [30] F. T. Liu, K. M. Ting, and Z.-H. Zhou, “Isolation forest,” 2008.
- [31] Y. Romano, E. Patterson, and R. J. Candès, “Conformalized quantile regression,” 2019.
- [32] I. Gibbs and E. J. Candès, “Adaptive conformal inference under distribution shift,” 2021.
- [33] V. Vovk, A. Gammerman, and G. Shafer, “Algorithmic learning in a random world,” 2005.
- [34] T. Chen and C. Guestrin, “Xgboost: A scalable tree boosting system,” 2016.
- [35] G. Ke *et al.*, “Lightgbm: A highly efficient gradient boosting decision tree,” 2017.
- [36] B. N. Oreshkin, D. Carpio, N. Chapados, and Y. Bengio, “N-beats: Neural basis expansion analysis for interpretable time series forecasting,” 2020.
- [37] B. Lim, S. Arık, N. Loeff, and T. Pfister, “Temporal fusion transformers for interpretable multi-horizon time series forecasting,” 2021.
- [38] Y. Nie *et al.*, “A time series is worth 64 words: Long-term forecasting with transformers,” 2023.
- [39] A. Wachi, Y. Sui, Y. Yue, and M. Ono, “Safe reinforcement learning in constrained markov decision processes,” 2024.
- [40] D. Amodei *et al.*, “Concrete problems in ai safety,” 2016.
- [41] A. Bemporad and M. Morari, “Robust model predictive control: A survey,” 2007.
- [42] A. A. Cardenas, S. Amin, Z.-S. Lin, Y.-L. Huang, C.-Y. Huang, and S. Sastry, “Attacks against process control systems: Risk assessment, detection, and response,” 2011.
- [43] Y. Liang, C.-Z. Zhao, H. Yuan, Z. Chen, W. Zhang, J.-Q. Huang, D. Yu, Y. Liu, M.-M. Titirici, Y.-L. Chueh, H. Yu, and Q. Zhang, “A review of rechargeable batteries for portable electronic devices,” 2019.
- [44] A. Shapiro, D. Dentcheva, and A. Ruszczyński, “Lectures on stochastic programming: Modeling and theory,” 2021.
- [45] H. Sun, J. Chen, Q. Ye, and J. Shao, “Adversarial attacks on neural networks for logical reasoning,” 2007.
- [46] C. Xu and Y. Xie, “Conformal prediction interval for dynamic time-series,” 2021.
- [47] J. García and F. Fernández, “A comprehensive survey on safe reinforcement learning,” 2015.
- [48] R. J. Tibshirani, B. Efron, R. F. Barber, E. J. Candès, and A. Ramdas, “Conformal prediction under covariate shift,” 2019.
- [49] H. Papadopoulos, K. Proedrou, V. Vovk, and A. Gammerman, “Inductive confidence machines for regression,” 2002.
- [50] K. Stankeviciute, A. M’Charrak, R. Cornish, and T. Rainforth, “Conformal time-series forecasting,” 2021.
- [51] S. Bates, E. J. Candès, L. Lei, Y. Romano, and M. Sesia, “Testing for outliers with conformal p-values,” 2023.
- [52] S. Prajna and A. Jadbabaie, “Safety verification of hybrid systems using barrier certificates,” 2004.
- [53] A. J. Taylor and A. D. Ames, “Adaptive safety with control barrier functions,” 2020.
- [54] P. Wieland and F. Allgöwer, “Constructive safety using control barrier functions,” 2007.
- [55] S. V. Rakočić, E. C. Kerrigan, K. I. Kouramas, and D. Q. Mayne, “Invariant approximations of the minimal robust positively invariant set,” 2005.
- [56] M. Lorenzen, F. Dabbene, R. Tempo, and F. Allgöwer, “Constraint-tightening and stability in stochastic model predictive control,” 2019.
- [57] J. Achiam, D. Held, A. Tamar, and P. Abbeel, “Constrained policy optimization,” 2017.
- [58] L. Brunke, M. Greeff, A. W. Hall, Z. Yuan, S. Zhou,

- J. Panerati, and A. P. Schoellig, “Safe learning in robotics: From learning-based control to safe reinforcement learning,” 2022.
- [59] T. M. Moldovan and P. Abbeel, “Safe exploration in markov decision processes,” 2012.
- [60] A. Ray, J. Achiam, and D. Amodei, “Benchmarking safe exploration in deep reinforcement learning,” 2010.
- [61] E. Altman, “Constrained markov decision processes,” 1999.
- [62] Y. Falcone, K. Havelund, and G. Reger, “A tutorial on runtime verification,” 2013.
- [63] E. Bartocci, Y. Falcone, A. Francalanza, and G. Reger, “Introduction to runtime verification,” 2018.
- [64] A. Pnueli and A. Zaks, “Psl model checking and runtime verification via testers,” 2006.
- [65] O. Maler and D. Nickovic, “Monitoring temporal properties of continuous signals,” 2004.
- [66] D. I. Urbina, J. A. Giraldo, A. A. Cardenas, N. O. Tippenhauer, J. Valente, M. Faisal, J. Ruths, R. Candell, and H. Sandberg, “Limiting the impact of stealthy attacks on industrial control systems,” 2016.
- [67] Y. Mo and B. Sinopoli, “Secure estimation in the presence of integrity attacks,” 2015.
- [68] J. A. Giraldo, D. Urbina, A. A. Cardenas, J. Valente, M. Faisal, J. Ruths, N. O. Tippenhauer, H. Sandberg, and R. Candell, “A survey of physics-based attack detection in cyber-physical systems,” 2018.
- [69] X. Hu, S. Li, and H. Peng, “A comparative study of equivalent circuit models for li-ion batteries,” 2012.
- [70] J. M. Tarascon and M. Armand, “Issues and challenges facing rechargeable lithium batteries,” 2001.
- [71] B. Dunn, H. Kamath, and J.-M. Tarascon, “Electrical energy storage for the grid: A battery of choices,” 2011.
- [72] S. Bose, S. Fattahi, and A. Swami, “Stochastic battery operation in a competitive electricity market,” 2019.
- [73] D. M. Rosewater, D. A. Copp, T. A. Nguyen, R. H. Byrne, and S. Santoso, “Battery energy storage models for optimal control,” 2019.
- [74] T. Hong and S. Fan, “Probabilistic electric load forecasting: A tutorial review,” 2016.
- [75] T. Gneiting and A. E. Raftery, “Strictly proper scoring rules, prediction, and estimation,” 2007.
- [76] S. Makridakis, E. Spiliotis, and V. Assimakopoulos, “The m4 competition: 100,000 time series and 61 forecasting methods,” 2020.
- [77] J. W. Taylor, “Short-term electricity demand forecasting using double seasonal exponential smoothing,” 2003.
- [78] P. Koopman and M. Wagner, “Autonomous vehicle safety: An interdisciplinary challenge,” 2016.
- [79] S. Shalev-Shwartz, S. Shammah, and A. Shashua, “On a formal model of safe and scalable self-driving cars,” 2017.
- [80] M. Althoff and J. M. Dolan, “Online verification of automated road vehicles using reachability analysis,” 2014.
- [81] T. Nolte, M. Sjodin, K. Akesson, K. Lundqvist, M. Behnam, M. Lindberg, and C. Normann, “Towards bringing research on functional safety in automotive systems to practice,” 2017.
- [82] I. Lee and O. Sokolsky, “Medical cyber physical systems,” 2010.
- [83] M. Saeed, M. Villarroel, A. T. Reisner, G. Clifford, L. W. Lehman, G. Moody, T. Heldt, T. H. Kyaw, B. Moody, and R. G. Mark, “Multiparameter intelligent monitoring in intensive care ii (mimic-ii): A public-access intensive care unit database,” 2011.
- [84] G. D. Clifford, W. J. Long, G. B. Moody, and P. Szolovits, “Robust parameter extraction for decision support using multimodal intensive care data,” 2009.
- [85] U.S. Food and Drug Administration, “Software as a medical device (samd): Clinical evaluation,” 2019.
- [86] RTCA, “Do-178c: Software considerations in airborne systems and equipment certification,” 2011.
- [87] J. Rushby, “Formal methods and the certification of critical systems,” 1993.
- [88] I. Habli and T. Kelly, “Model-based assurance of control system safety and security,” 2010.
- [89] L. P. Kaelbling, M. L. Littman, and A. R. Cassandra, “Planning and acting in partially observable stochastic domains,” 1998.
- [90] S. J. Julier and J. K. Uhlmann, “Unscented filtering and nonlinear estimation,” 2004.
- [91] R. T. Rockafellar and S. Uryasev, “Optimization of conditional value-at-risk,” 2000.
- [92] S. J. Qin and T. A. Badgwell, “A survey of industrial model predictive control technology,” 2003.
- [93] S. Yin, S. X. Ding, X. Xie, and H. Luo, “A review on basic data-driven approaches for industrial process monitoring,” 2014.
- [94] R. Isermann, “Model-based fault-detection and diagnosis—status and applications,” 2005.
- [95] V. Chandola, A. Banerjee, and V. Kumar, “Anomaly detection: A survey,” 2009.
- [96] A. B. Tsybakov, “Introduction to nonparametric estimation,” 2009.
- [97] B. Yu, “Assouad, fano, and le cam,” 1997.
- [98] M. Blanke, M. Kinnaert, J. Lunze, and M. Staroswiecki, “Diagnosis and fault-tolerant control,” 2006.
- [99] Y. Zhang and J. Jiang, “Bibliographical review on reconfigurable fault-tolerant control systems,” 2008.
- [100] L. Lindemann, M. Cleaveland, G. Shim, and G. J. Pappas, “Safe planning in dynamic environments using conformal prediction,” 2023.
- [101] A. Dixit, L. Lindemann, S. X. Wei, M. Cleaveland, G. J. Pappas, and J. V. Deshmukh, “Adaptive conformal prediction for motion planning among dynamic agents,” 2023.
- [102] A. N. Angelopoulos, S. Bates, E. J. Candès, M. I. Jor-

- dan, and L. Lei, "Conformal risk control," 2024.
- [103] Xu, Chen and Xie, Yao, "Conformal prediction interval for dynamic time-series," 2021, proceedings of ICML.
- [104] Sesia, Matteo and Candes, Emmanuel J. and Romano, Yaniv, "A comparison of some conformal quantile regression methods," 2021, stat.
- [105] Bates, Stephen and Candes, Emmanuel J. and Lei, Lihua and Romano, Yaniv, "Distribution-free, risk-controlling prediction sets," 2021, journal of the ACM.
- [106] Chernozhukov, Victor and Wuthrich, Kaspar and Zhu, Yinchu, "Distributional conformal prediction," 2021, proceedings of the National Academy of Sciences.
- [107] Feldman, Shai and Bates, Stephen and Candes, Emmanuel J., "Calibrated multiple-output quantile regression with conformal prediction," 2021, arXiv preprint.
- [108] Fisch, Adam and Grimson, W. Eric and Vovk, Vladimir and Katz-Samuels, Julian, "Online conformal prediction with decaying step sizes," 2021, arXiv preprint.
- [109] Schierman, John D. and Wang, Xun J. and Griffith, Dustin and DeCastro, Jonathan A. and Dutle, Aaron and Pike, Lee and Johnson, Taylor T., "Runtime assurance for complex systems," 2015, aIAA Infotech@Aerospace.
- [110] Bak, Stanley and Manamcheri, Krishna and Mitra, Sayan and Caccamo, Marco, "Sandboxing controllers for cyber-physical systems," 2011, proceedings of ICCPS.
- [111] Cofer, Darren and Rangarajan, Anand and Janssen, Wil and Havelund, Klaus and Perez, Ivan, "Assurance cases for runtime assurance frameworks," 2012, nASA Technical Report.
- [112] Herbert, Sylvia and Chen, Mo and Han, Siyuan and Bansal, Somil and Fisac, Jaime F. and Tomlin, Claire, "Fastrack: A modular framework for fast and guaranteed safe motion planning," 2017, proceedings of CDC.
- [113] Sen, Koushik and Vardhan, Abhay and Agha, Gul and Rosu, Grigore, "Efficient decentralized monitoring of safety in distributed systems," 2004, proceedings of ICSE.
- [114] d'Angelo, Ben and Sankaranarayanan, Sriram and Sanchez, Cesar and Robinson, Will and Finkbeiner, Bernd and Sipma, Henny and Mehrotra, Sharad and Manna, Zohar, "Lola: Runtime monitoring of synchronous systems," 2005, tIME.
- [115] Nguyen, Quan and Sreenath, Koushil, "Exponential control barrier functions for enforcing high relative-degree safety-critical constraints," 2016, proceedings of ACC.
- [116] Xu, Xiangru and Tabuada, Paulo and Grizzle, Jessy W. and Ames, Aaron D., "Robustness of control barrier functions for safety critical control," 2015, iFAC-PapersOnLine.
- [117] Wang, Li and Ames, Aaron D. and Egerstedt, Magnus, "Safety barrier certificates for collisions-free multirobot systems," 2017, iEEE Transactions on Robotics.
- [118] Cheng, Richard and Orosz, Gyorgy and Murray, Richard M. and Burdick, Joel W., "End-to-end safe reinforcement learning through barrier functions for safety-critical continuous control tasks," 2019, proceedings of AAAI.
- [119] Cheng, Richard and Nurnberger, Andrew and Scheibler, Karol and Ivanov, Radoslav and Pappas, George J., "Cautious adaptation for reinforcement learning under safety constraints," 2021, proceedings of L4DC.
- [120] Wabersich, Kai P. and Zeilinger, Melanie N. and Hewing, Lukas and Carron, Andrea, "Data-driven safety filters: A survey and perspective," 2023, annual Review of Control, Robotics, and Autonomous Systems.
- [121] Ben-Tal, Aharon and El Ghaoui, Laurent and Nemirovski, Arkadi, "Robust optimization," 2009, princeton University Press.
- [122] Koller, Torsten and Berkenkamp, Felix and Turchetta, Matteo and Krause, Andreas, "Learning-based model predictive control for safe exploration," 2018, proceedings of CDC.
- [123] Berkenkamp, Felix and Turchetta, Matteo and Schoellig, Angela P. and Krause, Andreas, "Safe model-based reinforcement learning with stability guarantees," 2017, advances in Neural Information Processing Systems.
- [124] Mo, Yilin and Sinopoli, Bruno, "False data injection attacks in control systems," 2012, proceedings of the First Workshop on Secure Control Systems.
- [125] Urbina, David I. and Giraldo, Jairo and Cardenas, Alvaro A. and Tippenhauer, Nils Ole, "Survey and new directions for physics-based attack detection in control systems," 2016, proceedings of CPS-SPC.
- [126] Khraisat, Ansam and Gondal, Iqbal and Vamplew, Peter and Kamruzzaman, Joarder, "Survey of intrusion detection systems: Techniques, datasets and challenges," 2019, cybersecurity.
- [127] Salinas, David and Flunkert, Valentin and Gasthaus, Jan and Januschowski, Tim, "Deepar: Probabilistic forecasting with autoregressive recurrent networks," 2020, international Journal of Forecasting.
- [128] Wu, Haixu and Xu, Jiehui and Wang, Jianmin and Long, Mingsheng, "Autoformer: Decomposition transformers with auto-correlation for long-term series forecasting," 2021, advances in Neural Information Processing Systems.
- [129] Zhou, Haoyi and Zhang, Shanghang and Peng, Jieqi and Zhang, Shuai and Li, Jianxin and Xiong, Hui and Zhang, Wancai, "Informer: Beyond efficient transformer for long sequence time-series forecasting," 2021, proceedings of AAAI.
- [130] Zhou, Tian and Ma, Ziqing and Wen, Qingsong and Wang, Xue and Sun, Liang and Jin, Rong and

- Zhou, Xiaokang and Qian, Wending and Li, Xue and Zhang, Weifeng and others, "Fedformer: Frequency enhanced decomposed transformer for long-term series forecasting," 2022, proceedings of ICML.
- [131] Wu, Haixu and Hu, Tao and Liu, Yong and Zhou, Hang and Wang, Jianmin and Long, Mingsheng, "Timesnet: Temporal 2d-variation modeling for general time series analysis," 2023, proceedings of ICLR.
- [132] Zeng, Ailing and Chen, Muxin and Zhang, Lei and Xu, Qiang, "Are transformers effective for time series forecasting?" 2023, proceedings of AAAI.
- [133] Dosovitskiy, Alexey and Beyer, Lucas and Kolesnikov, Alexander and others, "An image is worth 16x16 words: Transformers for image recognition at scale," 2021, proceedings of ICLR.
- [134] Vaswani, Ashish and Shazeer, Noam and Parmar, Niki and Uszkoreit, Jakob and Jones, Llion and Gomez, Aidan and Kaiser, Lukasz and Polosukhin, Illia, "Attention is all you need," 2017, advances in Neural Information Processing Systems.
- [135] Severson, Kristen A. and Attia, Peter M. and Jin, Norman and Perkins, Nicholas and Jiang, Ben and Yang, Zi and Chen, Michael and Aykol, Muratahan and Herring, Patrick K. and Fraggedakis, Dimitrios and others, "Data-driven prediction of battery cycle life before capacity degradation," 2019, nature Energy.
- [136] Berecibar, Mikel and Gandiaga, Iker and Villarreal, Ibon and Omar, Noshin and Van Mierlo, Joeri and Van den Bossche, Peter, "Critical review of state of health estimation methods of li-ion batteries for real applications," 2016, renewable and Sustainable Energy Reviews.
- [137] He, Hongwen and Xiong, Rui and Fan, Jianxin, "Evaluation of lithium-ion battery equivalent circuit models for state of charge estimation," 2011, applied Energy.
- [138] Hu, Xiaosong and Li, Shengbo Eben and Peng, Huei, "A comparative study of equivalent circuit models for li-ion batteries," 2012, journal of Power Sources.
- [139] Weitzel, Tim and Glock, Christoph H., "Energy storage dispatch in energy markets: A review," 2018, energy.
- [140] Xu, Bing and Oudalov, Alexandre and Ulbig, Andreas and Andersson, Goran and Kirschen, Daniel S., "Modeling of lithium-ion battery degradation for cell life assessment," 2020, IEEE Transactions on Smart Grid.
- [141] Wang, Yi and Zhang, Dong and Li, Xinyu, "Battery storage operation under uncertainty: A distributionally robust perspective," 2022, IEEE Transactions on Smart Grid.
- [142] Hossain, Eklas and Faruque, Hasan and Sunny, Md Shahjalal and Mohammad, Naimul and Nawar, Nurul, "A comprehensive review on grid-connected battery energy storage systems," 2020, energies.
- [143] Plett, Gregory L., "Battery management systems, volume ii: Equivalent-circuit methods," 2015, artech House.
- [144] —, "Battery management systems, volume iii: Battery state estimation," 2016, artech House.
- [145] Srinivasan, Ranjani and Sanner, Scott and others, "Runtime assurance for safe learning-enabled aerospace systems," 2020, aIAA Scitech.
- [146] Herbert, Sylvia and Chen, Mo and Han, Siyuan and Bansal, Somil and Fisac, Jaime and Tomlin, Claire, "Fastrack: A modular framework for fast and guaranteed safe motion planning," 2017, proceedings of CDC.
- [147] Ames, Aaron D. and Molnar, Tamas and Orosz, Gabor and Sreenath, Koushil and Tabuada, Paulo, "Safety-critical control for autonomous systems," 2021, annual Review of Control, Robotics, and Autonomous Systems.
- [148] Lederer, Johannes and Kohler, Jonas and Berkenkamp, Felix and Zeilinger, Melanie N., "A survey on safe reinforcement learning," 2019, arXiv preprint.
- [149] Ulbrich, Simon and Reschka, Andreas and Rieken, Jonathan and Ernst, Sebastian and Bagschik, Gereon and Damm, Werner and Maurer, Markus, "Towards a functional system architecture for automated vehicles," 2015, proceedings of ITSC.
- [150] Schwarting, Wilko and Alonso-Mora, Javier and Rus, Daniela, "Planning and decision-making for autonomous vehicles," 2018, annual Review of Control, Robotics, and Autonomous Systems.
- [151] Badue, Claudine and Guidolini, Rafael and Carneiro, Raphael and Berriel, Rodrigo and Paixao, Thiago and Mutz, Filipe and de Paula Veronese, Luis and Oliveira-Santos, Thiago, "Self-driving cars: A survey," 2021, expert Systems with Applications.
- [152] Paszke, Adam and Gross, Sam and Massa, Francisco and others, "Pytorch: An imperative style, high-performance deep learning library," 2019, advances in Neural Information Processing Systems.
- [153] Sutton, Richard S. and Barto, Andrew G., "Reinforcement learning: An introduction," 2018, mIT Press.
- [154] Kulkarni, Saurabh and Sahoo, Sarmila and Mohanty, Sasmita, "Industrial process monitoring and fault diagnosis: A survey," 2022, computers and Chemical Engineering.
- [155] Yin, Shen and Ding, Steven and Xie, Xiaoxia and Luo, Haifeng, "A review on basic data-driven approaches for industrial process monitoring," 2014, IEEE Transactions on Industrial Electronics.
- [156] Chiang, Leo H. and Russell, Edwin L. and Braatz, Richard D., "Fault detection and diagnosis in industrial systems," 2001, springer.
- [157] Clifton, Lei and Clifton, David A. and Pimentel, Marco A. F. and Watkinson, Peter J. and Tarassenko, Lionel, "Gaussian processes for personalized e-health monitoring," 2012, IEEE Transactions on Biomedical Engineering.
- [158] Goldstein, Benjamin A. and Navar, Ann Marie and

- Pencina, Michael J. and Ioannidis, John P. A., "Opportunities and challenges in developing risk prediction models with electronic health records data," 2017, journal of the American Medical Informatics Association.
- [159] Johnson, Alistair E. W. and Pollard, Tom J. and Shen, Lu and Lehman, Li-wei H. and Feng, Mengling and Ghassemi, Mohammad and Moody, Benjamin and Szolovits, Peter and Celi, Leo Anthony and Mark, Roger G., "Mimic-iii, a freely accessible critical care database," 2016, scientific Data.
- [160] Esteva, Andre and Robicquet, Alexandre and Ram-sundar, Bharath and Kuleshov, Volodymyr and De-Pristo, Mark and Chou, Katie and Cui, Claire and Corrado, Greg and Thrun, Sebastian and Dean, Jeff, "A guide to deep learning in healthcare," 2019, nature Medicine.
- [161] Lee, Edward A., "The past, present and future of cyber-physical systems: A focus on models," 2015, sensors.
- [162] Rajkumar, Ragunathan and Lee, Insup and Sha, Lui and Stankovic, John, "Cyber-physical systems: The next computing revolution," 2010, proceedings of DAC.
- [163] Derler, Patricia and Lee, Edward A. and Sangiovanni-Vincentelli, Alberto, "Modeling cyber-physical systems," 2012, proceedings of the IEEE.
- [164] Baheti, Radhakisan and Gill, Helen, "Cyber-physical systems," 2011, the Impact of Control Technology.
- [165] Wolf, Mark T. and others, "Safety and trustworthiness in physical ai systems," 2018, aAAI Spring Symposium.
- [166] Gernstedt, Erik and Nilsson, Ola and Smith, Reid, "Physical ai systems: Architectures, risks, and safety layers," 2024, arXiv preprint.
- [167] Astrom, Karl Johan and Murray, Richard M., "Feed-back systems: An introduction for scientists and engineers," 2008, princeton University Press.
- [168] Zhou, Kemin and Doyle, John C. and Glover, Keith, "Robust and optimal control," 1996, prentice Hall.
- [169] Bertsekas, Dimitri P., "Dynamic programming and optimal control," 2005, athena Scientific.
- [170] UL Standards and Engagement, "UL 4600: Standard for safety for the evaluation of autonomous products," 2023, underwriters Laboratories standard.
- [171] National Institute of Standards and Technology, "Artificial intelligence risk management framework (ai rmf 1.0)," 2023, nIST AI governance framework.
- [172] International Organization for Standardization, "Iso 26262 road vehicles – functional safety," 2018, international functional-safety standard.
- [173] International Electrotechnical Commission, "Iec 61508 functional safety of electrical, electronic, and programmable electronic safety-related systems," 2010, cross-sector functional-safety standard.
- [174] Vovk, Vladimir and Gammernan, Alex and Shafer, Glenn, "Conformal predictors and confidence machines," 2005, algorithmic Learning in a Random World.
- [175] Shafer, Glenn and Vovk, Vladimir, "A tutorial on conformal prediction," 2008, journal of Machine Learning Research.
- [176] Barber, Rina Foygel and Candes, Emmanuel J. and Ramdas, Aaditya and Tibshirani, Ryan J., "Conformal prediction beyond exchangeability," 2023, annals of Statistics.
- [177] Kiyani, Shayan and Pappas, George J. and Roth, Aaron and Hassani, Hamed, "Decision theoretic foundations for conformal prediction: Optimal uncertainty quantification for risk-averse agents," in *Proceedings of the 42nd International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, vol. 267, 2025, pp. 30 943–30 965.
- [178] Kwon, Minjae and Ingebrand, Tyler and Topcu, Ufuk and Feng, Lu, "Runtime safety through adaptive shielding: From hidden parameter inference to provable guarantees," 2025, arXiv:2506.11033.
- [179] Alshiekh, Mohammad and Bloem, Roderick and Ehlers, Ruediger and Konighofer, Bettina and Niekum, Scott and Topcu, Ufuk, "Safe reinforcement learning via shielding," 2018, proceedings of AAAI.
- [180] Konighofer, Bettina and Alshiekh, Mohammad and Bloem, Roderick, "Shielding techniques for safe reinforcement learning," 2020, formal Methods in System Design.
- [181] Bastani, Osbert, "Safe reinforcement learning: A control-theoretic perspective," 2021, foundations and Trends in Machine Learning.
- [182] Bertsimas, Dimitris and Brown, David B. and Caramanis, Constantine, "Theory and applications of robust optimization," 2011, sIAM Review.
- [183] Birge, John R. and Louveau, Francois, "Introduction to stochastic programming," 2011, springer.
- [184] Borisov, Vadim and Leemann, Tobias and SeSSLer, Kathrin and Haug, Johannes and Pawelczyk, Martin and Kasneci, Gjergji, "Deep neural networks and tabular data: A survey," 2022, iEEE Transactions on Neural Networks and Learning Systems.
- [185] Chen, Min and Lu, Shengnan and Wang, Zhen and Wang, Xin, "A review of battery energy storage system safety and reliability," 2020, journal of Energy Storage.
- [186] Dibaji, Seyed Mohammad and Pirani, Mohsen and Flamholz, David and Annaswamy, Anuradha and Johansson, Karl Henrik and Chakraborty, Aranya, "A systems and control perspective of cps security," 2019, annual Reviews in Control.
- [187] Lamport, Leslie and Shostak, Robert and Pease, Marshall, "The byzantine generals problem," 1982, aCM Transactions on Programming Languages and Systems.
- [188] Lundberg, Scott M. and Lee, Su-In, "A unified approach to interpreting model predictions," 2017, advances in Neural Information Processing Systems.
- [189] Pineau, Joelle and Vincent-Lamarre, Philippe and

- Sinha, Koustuv and Lariviere, Vincent and Beygelzimer, Alina and d'Alche-Buc, Florence and Fox, Emily and Larochelle, Hugo, "Improving reproducibility in machine learning research," 2021, journal of Machine Learning Research.
- [190] Rosewater, David and Ferraro, Paul and Byrne, Raymond and Santoso, Surya, "Risk-averse battery dispatch under forecast error," 2020, IEEE Transactions on Smart Grid.
- [191] Leveson, Nancy, "Engineering a safer world: Systems thinking applied to safety," 2011, mIT Press.
- [192] Koopman, Philip and Wagner, Michael, "Challenges in autonomous vehicle testing and validation," 2016, SAE International Journal of Transportation Safety.
- [193] Dawid, A. Philip, "The well-calibrated bayesian," 1982, journal of the American Statistical Association.
- [194] Platt, John, "Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods," 1999, advances in Large Margin Classifiers.
- [195] Guo, Chuan and Pleiss, Geoff and Sun, Yu and Weinberger, Kilian Q., "On calibration of modern neural networks," 2017, proceedings of ICML.
- [196] Kuleshov, Volodymyr and Fenner, Nathan and Ermon, Stefano, "Accurate uncertainties for deep learning using calibrated regression," 2018, proceedings of ICML.
- [197] Vaicnavicius, Joris and Widmann, David and Andersson, Carl and Lindsten, Fredrik and Roll, Jacob and Sch"on, Thomas B., "Evaluating model calibration in classification," 2019, proceedings of AISTATS.
- [198] Nixon, Jeremy and Dusenberry, Michael W. and Zhang, Linchuan and Jerfel, Ghassen and Tran, Dustin, "Measuring calibration in deep learning," 2019, proceedings of CVPR Workshops.
- [199] Lee, Edward A., "Cyber physical systems: Design challenges," 2008, proceedings of ISORC.
- [200] Kopetz, Hermann, "Real-time systems: Design principles for distributed embedded applications," 2011, springer.
- [201] Alur, Rajeev, "Principles of cyber-physical systems," 2015, mIT Press.
- [202] Blanchini, Franco, "Set invariance in control," 1999, automatica.
- [203] Aubin, Jean-Pierre, "Viability theory," 1991, birkh"auser.
- [204] Mitchell, Ian M. and Bayen, Alexandre M. and Tomlin, Claire J., "A time-dependent hamilton-jacobi formulation of reachable sets for continuous dynamic games," 2005, IEEE Transactions on Automatic Control.
- [205] Baier, Christel and Katoen, Joost-Pieter, "Principles of model checking," 2008, mIT Press.
- [206] Clarke, Edmund M. and Grumberg, Orna and Peled, Doron, "Model checking," 1999, mIT Press.
- [207] Geiger, Andreas and Lenz, Philip and Urtasun, Raquel, "Are we ready for autonomous driving? the kitti vision benchmark suite," 2012, proceedings of CVPR.
- [208] Caesar, Holger and Bankiti, Varun and Lang, Alex H. and Vora, Sourabh and Liong, Venice Erin and Xu, Qiang and Krishnan, Anush and Pan, Yu and Baldan, Giancarlo and Beijbom, Oscar, "nusenes: A multimodal dataset for autonomous driving," 2020, proceedings of CVPR.
- [209] Ettinger, Scott and Cheng, Shuyang and Caine, Benjamin and Liu, Chenxi and Zhao, Hang and Pradhan, Saurabh and Chai, Yuning and Sapp, Benjamin and Qi, Charles R. and Zhou, Yin and others, "Large scale interactive motion forecasting for autonomous driving: The waymo open motion dataset," 2021, proceedings of ICCV.
- [210] Caesar, Holger and Kabzan, Juraj and Tan, Kok Seang and Fong, Whye Kit and Wolff, Eric and Lang, Alex and Fletcher, Luke and Beijbom, Oscar and Omari, Sammy, "nuplan: A closed-loop ml-based planning benchmark for autonomous vehicles," 2021, arXiv:2106.11810.
- [211] Karnchanachari, Napat and others, "Towards learning-based planning: The nuplan benchmark for real-world autonomous driving," 2024, arXiv:2403.04133.
- [212] Peng, Mingxing and Yao, Ruoyu and Guo, Xusen and Ma, Jun, "nuplan-r: A closed-loop planning benchmark for autonomous driving via reactive multi-agent simulation," 2025, arXiv:2511.10403.
- [213] Wilson, Benjamin and Qi, Wenda and Agarwal, Tushar and Lambert, John and Singh, Jitendra and Khandelwal, Siddhesh and Pan, Bowen and Wang, Cong and D'Sa, Sharath and Lee, Seung and others, "Argoverse 2: Next generation datasets for self-driving perception and forecasting," 2021, arXiv preprint.
- [214] Downs, James J. and Vogel, Ernest F., "A plant-wide industrial process control problem," 1993, computers and Chemical Engineering.
- [215] Rieth, Cory A. and Amsel, Ben D. and Tran, Rene and Cook, Miles B., "Revisiting the tennessee eastman process simulation benchmark for industrial machine learning," 2017, proceedings of the AIChE Annual Meeting.
- [216] Goldberger, Ary L. and Amaral, Luis A. N. and Glass, Leon and Hausdorff, Jeffrey M. and Ivanov, Plamen Ch. and Mark, Roger G. and Mietus, Joseph E. and Moody, George B. and Peng, C.-K. and Stanley, H. Eugene, "Physiobank, physiobank, and physionet: Components of a new research resource for complex physiologic signals," 2000, circulation.
- [217] Johnson, Alistair E. W. and Bulgarelli, Lucas and Shen, Lu and Gayles, Anna and Shammout, Ahmad and Horng, Steven and Pollard, Tom and Hao, Shirley and Moody, Benjamin and Gow, Brian and others, "Mimic-iv, a freely accessible electronic health record dataset," 2023, scientific Data.

- [218] Johnson, Alistair E. W. and Bulgarelli, Lucas and Pollard, Tom J. and Celi, Leo Anthony and Mark, Roger G., "Mimic-iv version 3.1," 2024, physioNet.
- [219] Huang, Yong and Yang, Zhongqi and Rahmani, Amir M., "Mimic-sepsis: A curated benchmark for modeling and learning from sepsis trajectories in the icu," 2025, iEEE BHI / OpenReview.
- [220] Pollard, Tom J. and Johnson, Alistair E. W. and Raffa, Jonathan D. and Celi, Leo Anthony and Mark, Roger G. and Badawi, Omar, "The eicu collaborative research database, a freely available multi-center database for critical care research," 2018, scientific Data.
- [221] Collins, Gary S. and Moons, Karel G. M. and Dhiman, Paula and Riley, Richard D. and Beam, Andrew L. and Van Calster, Ben and Ghassemi, Marzyeh and Liu, Xiaoxuan and Reitsma, Johannes B. and van Smeden, Maarten, "Tripod+ai statement: Updated guidance for reporting clinical prediction models that use regression or machine learning methods," 2024, *BMJ* 2024;385:e078378.
- [222] Moons, Karel G. M. and Damen, Johanna A. A. and Kaul, Tabea and Hooft, Lotty and Andaur Navarro, Constanza and Dhiman, Paula and Beam, Andrew L. and Van Calster, Ben and Celi, Leo Anthony, "Probast+ai: An updated quality, risk of bias, and applicability assessment tool for prediction models using regression or artificial intelligence methods," 2025, *BMJ* 2025;388:e082505.
- [223] Liu, Xiaoxuan and Rivera, Samantha Cruz and Moher, David and Calvert, Melanie J. and Denniston, Alastair K., "Reporting guidelines for clinical trial reports for interventions involving artificial intelligence: The consort-ai extension," 2020, *nature Medicine* 26, 1364-1374.
- [224] Vasey, Bonnie and Nagendran, Myura and Campbell, Bruce and Clifton, David A. and Collins, Gary S. and Denaxas, Spiros and Denniston, Alastair K. and Faes, Livia and Geerts, Bart and Ibrahim, Mohammed and Liu, Xiaoxuan and Mateen, Bilal A. and McCradden, Melissa D. and McLennan, Stuart and Morgan, Laura and Ordish, Johan and Rogers, Campbell and Saria, Suchi and Ting, Daniel S. W. and Watkinson, Peter and Weber, Wim and Wheatstone, Peter and McCulloch, Peter, "Reporting guideline for the early-stage clinical evaluation of decision support systems driven by artificial intelligence: Decide-ai," 2022, *nature Medicine* 28, 924-933.
- [225] STARD-AI Steering Group, "Stard-ai reporting guideline for diagnostic accuracy studies involving artificial intelligence," 2025, *nature Medicine*.
- [226] Saxena, Abhinav and Goebel, Kai, "Turbofan engine degradation simulation data set," 2008, nASA Ames Prognostics Data Repository.
- [227] Ramasso, Emmanuel and Saxena, Abhinav, "Performance benchmarking and analysis of prognostic methods for c-mapss datasets," 2014, *international Journal of Prognostics and Health Management*.
- [228] Sch" afer, Matthias and Strohmeier, Martin and Olive, Xavier and Smith, Matthew and Lenders, Vincent and Martinovic, Ivan, "Opensky report 2020: Analysing in-flight emergency events using big data," 2020, openSky Network Report.
- [229] Strohmeier, Martin and Sch" afer, Matthias and Lenders, Vincent and Martinovic, Ivan, "On perception and reality in wireless air traffic communication security," 2014, *proceedings of IEEE CNS*.
- [230] Peng, Roger D., "Reproducible research in computational science," 2011, *science*.
- [231] Sandve, Geir Kjetil and Nekrutenko, Anton and Taylor, James and Hovig, Eivind, "Ten simple rules for reproducible computational research," 2013, *pLoS Computational Biology*.
- [232] Wilson, Greg and Aruliah, D. A. and Brown, C. Titus and Chue Hong, Neil P. and Davis, Matt and Guy, Richard T. and Haddock, Steven H. D. and Huff, Kathryn D. and Mitchell, Ian M. and Plumbly, Mark D. and others, "Best practices for scientific computing," 2014, *pLoS Biology*.
- [233] Amershi, Saleema and Begel, Andrew and Bird, Christian and DeLine, Robert and Gall, Harald and Kamar, Ece and Nagappan, Nachiappan and Nushi, Besmira and Zimmermann, Thomas, "Software engineering for machine learning: A case study," 2019, *proceedings of ICSE-SEIP*.
- [234] Doshi-Velez, Finale and Kim, Been, "Towards a rigorous science of interpretable machine learning," 2017, *arXiv preprint*.
- [235] Rudin, Cynthia, "Stop explaining black box machine learning models for high stakes decisions and use interpretable models instead," 2019, *nature Machine Intelligence*.
- [236] Lipton, Zachary C., "The mythos of model interpretability," 2018, *queue*.
- [237] Russell, Stuart, "Human compatible: Artificial intelligence and the problem of control," 2019, *viking*.
- [238] Koopman, Philip and Wagner, Michael, "Autonomous vehicle safety: An interdisciplinary challenge," 2017, *iEEE Intelligent Transportation Systems Magazine*.
- [239] Favarò, Francesca M. and Eurich, Samantha O. and Nader, Nadine and Tripp, Matthew, "Examining accident reports involving autonomous vehicles in california," 2017, *pLoS ONE*.
- [240] Kalra, Nidhi and Paddock, Susan M., "Driving to safety: How many miles of driving would it take to demonstrate autonomous vehicle reliability?" 2016, *transportation Research Part A*.
- [241] Cummings, Mary L., "Managing the complexity of autonomous vehicle safety," 2017, *journal of Transportation Safety and Security*.
- [242] Dolan, John M. and others, "The field benchmark problem and evaluation standard for mobile robots," 1994, *aAAI Mobile Robot Workshop*.
- [243] Thrun, Sebastian and Montemerlo, Michael and

- Dahlkamp, Hendrik and Stavens, David and Aron, Alex and Diebel, James and Fong, Philip and Gale, John and Halpenny, Morgan and Hoffmann, Gabe and others, "Stanley: The robot that won the darpa grand challenge," 2006, journal of Field Robotics.
- [244] Montemerlo, Michael and Becker, Jan and Bhat, Surya and Dahlkamp, Hendrik and Dolgov, Dmitri and Ettinger, Scott and Haehnel, Dirk and Hilden, Tim and Hoffmann, Gabe and Huhnke, Burkhard and others, "Junior: The stanford entry in the urban challenge," 2008, journal of Field Robotics.
- [245] Broggi, Alberto and Medici, Paolo and Cardarelli, Emanuele and Cerri, Pietro, "The vislab intercontinental autonomous challenge: An extensive road test of autonomous driving technologies," 2013, iEEE Intelligent Transportation Systems Magazine.
- [246] LaValle, Steven M., "Planning algorithms," 2006, cambridge University Press.
- [247] Kochenderfer, Mykel J., "Decision making under uncertainty: Theory and application," 2015, mIT Press.
- [248] Rasmussen, Carl Edward and Williams, Christopher K. I., "Gaussian processes for machine learning," 2006, mIT Press.
- [249] Murphy, Kevin P., "Machine learning: A probabilistic perspective," 2012, mIT Press.
- [250] Jordan, Michael I. and Mitchell, Tom M., "Machine learning: Trends, perspectives, and prospects," 2015, science.
- [251] Hendrycks, Dan and Gimpel, Kevin, "A baseline for detecting misclassified and out-of-distribution examples in neural networks," 2017, international Conference on Learning Representations.
- [252] Hespanha, Joao P. and Naghshtabrizi, Payam and Xu, Yonggang, "A survey of recent results in networked control systems," 2007, proceedings of the IEEE.