

ORIOUS

Observation–Reality Integrity for Universal Safety

Runtime Safety Under Degraded Observation

Dissertation-Style Research Manuscript

Pratik Niroula

Minnesota State University, Mankato

CIT Major, Mathematics Minor

*Canonical ORIOUS manuscript built from one active thesis surface with
tiered cross-domain evidence and manifest-backed claim discipline.*

ORIOUS Safety Principle

“Physical AI safety is bounded by observation reliability.”

May 2026

Abstract

Physical AI systems do not act on reality directly. They act on measurements, filtered state estimates, delayed telemetry, and learned summaries of the world. When that observation surface degrades, an action can remain admissible on the observed state while already being unsafe on the true state. This dissertation treats that discrepancy as an action-semantics safety problem rather than as a prediction problem alone, and names it the observation–action safety gap (OASG).

Existing lines of work address important parts of this failure mode: conformal uncertainty quantifies predictive spread, safe control constrains actions, runtime assurance supervises release, and fault monitoring detects degraded telemetry. What is still missing is the integrated runtime object that connects those pieces at the moment of physical release. ORIUS, Observation–Reality Integrity for Universal Safety, is proposed as that missing object: a certificate-aware runtime safety layer for physical AI under degraded observation.

The dissertation develops ORIUS through three implementation-grounded surfaces. First, **DomainAdapter** makes the domain safety predicate, fallback semantics, and repair geometry explicit. Second, a universal benchmark schema turns degraded-observation replay into a common evaluation contract across domains. Third, CertOS provides certificate lifecycle, audit continuity, and release semantics so that repaired action can be governed as a runtime object rather than reported only as controller-side logic. The formal contribution is a domain-agnostic theorem ladder connecting degraded observation, reliability-conditioned uncertainty, repaired action, fallback, and certificate semantics under explicit assumptions. The quantitative safety guarantees are stated as finite-sample *marginal* coverage bounds rather than as per-observation conditional guarantees: they control aggregate true-state risk over the joint distribution of observations, faults, and calibration residuals.

Empirically, the dissertation establishes one deep witness row and two bounded promoted rows without changing what ORIUS means at the release boundary. Battery provides the strongest theorem-to-runtime closure. The autonomous-vehicle row is bounded to completed nuPlan replay artifacts under the TTC plus predictive-entry-barrier contract. Healthcare is bounded to retrospective MIMIC-backed monitoring and alert-release semantics. The result

is a thesis that is universal in architecture and tiered in evidence. The manuscript therefore separates what is proved, what is validated, and what remains outside the current public claim surface. That separation is enforced not only by prose but by a claim-locked publication surface in which runtime and governance tables, theorem audits, and manifest-backed figures govern what the dissertation is allowed to say.

Acknowledgments

This thesis was assembled as a solo research project around one question: how should a physical-AI system behave when the observation channel is no longer trustworthy, but the consequences of action remain physical and irreversible? ORIUS was built to answer that question with a runtime contract rather than a domain-specific controller.

The deepest current evidence still comes from the battery witness row because it offers the cleanest theorem-to-code-to-artifact surface in the repository. The broader contribution of the book, however, depends on carrying that safety grammar across bounded nuPlan AV replay and retrospective healthcare monitoring without letting the prose outrun the evidence.

The open-source systems, datasets, and research communities behind the methods cited in this book made that synthesis possible. Their work is visible throughout the bibliography, the artifact index, and the reviewer-facing appendices that helped harden the final manuscript.

Contents

Abstract	i
Acknowledgments	iii
I Problem, Principle, and Gap	1
1 Introduction and Thesis Claims	2
1.1 The hidden law of physical AI	2
1.2 Why degraded observation is a release-boundary problem	3
1.3 Thesis claim	4
1.4 Why this dissertation is a monograph, not a stitched paper bundle	4
1.5 Main contributions	5
1.6 Claim discipline	6
1.7 Why ORIUS matters beyond one witness domain	6
1.8 Dissertation organization	7
2 Related Work and the Novelty Gap	8
2.1 Conformal prediction and adaptive calibration	8
2.2 Runtime assurance and Simplex-style supervision	9
2.3 Safety filters, shields, and robust safe control	9
2.4 Anomaly, drift, and degraded-observation detection	10
2.5 Domain-specific cyber-physical safety systems	10
2.6 Safety-case governance and lifecycle discipline	10
2.7 Benchmark, replay, and reproducibility gaps	11
2.8 Why ORIUS is not another forecasting or perception paper	11
2.9 The novelty gap ORIUS addresses	12

3	Formal Problem Formulation	13
3.1	State, observation, action, and constraints	13
3.2	Observation degradation operator and fault process	13
3.3	Fault taxonomy	14
3.4	Definition of the observation–action safety gap	15
3.5	Observation-consistent state sets	15
3.6	Repair, fallback, and release	16
3.7	Dual-trajectory evaluation	16
3.8	DomainAdapter grounding	16
3.9	Scope and assumptions	17
II	ORIOUS Framework and Theorem Ladder	18
4	ORIOUS Architecture: DC3S Runtime Layer	19
4.1	Why a runtime layer is the right object	19
4.2	The DC3S release step	20
4.3	Universal runtime flow	20
4.4	DomainAdapter as the universal boundary	20
4.5	Relationship to supervision and safety filters	22
4.6	Certificate emission and operational semantics	23
4.7	Runtime cost and operational realism	23
4.8	Architectural takeaway	23
5	Mathematical Foundations and Safety Guarantees	24
5.1	Why the theorem surface is intentionally narrow	24
5.2	Notation and assumption discipline	25
5.3	Definition of OASG as a formal event	25
5.4	The paper-facing theorem surface	26
5.5	Risk envelope interpretation	28
5.6	Why the proofs stay witness-backed but the objects stay universal	28
5.7	Theorem-to-artifact linkage	28
5.8	What is proved, validated, and still open	29
6	Benchmark and Reproducibility Protocol	30
6.1	Why ORIOUS-Bench is central to the thesis	30
6.2	Dual-trajectory evaluation	30
6.3	Deterministic replay and fault schedules	31

6.4	Metrics that match the runtime object	31
6.5	Strict universal mode and legacy compatibility	31
6.6	Artifact provenance and claim locking	32
6.7	Reproducibility checklist	32
6.8	Benchmark takeaway	33
III	Witness Domain Evidence	34
7	Witness Domain Implementation	35
7.1	Why battery is the witness domain	35
7.2	Witness-domain problem and plant semantics	35
7.3	Data surfaces and dataset cards	36
7.4	Forecasting, calibration, and reliability surfaces	36
7.5	Control baselines and adapter implementation	38
7.6	Locked artifact protocol	39
7.7	Witness-row scope	39
8	Witness Results and Failure Analysis	40
8.1	The governing empirical question	40
8.2	Experiment design and baseline logic	40
8.3	Headline result	41
8.4	Mechanistic interpretation	41
8.5	Ablations and where the gain comes from	42
8.6	Grouped calibration and reliability slices	42
8.7	Operational traces and certificate-aware behavior	42
8.8	Bounded deep-learning novelty inside the witness row	43
8.9	What fails without the runtime layer	44
8.10	Witness-row claim boundary	44
IV	Universalization and Extensions	45
9	Universal ORIUS Across Domains	46
9.1	What universality means in this dissertation	46
9.2	A transfer recipe rather than a portability slogan	46
9.3	Cross-domain status at a glance	47
9.4	Final three-domain release results	47

9.5	Domain synthesis	49
9.6	Architecture universality versus promoted-row closure	51
9.7	Default proof mode for the dissertation	52
9.8	Boundary of the universality claim	52
10	Temporal Certificates and Graceful Degradation	53
10.1	Certificate horizon as a runtime object	53
10.2	Why temporal safety belongs inside runtime assurance	53
10.3	Expiration, half-life, and safe-hold	54
10.4	Severe-blackout protocol	54
10.5	Graceful degradation policies	55
10.6	Temporal extension	55
11	Compositional Safety	57
11.1	Why local safety does not automatically compose	57
11.2	A simple counterexample	57
11.3	Shared constraints as the right composition surface	57
11.4	Multi-agent certificates and negotiation	58
11.5	Why composition remains bounded in the dissertation	58
11.6	Compositional takeaway	58
12	CertOS Runtime Assurance	60
12.1	Why CertOS belongs in the main argument	60
12.2	Lifecycle operations	60
12.3	Runtime invariants	61
12.4	Audit chain and post hoc verification	61
12.5	Governance alignment, not compliance	62
12.6	Universality and evidence-tiered maturity	62
12.7	Governance boundary	62
V	Submission Boundary	63
13	Governance, Reproducibility, and Conclusion	64
13.1	What is proved	64
13.2	What is empirically validated	64
13.3	What is out of scope	65
13.4	Reproducibility discipline	66

13.5 Threats to validity	66
13.6 Conclusion	66
VI Archival Provenance Index	68
14 Archival Source Register	69
14.1 Tracked archive locations	69
14.2 Reader-facing rule	70
VII System Appendices	71
a Notation Sheet	72
a.1 Telemetry and State Variables	72
a.2 Quality, Drift, and Sensitivity Indicators	72
a.3 Uncertainty and Action Objects	73
a.4 Battery Plant Parameters	73
a.5 Metrics	74
a.6 Tuning Constants and Greek Letters	74
a.7 Set Notation and Operators	75
b Assumption Register	76
c Proofs	78
c.1 Proof map	78
c.2 Proof sketches	79
c.3 Interpretation	81
d Artifact and Claim Index	82
d.1 Claim classes and governing surfaces	82
d.2 Headline claim families	83
d.3 Canonical manuscript-facing evidence	83
d.4 Interpretation	84
e Safety-Case Bundle Index	85
e.1 How to read the bundle index	86
e.2 Interpretive boundary	86

f	Extended Battery Result Tables	87
f.1	Extended Controller Comparison	87
f.2	Metric Definitions	87
f.3	Regional Impact Breakdown	88
f.4	Fault-Profile Sensitivity	88
f.5	Artifact Provenance	88
g	Reliability-Group Coverage Audits	90
g.1	Audit Logic	90
g.2	Illustrative Audit Table	91
g.3	Coverage–Width Trade-off	91
g.4	CQR Group Coverage Visualisation	91
g.5	Audit Protocol	93
h	Benchmark Specifications	94
h.1	Canonical fault families	94
h.2	Combined profile: <code>drift_combo</code>	96
h.3	Replay and fault-injection implementation	96
h.4	Severity mapping	96
i	Battery Adapter Interface Specification	98
i.1	Interface Overview	98
i.2	Method Contracts	99
i.3	Pipeline Diagram	99
i.4	How to Implement a New Domain Adapter	99
j	Simulated 15-Reviewer Audit Protocol	101
j.1	Reviewer-Role Matrix	101
j.2	Master Vulnerability Register	103
j.3	Response / Concession Matrix	106
j.4	Final Disposition Summary	106
k	Audited Theorems and Gap Register	107
k.1	Theorem Statements	107
k.1.1	Theorem 1: OASG Existence	107
k.1.2	Theorem 2: One-Step Safety Preservation	108
k.1.3	Theorem 3a: ORIUS Core Envelope Derivation	108
k.1.4	Corollary 3b: ORIUS Core Aggregation Corollary	108

k.1.5	Theorem 4: No Free Safety	108
k.1.6	Definition 5: Certificate Validity Horizon	109
k.1.7	Theorem 6: Certificate Expiration Bound	109
k.1.8	Theorem 7: Feasible Fallback Existence	109
k.1.9	Theorem 8: Graceful Degradation Dominance	109
k.1.10	Theorem 9: Universal Impossibility	110
k.1.11	Theorem 10: Boundary-Indistinguishability Lower Bound	110
k.1.12	Theorem 11: Typed Structural Transfer	110
k.2	Gap Audit Table	111
k.3	Remaining Claim-Boundary Gaps	111
k.4	Locked Empirical Evidence	111
k.5	Integrated Release-Gate Status	111
l	Locked Artifact Hash Registry	116
m	Claim-to-Evidence and Coverage Registers	119
m.1	Primary Claim Matrix	119
m.2	Theorem and Chapter Register	121
m.3	Chapter Coverage Register	122
m.4	Full Claim–Evidence Matrix	125
n	CertOS Artifact Audit	127
n.1	Artifact Bundle	127
n.2	Lifecycle Counts by Status	128
n.3	Certificate Field Schema	128
n.4	Audit Ledger Record Format	128
n.5	Latency and Recovery Audit	128
n.6	Bounded Second-Domain Runtime Surface	129
o	Full CertOS Lifecycle Log	131
p	Dataset Cards	135
q	Reproducibility Appendix	136
q.1	Release Manifest Summary	136
q.2	Source Runs	136
q.3	Tracked Publication Artifacts	137

List of Tables

4.1	Domain state schema for ORIUS universal framework.	22
4.2	Safety predicates per domain.	22
4.3	Universal DC3S pipeline stages.	23
5.1	Counterexamples and assumption boundaries for the observation-ambiguity theorem.	27
5.2	DC3S Theorem Dependency and Proof Completeness. Each formal result is implemented in code and tested by at least one automated test.	27
6.1	Cross-Domain OASG (Observed-Action Safety Gap) Rates.	32
6.2	Claim-governing three-domain runtime evidence. These rows are the head-line safety rows allowed by the claim validators; validation-harness rows are diagnostic only.	32
7.1	Dataset-card inventory for the four evaluation regions.	36
7.2	Six-model forecast comparison (DE & US). GBM dominates across all targets. PICP/Width shown only for GBM (production conformal intervals). Negative R^2 indicates worse-than-mean predictions.	37
7.3	Battery DeepOQE diagnostics on the held-out degraded-telemetry episodes.	38
8.1	Main controller comparison under aggregate telemetry faults. DC ³ S variants achieve zero true-SoC violations across all seeds and scenarios. 95% CI via percentile bootstrap over seeds $\times$ scenarios.	41
8.2	Ablation study: DC ³ S violation-rate reduction vs. baselines under telemetry faults. Two-gate verification: relative reduction $\geq 10\%$ <i>and</i> Wilcoxon signed-rank $p < 0.01$	42
8.3	Regime-aware CQR coverage and interval width by target and volatility group.	43
9.1	Canonical Battery + AV + Healthcare closure matrix.	47

9.2	Canonical runtime release contract witnesses. Each promoted domain is normalized to the same release evidence schema: <code>observe>reliability>uncertainty>repair_or_fa</code>	
9.3	Final CPU release model quality used by the manuscript-facing ORIUS package.	48
9.4	Final claim-governing runtime safety evidence across the three promoted ORIUS domains.	48
9.5	Utility-preserving safety scorecard. ORIUS is compared with a degenerate fail-safe reference rather than only with unsafe persistence.	48
9.6	Final freeze and validation gates for the locked CPU release package.	48
11.1	Multi-Agent Fleet Coordination: joint safety violations and useful work under shared transformer capacity constraint (limit = 80 MW, two batteries at 60 MW each = 120 MW proposed). Independent protocol violates; ORIUS-coordinated eliminates violations.	58
13.1	Validation boundary for the current ORIUS evidence package.	65
a.1	Telemetry and state symbols.	72
a.2	Quality and detector symbols.	73
a.3	Uncertainty-set and action symbols.	73
a.4	Battery plant parameters.	74
a.5	Evaluation metrics.	74
a.6	Tuning constants and thresholds.	74
a.7	Set notation and mathematical operators.	75
b.1	ORIUS dissertation assumption register.	76
c.1	Theorem-to-proof map for the ORIUS dissertation.	78
d.1	Claim classes and their governing artifact surfaces.	82
d.2	Headline claim families and their canonical evidence surfaces.	83
e.1	Safety-case bundle index for the ORIUS dissertation.	85
f.1	Extended controller comparison (DE, drift_combo, 20 % dropout). Source: <code>dc3s_main_table.csv</code>	87
f.2	Regional impact breakdown for DC3S + FTIT. Source: <code>dc3s_main_table.csv</code> , <code>48h_trace_final_de.csv</code> , <code>48h_trace_final_us.csv</code>	88
f.3	DC3S + FTIT under individual fault families (DE, 20 % dropout). Source: <code>fault_performance_table.csv</code>	89

g.1	Coverage by reliability bin	91
h.1	Witness-domain fault families used to instantiate the universal ORIUS-Bench telemetry-degradation contract.	95
h.2	<code>drift_combo</code> default parameter values.	96
h.3	Fault-to-pipeline severity mapping.	97
i.1	<code>BatteryDomainAdapter</code> interface methods.	98
k.1	Historical T1–T11 theorem audit summary before promotion-gate closure . .	112
k.2	Audited extension laws and depth-theorem surfaces	113
k.3	Integrated 18-row hard-gated proof-surface traceability summary. The unique surface set contains the two supporting Chapter 4 statements and the T1–T8 surface with active defense tiers imported from the theorem audit; Appendix C restatements are checked against that unique surface.	114
k.4	Integrated theorem traceability-gate matrix.	114
l.1	Locked artifact hash registry from the final full-thesis package.	116
m.1	Primary claim-to-evidence matrix from the locked publication register. . . .	119
m.2	Registry-canonical defended-core theorem and chapter evidence register for the ORIUS defense package.	121
m.3	Chapter coverage register from the thesis coverage audit.	122
m.4	Claim–Evidence Matrix: all 22 thesis claims with locked artifact pointers. Status <i>locked</i> means the artifact file exists and contains real numerical data. .	126
n.1	CertOS artifact bundle and its role in the thesis.	127
n.2	Status counts derived from the locked CertOS summary artifact.	128
n.3	Certificate and runtime-state fields used by CertOS.	129
n.4	Audit ledger record fields from the CertOS audit-ledger module.	130
n.5	Latency and recovery audit from the CertOS runtime artifacts.	130
n.6	Vehicle-shaped CertOS toy artifact.	130
o.1	Full 96-step CertOS lifecycle log.	131
p.1	Source dataset cards and manifests for the active three-domain ORIUS claim surface.	135

List of Figures

1.1	ORIUS as a universal runtime safety layer between nominal intelligence and physical release. The five-stage kernel recurs across domains even though the domains do not share a plant model, nominal controller family, or optimization objective.	5
4.1	ORIUS runtime flow from degraded telemetry to repaired action and governed certificate. The figure is generated from the same publication assets that feed the manuscript-facing parity and governance surfaces.	21
5.1	The integrated theorem gate keeps theorem language aligned to typed contracts, assumptions, and locked artifacts rather than to informal prose alone.	29
7.1	Witness-domain reliability baselines. The engineered forecasting surface and the learned reliability surface are intentionally reported together so the dissertation does not confuse forecasting novelty with degraded-observation novelty.	38
8.1	Battery learned-reliability safety surface. The deep model is judged by true-state safety, intervention burden, and certificate-valid release behavior rather than by forecast error alone.	43
9.1	Final claim-governing runtime TSVR across Battery, AV, and Healthcare.	49
9.2	Final release-model PICP90 values for the manuscript-facing model-quality gate.	50
9.3	Useful work preserved by ORIUS over each domain’s degenerate fail-safe reference.	50
9.4	Cross-domain ORIUS validation surface. One runtime grammar carries across the active three-domain program, while each row remains limited by its own defended evidence surface.	51

10.1	Blackout half-life surface used to motivate certificate expiration, bounded safehold, and temporal intervention discipline.	54
10.2	Four-policy graceful degradation comparison. The dissertation uses this family to treat fallback quality as a benchmarked layer instead of an unexamined emergency default.	55
12.1	CertOS lifecycle and runtime-governance surface across ORIUS domains. The figure shows how certificate issue, validation, expiry, fallback, and audit continuity are exposed as part of runtime safety rather than as post hoc reporting.	61
g.1	Coverage vs. interval width by reliability bin. Lower reliability $\Rightarrow$ wider intervals $\Rightarrow$ higher coverage. Source: <code>fig_coverage_width.png</code>	92
g.2	CQR group coverage across reliability bins. All bins exceed the 90% target (dashed line). Source: <code>fig_coverage_width_tradeoff.png</code>	92

Part I

Problem, Principle, and Gap

Chapter 1

Introduction and Thesis Claims

A battery dispatch stack can continue releasing power while feeder telemetry has already gone stale enough to hide a true state-of-charge or network violation. A vehicle controller can preserve lane-centering behavior even after the scene estimate has fallen behind the physical traffic geometry. A bedside monitoring policy can remain internally consistent while its input stream smooths away the very transition that should have triggered intervention. These systems look different at the level of plant dynamics and domain semantics, but they share one hidden law: *physical AI acts on observations, not on reality*. Once the observation surface degrades, a system can continue to compute coherent actions while the physical basis for those actions has already failed.

This dissertation is organized around that release-boundary hazard. Its central claim is that physical-AI safety under degraded observation requires a dedicated runtime layer between nominal intelligence and actuation, and that this layer can be expressed through one reusable grammar: *Detect*, *Calibrate*, *Constrain*, *Shield*, and *Certify*. ORIUS is the implemented form of that grammar. It is not proposed as a universal controller, not as a replacement for domain-specific estimation or control, and not as a blanket claim that every supported domain has already reached equal empirical maturity. It is proposed as the right runtime object for deciding what action remains admissible once observation trust degrades.

1.1 The hidden law of physical AI

Physical AI systems never act on truth directly. They act on packets, sensor fusion products, state estimators, learned summaries, and communication layers that compress the world into a runtime observation object. That compression is unavoidable. It is what allows downstream forecasting, planning, optimization, and control to happen at all. But it also creates a standing safety debt: the legality of an action is usually evaluated on the observed

state even though the true physical state may already have diverged.

The dissertation therefore takes a position that differs from the standard performance-first framing. The deepest safety problem is not only forecast error, perception error, or nominal control error. It is the possibility that a system remains computationally coherent while its observation surface is no longer trustworthy enough to justify physical release. In ORIUS, that discrepancy is named the *observation–action safety gap* (OASG). The OASG is the primary scientific object because it marks the exact point at which internally legal behavior can become physically unsafe without obvious software failure.

The flagship novelty sentence used throughout the defense package is:

ORIUS provides a reliability-aware runtime safety layer for physical AI under degraded observation, enforcing certificate-backed action release through uncertainty coverage, repair, and fallback.

The novelty claim is architectural and runtime-semantic: ORIUS does not claim a new conformal method, a new robust-optimization primitive, a new universal controller, or a new conditional-coverage theorem.

1.2 Why degraded observation is a release-boundary problem

Most deployed stacks already include pieces of the response: anomaly scores, confidence thresholds, conservative margins, or fallback modes. What they often lack is a single runtime discipline that ties those pieces together. Detection alone is not enough, because a system that detects degraded telemetry must still decide how uncertainty changes. Wider uncertainty is not enough, because the admissible action set must be recomputed under that widened state set. Recomputed admissibility is not enough, because unsafe candidate actions must be repaired or replaced before release. And none of that is enough if the runtime cannot later explain why an action was allowed, for how long, and under which fallback or certificate state.

ORIUS is organized around that full loop. The contribution of the dissertation is therefore not “better forecasting” and not “better control” in isolation. It is a runtime safety layer that accepts the reality of degraded observation and then turns that reality into governed release semantics. This is why the central ORIUS question is not “what prediction model is best?” but “what action remains safe to release, with what uncertainty, with what repair, and under what certificate, once observation quality degrades?”

1.3 Thesis claim

The thesis claim can be stated compactly:

Physical AI under degraded observation requires a certificate-aware runtime safety layer that mediates between nominal intelligence and physical release; that layer can be expressed through one reusable contract of reliability assessment, uncertainty inflation, admissible-set tightening, repaired or fallback action, and governed certificate emission.

ORIOUS operationalizes that claim through a typed Detect–Calibrate–Constrain–Shield–Certify kernel.

- **Detect** scores the reliability of the observation channel.
- **Calibrate** widens the observation-consistent state set as reliability falls.
- **Constrain** recomputes admissible action under that widened uncertainty surface.
- **Shield** repairs unsafe candidate action or escalates to a fallback policy.
- **Certify** records why release remains admissible, for how long, and under what governance state.

The key point is architectural. ORIOUS is written as a post-nominal layer. It wraps inherited domain controllers rather than assuming they can all be redesigned from first principles. That choice is what makes the framework portable across the active Battery, AV, and Healthcare lane without pretending that those domains share one plant model or one optimal nominal controller.

1.4 Why this dissertation is a monograph, not a stitched paper bundle

The argument of the dissertation is structurally single-threaded. It begins with one hidden hazard, develops one runtime object, formalizes one theorem ladder, measures one benchmark contract, and then asks how far that same safety grammar travels across multiple domains. Because the argument is unified, the correct dissertation form is a doctrine-driven monograph rather than a stitched collection of papers. Battery remains the deepest witness row, but it is not the identity of the thesis. Its role is to provide the deepest theorem-to-code-to-artifact surface against which the broader universal claim can be calibrated.




Figure 1.1: ORIUS as a universal runtime safety layer between nominal intelligence and physical release. The five-stage kernel recurs across domains even though the domains do not share a plant model, nominal controller family, or optimization objective.

Three implementation-grounded objects make that monograph possible. First, the `DomainAdapter` contract fixes what every domain must expose: telemetry semantics, constraint predicates, repair hooks, and certificate payloads. Second, ORIUS-Bench fixes the dual-trajectory replay contract through which observed safety and true-state safety can be compared. Third, CertOS turns repaired action into a governed release object through lifecycle transitions, certificate horizons, and audit records. Taken together, those objects let the dissertation defend a reusable runtime layer rather than a loose set of case studies.

1.5 Main contributions

The dissertation makes four contributions.

1. **A problem contribution.** It defines degraded observation as a universal release-boundary hazard for physical AI through the observation–action safety gap.
2. **A systems contribution.** It presents ORIUS as a typed runtime layer that connects reliability assessment, uncertainty inflation, safe-set tightening, repair or fallback, and governed certificate emission.

3. **A formal contribution.** It develops a theorem bridge from reliability-aware uncertainty and repaired action to bounded reduction in true-state safety violation under explicit assumptions.
4. **An empirical contribution.** It assembles a claim-locked benchmark and multi-domain evidence program that distinguishes a deepest witness row from defended bounded peers, portability rows, and explicitly gated rows rather than flattening them into a false equal-maturity claim.

1.6 Claim discipline

The dissertation uses three claim classes and keeps them visibly separate from the first chapter forward.

Proved claims are theorem-grade statements under explicit assumptions about degraded observation, uncertainty inflation, repair, and bounded violation reduction.

Validated claims are empirical statements tied to replayable artifacts, runtime traces, and governed benchmark outputs.

Roadmap claims are structurally motivated extensions that are intentionally not promoted as defended evidence.

This distinction is not only stylistic caution. It prevents a universal architectural argument from being mistaken for equal empirical parity, and it keeps the manuscript aligned with the repository’s claim-lock discipline. Reviewers should be able to tell, at any point in the book, whether a statement is theorem-backed, artifact-backed, or explicitly left as future work.

1.7 Why ORIUS matters beyond one witness domain

The importance of ORIUS is not that batteries, vehicles, and healthcare monitoring systems secretly reduce to one plant. The importance is that they all face the same release-boundary question once observation trust degrades: what action can still be justified, with what uncertainty, with what fallback, and under what certificate? ORIUS freezes the semantics of that question so that cross-domain comparison happens at the level of runtime safety rather than only at the level of application area.

That framing also changes how the thesis relates to adjacent literature. Conformal prediction and adaptive calibration provide disciplined uncertainty surfaces [1–3]. Runtime assurance contributes supervisory release logic [4, 5]. Safety filters and barrier methods provide computable constraint-preserving action geometry [6–8]. Safety-case governance and risk management standards clarify why claims, lifecycle discipline, and auditable evidence matter even when formal compliance is not being claimed [9–12]. ORIUS brings those strands together by centering the missing object they all touch but do not fully stabilize: the runtime safety layer for degraded observation.

1.8 Dissertation organization

The rest of the dissertation follows the doctrine directly. Chapters 2 and 3 define the novelty gap and the formal OASG object. Chapters 4 through 6 present the ORIUS architecture, theorem ladder, and benchmark contract. Chapters 7 and 8 defend the battery witness row. Chapters 9 through 12 ask how far the same runtime grammar carries across domains, temporal certificate semantics, composition, and governance. Chapter 13 then states the precise submission boundary, reproducibility discipline, and final thesis claim.

Boundary of this chapter. This chapter claims only the dissertation’s governing doctrine and contribution structure. It does not yet prove the theorem ladder or validate broader future-domain closure; those later claims remain tied to the benchmark, artifact, and governance surfaces that follow.

Chapter 2

Related Work and the Novelty Gap

ORIUS enters a crowded intellectual landscape, but the relevant literatures are crowded in a revealing way. Each contributes a necessary piece of the final argument: uncertainty quantification explains what can be said about predictive trust, runtime assurance explains why intervention must remain separate from nominal control, safe-control methods explain how constraints can reshape action, fault-detection work explains how degradation is recognized, and governance frameworks explain how claims and evidence must be bounded. The novelty gap appears only when those families are read together. They are rich in local technique yet still leave the release-boundary object of degraded observation split across separate traditions.

2.1 Conformal prediction and adaptive calibration

Conformal prediction and related distribution-free uncertainty methods provide one of the few disciplined ways to wrap learned models with coverage-aware uncertainty without assuming a perfectly specified error law [1–3, 13–15]. Conformalized quantile regression and adaptive conformal variants extend that logic to heteroskedastic and nonstationary settings, which makes them especially relevant once a physical system moves through regime changes or degraded telemetry states [2, 3, 16–20].

What remains missing is the release-boundary question that follows from calibrated uncertainty. Coverage guarantees do not by themselves decide how the admissible action set should change, which fallback is legitimate, or what evidence must accompany release once the observation surface becomes unreliable. ORIUS therefore treats calibration as an input to runtime safety rather than as the full safety object.

2.2 Runtime assurance and Simplex-style supervision

Runtime assurance architectures, supervisory controllers, shielding, and runtime verification supply the language of intervention, veto, fallback, and monitored execution [4, 5, 21–27]. That line of work is indispensable because it treats safe release as a separate problem from nominal policy design.

Yet most supervisory schemes enter after the nominal controller has already proposed an action under a state estimate that is treated as sufficiently meaningful. Crucially, classic Simplex architectures hard-code the boundary between unverified and verified controllers based on static state-space triggers, implicitly assuming the physical state observation itself remains perfectly reliable. ORIUS addresses an earlier, more insidious fault line: once observation quality degrades invisibly (the OASG), state-triggered Simplex fails because it trusts a corrupted monitor. The DC_SS assurance layer must reason jointly about epistemic reliability, dynamically widened state uncertainty, repaired action, fallback, and certificate semantics. A supervisor without that observation-side, dynamically calibrated logic is still missing the central OASG object.

2.3 Safety filters, shields, and robust safe control

Control barrier functions, predictive safety filters, robust MPC, and tube-based control translate model uncertainty into admissible control sets and corrective intervention rules [6–8, 28–35]. These methods are the closest classical relatives of ORIUS because they formalize the difference between nominal control intent and safe control execution.

The remaining gap is that this family usually assumes the uncertainty set and safety object have already been defined in a domain-specific, often static or highly conservative way (e.g., assuming a fixed worst-case bounded disturbance $w \in W$). ORIUS does not replace that work with a universal controller. Instead, by architecturally composing Conformal Quantile Regression (CQR) directly with Robust Optimization at runtime, DC_SS dynamically inflates and deflates the robust feasible set based on active telemetry drift. It eliminates the crippling conservativeness of standalone robust control while providing coverage-backed safety guarantees that adapt as sensors fail. The novelty is not the existence of safe control or CQR alone; it is the runtime composition that dynamically bridges data-driven epistemic uncertainty directly into physical constraint-tightening, making safe control portable across domains.

2.4 Anomaly, drift, and degraded-observation detection

Fault diagnosis, change detection, cyber-physical security, and anomaly monitoring provide the vocabulary for stale, delayed, reordered, dropped, or corrupted observations [36–46]. This literature is essential because degraded observation must first be recognized before it can influence action release.

Detection alone does not specify how action semantics should change after trust degrades. A drift alarm can say that something is wrong without saying how uncertainty inflates, how the safe set tightens, what repair is permitted, or what certificate makes the next step auditable. ORIUS uses detection as the opening move in a runtime safety sequence, not as the full response.

2.5 Domain-specific cyber-physical safety systems

Energy systems, autonomous mobility stacks, healthcare monitoring, and other cyber-physical systems already contain rich safety practice in their own local forms [5, 6, 36, 38–40]. This practice matters because any claim of universality that ignores domain structure is empty from the start.

What they usually do not expose is a common contract through which degraded-observation replay, repaired action, fallback semantics, certificate validity, and latency can be compared across domains. ORIUS addresses that missing cross-domain runtime object rather than trying to erase domain structure.

2.6 Safety-case governance and lifecycle discipline

The dissertation also inherits an important lesson from safety-case and lifecycle practice: strong safety claims require structured argument, explicit assumptions, traceable evidence, and disciplined lifecycle boundaries. ORIUS does not claim compliance with sector-specific standards, but its claim-lock pipeline, artifact provenance, and certificate-governance surfaces are best understood in conversation with claim-based safety arguments and lifecycle-oriented risk management [9–12].

Governance frameworks say how evidence should be organized and how claims should be bounded; they do not themselves provide the runtime object that converts degraded observation into repaired or denied release. ORIUS borrows the discipline of explicit evidence

structure and lifecycle boundaries, but its technical novelty remains the runtime layer that sits between observation degradation and actuation.

2.7 Benchmark, replay, and reproducibility gaps

There is also a measurement gap that cuts across the technical literatures above. Many papers report prediction error, control regret, collision rate, or intervention count, but do not record the dual-trajectory evidence needed to tell whether an action was only observed-safe or was truly safe on the latent physical state. That omission matters because the OASG is invisible unless the benchmark keeps both trajectories in view at once.

Without replayable fault schedules, governed intervention traces, and a locked mapping from headline claims to artifacts, a physical-AI safety result is difficult to audit and easy to overstate. ORIUS therefore treats measurement discipline as part of the contribution rather than as packaging around the contribution. The benchmark contract, claim matrix, and manifest surfaces are necessary because they keep the monograph grounded in true-state safety rather than nominal software coherence.

2.8 Why ORIUS is not another forecasting or perception paper

The literature also invites a narrower but common misunderstanding: that ORIUS should be read as one more attempt to beat a baseline predictor, one more perception stack, or one more domain-specific robust controller. The dissertation rejects that framing directly. In the strongest witness row, forecasting quality matters because it affects calibrated uncertainty and therefore the release boundary, but the main scientific object is still the runtime layer that decides what action remains admissible once the observation surface becomes unreliable.

Why this matters for the novelty claim. This distinction prevents ORIUS from collapsing into a model-comparison paper. A new forecaster or perception backbone may improve nominal error while leaving the release-boundary question unchanged. ORIUS is novel only if it remains centered on degraded observation, admissible repair, governed fallback, and certificate semantics. That is why the related-work chapter is organized by missing runtime objects rather than by model families alone.

2.9 The novelty gap ORIUS addresses

The literature becomes most powerful when these families are read together, but it also reveals a persistent separation of levels. Uncertainty methods stop at calibrated sets. Assurance methods stop at supervisory intervention. Safe-control methods stop at plant-specific repair geometry. Detection methods stop at identifying degraded telemetry. Domain systems stop at application-local safety practice. What is missing between them is a single typed runtime composition that starts with observation reliability and ends with repaired action, fallback, certificate semantics, and cross-domain replay under one common contract.

The novelty sentence introduced in Chapter 1 is therefore read narrowly. For traceability, the canonical framing remains: *ORIUS provides a reliability-aware runtime safety layer for physical AI under degraded observation, enforcing certificate-backed action release through uncertainty coverage, repair, and fallback.* This is not a claim of a new conformal method, a new robust-optimization primitive, a new universal controller, or a new conditional-coverage theorem. The promoted ML surface is grouped calibration by reliability bucket plus bounded runtime-safety deltas on the active three-domain lane.

That missing composition is ORIUS’s novelty claim. The dissertation does not argue that any one ORIUS component is unprecedented in isolation. It argues that the combination of observation reliability, action repair, certificate semantics, and universal replay has not previously been stabilized as one reusable runtime safety layer for physical AI. ORIUS is therefore not “better forecasting,” “better perception,” or “one more controller.” It is a release-boundary discipline for degraded-observation safety.

The closest plausible rival to ORIUS is not any single paper in these families. It is an imagined composition of them: adaptive conformal intervals for uncertainty, a safety filter or Simplex supervisor for intervention, anomaly detection for degraded telemetry, and runtime verification for logging and policy enforcement. Even that composite stack still stops short of ORIUS’s central object. It lacks a typed release contract that binds observed-state and true-state semantics, repaired action, fallback, certificate lifecycle, and cross-domain replay under one explicit interface. That final integration step is the novelty claim the dissertation defends.

Boundary of this chapter. This chapter establishes a literature gap, not a theorem or a parity result. Its role is to show why uncertainty, supervision, safe control, detection, and governance still leave the release-boundary object fragmented. The next chapters must still prove and validate that ORIUS closes enough of that fragmentation to justify a universal runtime safety-layer claim under bounded evidence.

Chapter 3

Formal Problem Formulation: OASG under Degraded Observation

3.1 State, observation, action, and constraints

Let x_t^* denote the true physical state of the plant at time t , let $\hat{x}_t$ denote the observation available to the nominal policy, and let u_t denote the candidate action proposed for release. Let $\mathcal{C}(x)$ denote the set of admissible actions when the state is x , or equivalently the safety predicate that separates legal from illegal actuation. ORIUS is concerned with regimes in which the software stack continues to evaluate legality on $\hat{x}_t$ even though the physically relevant state is x_t^* or a broader observation-consistent set around it.

This distinction matters because physical safety is defined on the plant, not on the software representation alone. If $\hat{x}_t$ lags, omits, or distorts the relevant state variables, then the nominal controller can continue producing actions that are legal with respect to the observed state while being unsafe with respect to the state that is actually realized. The formal problem of the dissertation is therefore not generic uncertainty. It is the mismatch between observation-conditioned legality and physical legality.

3.2 Observation degradation operator and fault process

Observation degradation is modeled as a runtime operator

$$\hat{x}_t = \mathcal{D}(x_t^*; \delta_t),$$

where δ_t is a fault process that summarizes the degradation mechanisms affecting the telemetry channel. In the repository and benchmark, this process includes stale packets, missing samples, delay, packet reordering, spikes, drift, transport corruption, and domain-specific state-estimation loss. The universal object of the dissertation is therefore not a plant alone; it is the coupled pair $(x_t^*, \hat{x}_t)$ together with the degradation process that can make those two states diverge.

Two implications follow. First, the fault process is not merely a nuisance variable to be averaged away in nominal performance statistics. It is part of the safety problem definition. Second, runtime safety cannot be reduced to point estimation of the true state. Under degraded observation, the runtime must reason over a set of plausible physical states rather than rely on one best guess.

3.3 Fault taxonomy

The fault process is treated generically, but the benchmark and adapters instantiate it through a shared taxonomy:

1. **Freshness faults:** stale or frozen telemetry that preserves coherence while erasing temporal validity.
2. **Availability faults:** missing packets, dropped samples, and partial channel loss.
3. **Timing faults:** delay, jitter, or reordering that breaks synchronization with the physical plant.
4. **Value faults:** spikes, outliers, drift, bias, or corruption that distort the observed state itself.
5. **Inference faults:** estimator or learned-model errors that turn degraded inputs into overconfident downstream summaries.

This taxonomy is intentionally broad enough to travel across all ORIUS domains. A battery controller and a healthcare monitoring runtime do not share the same raw sensors, but they do share the possibility that freshness, timing, value, or inference degradation breaks the semantics of observed legality.

3.4 Definition of the observation–action safety gap

The central event of the dissertation is the observation–action safety gap:

$$u_t \in \mathcal{C}(\hat{x}_t) \quad \text{but} \quad u_t \notin \mathcal{C}(x_t^*).$$

Equivalently, the action appears safe on the observed state but unsafe on the true physical state. This event is what the benchmark later measures through true-state safety violation and observed-state satisfaction. It is also what makes degraded observation a first-class runtime problem rather than an upstream modeling nuisance.

The OASG formulation is stronger than a generic claim that “uncertainty matters.” Uncertainty can be large without creating a safety violation if the action remains conservative enough. The OASG is the specific regime where the observation channel misleads the release decision itself. That is why ORIUS focuses on action mediation, not only on uncertainty reporting.

3.5 Observation-consistent state sets

ORIUS does not assume that the true state can be recovered exactly at runtime. Instead, it works with an observation-consistent state set

$$\mathcal{X}_t = \{x : x \text{ remains plausible given } \hat{x}_t \text{ and runtime reliability } w_t\},$$

where $w_t \in [0, 1]$ is the observation reliability score. As reliability degrades, $\mathcal{X}_t$ widens. This widening is the formal bridge between degraded telemetry and safe action. Once $\mathcal{X}_t$ expands, the admissible action set must contract:

$$\mathcal{U}_t^{\text{safe}} = \{u : u \in \mathcal{C}(x) \text{ for every } x \in \mathcal{X}_t\}.$$

The runtime question is then precise: given a candidate action u_t , does it belong to the current safe set $\mathcal{U}_t^{\text{safe}}$? If not, can it be repaired by projection or replacement? If not, which governed fallback policy must be triggered? This formulation is what lets ORIUS connect reliability scoring, uncertainty inflation, safe-set tightening, and certificate semantics into one runtime loop.

3.6 Repair, fallback, and release

Let Π_t denote the runtime repair map. Given a candidate action u_t , ORIUS computes either a repaired action

$$\tilde{u}_t = \Pi_t(u_t, \mathcal{X}_t)$$

that lies inside $\mathcal{U}_t^{\text{safe}}$, or else escalates to a fallback action u_t^{fb} specified by the domain governance policy. The formal point is that the runtime layer does not merely observe illegality; it mediates release. This distinguishes ORIUS from alarm-only monitoring and from nominal predictive systems that surface confidence but leave the release decision unchanged.

The release object is completed by a certificate

$$\kappa_t = (\hat{x}_t, w_t, \mathcal{X}_t, \tilde{u}_t, h_t, g_t),$$

where h_t denotes a validity horizon or half-life proxy and g_t denotes governance metadata such as lifecycle state, fallback reason, or audit hash context. Later chapters make that certificate operational through CertOS.

3.7 Dual-trajectory evaluation

The dissertation evaluates ORIUS on dual trajectories. The *observed trajectory* is the one visible to the nominal controller and therefore to the candidate action it proposes. The *true trajectory* is the one used to determine whether that candidate or repaired action actually violates the physical safety envelope. A benchmark that inspects only observed-state legality cannot see the OASG directly. ORIUS-Bench therefore records the candidate action, the repaired action, true-state violation, observed-state satisfaction, optional true and observed margins, fallback and intervention markers, latency, and certificate validity inside one replay schema.

This dual-trajectory formulation is the critical measurement move of the dissertation. It keeps the benchmark aligned with the formal problem object instead of allowing nominal performance metrics to impersonate safety evidence.

3.8 DomainAdapter grounding

The universal formulation is grounded in the repository through the `DomainAdapter` contract. Each domain must supply:

1. telemetry parsing and state semantics,

2. uncertainty or margin construction hooks,
3. admissible-action or repair logic,
4. true-state violation predicates,
5. observed-state satisfaction predicates, and
6. certificate payload specifics.

This boundary is what makes the universal claim scientifically credible. The dissertation does not claim that one plant model or one controller solves every domain. It claims that every domain can expose the same runtime safety objects and thereby enter the same benchmark and governance discipline.

3.9 Scope and assumptions

The formal problem is intentionally narrower than a claim of general AI safety. ORIUS assumes:

- an actionable physical system with explicit or computable safety constraints,
- an observation channel that can degrade in ways that matter to release semantics,
- a runtime interposition layer capable of intercepting, repairing, or replacing actions,
- and a governance layer capable of recording why a release remained admissible.

ORIUS does not assume perfect state recovery, universal optimal control, or equal empirical maturity across all domains. It also does not assume that uncertainty calibration alone proves safety. The later theorem and validation chapters are therefore written under explicit assumptions and are tied to specific artifact surfaces rather than left as narrative implication.

Boundary of this chapter. This chapter defines the formal objects that the rest of the dissertation uses: degraded observation, observation-consistent state sets, repaired release, and dual-trajectory evaluation. It does not yet prove the theorem ladder or claim equal cross-domain closure; those claims remain bounded by the assumptions and evidence surfaces introduced in the chapters that follow.

Part II

ORIOUS Framework and Theorem Ladder

Chapter 4

ORIOUS Architecture: DC3S Runtime Layer

ORIOUS enters the dissertation at the narrowest and most consequential point in the control loop: after nominal intelligence has proposed an action, but before that action is allowed to alter a physical plant. That placement is deliberate. Batteries, vehicles, healthcare monitors, and future additional domains do not share one estimator, one optimizer, or one plant model. What they can share is a disciplined runtime treatment of what happens when degraded observation changes the meaning of safe release.

4.1 Why a runtime layer is the right object

The architectural claim is therefore narrower and stronger than a proposal for a universal controller. Degraded observation creates a release-boundary problem that sits between inference and actuation. Once telemetry becomes stale, delayed, missing, or corrupted, the relevant question is no longer only “what would the nominal controller do?” It becomes “what action can still be justified for physical release under the observation quality now available?” ORIOUS is the layer that answers that second question.

That framing is what lets the book stay universal without drifting into abstraction. The runtime layer leaves domain-specific controllers, estimators, and plant models in place while freezing one common control-flow grammar for degraded observation. That grammar is Detect, Calibrate, Constrain, Shield, and Certify.

4.2 The DC3S release step

The universal runtime step has five stages.

Detect compute observation reliability from freshness, missingness, delay, ordering, spikes, and consistency markers;

Calibrate widen the observation-consistent state set as reliability falls;

Constrain recompute the admissible action set under the widened uncertainty set;

Shield repair the candidate action or replace it with a bounded fallback when no admissible repaired action remains;

Certify emit a causal certificate describing the reliability state, uncertainty summary, repaired action, fallback status, and lifecycle metadata that govern release.

What matters is not that each stage is unprecedented in isolation. The architectural move is the typed composition of all five at one runtime boundary. ORIUS treats degraded observation as something that must change release semantics immediately, not as a warning signal that can be logged while nominal actuation proceeds unchanged.

4.3 Universal runtime flow

Figure 4.1 makes the runtime object concrete. Reliability is not a dashboard number layered on top of an unchanged controller. It changes the admissible state set, which changes the admissible action set, which in turn changes whether nominal action may be released, whether it must be repaired, or whether the runtime must fall back to a safer policy. Certification closes the loop by recording what was released, why it remained admissible, and when that justification should expire.

4.4 DomainAdapter as the universal boundary

The `DomainAdapter` contract is the narrow interface that makes universality credible. The adapter owns the local semantics that should remain local:

- how raw telemetry is interpreted into state features;
- how true-state violation and observed-state satisfaction are defined;
- how candidate action is repaired or projected;




Figure 4.1: ORIOUS runtime flow from degraded telemetry to repaired action and governed certificate. The figure is generated from the same publication assets that feed the manuscript-facing parity and governance surfaces.

- what fallback means for the plant; and
- what certificate payload must be preserved for reviewable release.

The universal kernel owns what should remain shared:

- reliability handling;
- uncertainty inflation;
- admissible-set tightening;
- benchmark-schema compatibility; and
- certificate lifecycle discipline.

This separation is the implementation answer to the universality question. ORIOUS does not argue that every domain shares one plant geometry. It argues that each domain can expose its own geometry through one runtime contract while the safety logic surrounding degraded observation remains reusable.

Table 4.1: Domain state schema for ORIOUS universal framework.

Domain	State Fields	Unit	Description
Battery	soc_mwh, load_mw, renewables_mw	MWh, MW	Battery SOC and dispatch context
AV	ego_speed_mps, relative_gap_m, ttc	m/s, m, s	Bounded nuPlan longitudinal replay context
Healthcare	hr_bpm, spo2_pct, respiratory_rate	bpm, %, /min	Retrospective monitoring and alert context

Table 4.2: Safety predicates per domain.

Domain	Safety Predicate	Config Key
Battery	SOC $\in$ [min_soc, max_soc]; charge/discharge mutual exclusion	min_soc_mwh
AV	TTC and predictive-entry barrier remain satisfied on the bounded nuPlan replay state	ttc_min, gap_min
Healthcare	vital alert release remains inside the bounded monitoring envelope	spo2_min_pct, hr_bounds

4.5 Relationship to supervision and safety filters

ORIOUS should be read as compatible with, not opposed to, established safety-filter and runtime-assurance ideas. Control barrier functions, predictive safety filters, and robust MPC show how plant constraints can be translated into admissible control updates [6–8, 31]. Simplex-style runtime assurance shows how a high-performance controller can be supervised by a safer fallback authority when operating conditions invalidate the nominal control contract [4, 5].

ORIOUS differs in emphasis. It freezes degraded observation itself as the opening object of the runtime step. Reliability loss is not merely one more disturbance bound; it is the mechanism that tells the runtime when the meaning of observed-state legality has changed. That is why Detect and Calibrate appear before Shield. The runtime does not only ask how to project action back into a safe set. It asks how the safe set should be redefined once the observation channel can no longer be trusted at face value.

Table 4.3: Universal DC3S pipeline stages.

Stage	Name	Adapter Method	Output
1	Detect	compute_oqe	w_t , flags
2	Calibrate	build_uncertainty_set	$\mathcal{U}_t$, meta
3	Constrain	tighten_action_set	$\mathcal{A}_t$
4	Shield	repair_action	safe_action, repair_meta
5	Certify	emit_certificate	certificate

4.6 Certificate emission and operational semantics

Certify is not bookkeeping. Once repaired action is released, the runtime must still say what assumptions justified release, which fallback would be taken if the certificate expired, and how later review can reconstruct the decision. In ORIOUS, certificate metadata therefore sits inside the same runtime object as reliability, repair, and fallback.

This is what turns ORIOUS from a controller-side heuristic into a governed runtime layer. The release record must be causal, bounded, and auditable. That discipline is what makes it possible to compare domains through runtime artifacts rather than after-the-fact narrative.

4.7 Runtime cost and operational realism

ORIOUS is written for operational settings rather than theorem exposition alone. Runtime overhead, fallback behavior, and auditable certificates are therefore part of the scientific object. A safety layer that cannot act in time, cannot preserve a coherent lifecycle state, or cannot explain why it intervened is not a convincing architecture for physical AI.

That operational constraint is what leads naturally into the next two chapters. Once the runtime layer has been fixed as the central object, the next question is what can actually be proved about it and how those claims can be measured without confusing nominal coherence for physical safety.

4.8 Architectural takeaway

What this chapter establishes is architectural rather than empirical closure. ORIOUS provides a reusable runtime contract that translates degraded observation into reliability-conditioned uncertainty, tightened action legality, repaired or fallback release, and certificate-aware governance. Stronger claims about optimality, plant realism, or deployment maturity remain tied to the benchmark rows and governed evidence surfaces that follow.

Chapter 5

Mathematical Foundations and Safety Guarantees

5.1 Why the theorem surface is intentionally narrow

ORIUS is not defended by one grand theorem claiming safety for every plant, every sensor regime, and every domain. The mathematical surface is intentionally narrower and therefore more durable. The theorem ladder states what can be proved once degraded observation is modeled explicitly, once reliability-conditioned uncertainty is admitted as part of the runtime object, and once repaired action or fallback is required before release.

The monograph keeps that surface deliberately smaller than the dissertation’s full formal register. The broader theorem/lemma/proposition lattice is part of the thesis governance apparatus; the paper-facing surface only carries the claims that remain central after proof hardening and claim-scope reduction. Concretely, the paper-facing surface is organized around four flagship theorem roles: the observation–action safety gap, observation necessity, ORIUS covered release preservation, and typed transfer under domain obligations. The permanent T1–T11 labels are retained as audit identifiers, not as a requirement that every numbered item be promoted to a headline theorem in the paper.

That discipline matters because it keeps the dissertation from making a false choice between abstraction and implementation. The theorem surface is written against typed runtime objects, explicit assumptions, and artifact-backed diagnostics. As a result, the strongest claims are not merely conceptual and they are not merely empirical.

5.2 Notation and assumption discipline

The formal objects used in this chapter are simple but exact:

- a latent physical state x_t ;
- an observed state or telemetry-derived state z_t ;
- a candidate action a_t proposed by nominal intelligence;
- a repaired or fallback action $\tilde{a}_t$ released by ORIUS;
- a runtime reliability score $w_t \in [0, 1]$ summarizing observation trust; and
- a violation predicate $\mathcal{V}(x_t, a_t)$ defined through the domain adapter.

The score w_t is not interpreted as a probability by definition. Whenever a theorem turns w_t into a probabilistic risk or indistinguishability claim, that conversion is written as an extra theorem-local assumption.

The dissertation does not hide its assumptions. Every strong statement depends on the assumption register in Appendix b. The core assumptions are finite observation-consistent state inflation, availability of a feasible repaired action or bounded fallback, causal certificate issuance, dual-trajectory replay fidelity, adapter faithfulness, and evidence-tier discipline. If any of those assumptions is violated, the associated theorem is not silently inherited by the runtime row.

5.3 Definition of OASG as a formal event

Definition 5.1 (Observation–Action Safety Gap). An observation–action safety gap occurs at time t when the candidate action is admissible under the observed state but unsafe under the true physical state:

$$\mathcal{S}_{\text{obs}}(z_t, a_t) = 1 \quad \text{and} \quad \mathcal{V}(x_t, a_t) = 1.$$

This definition is deliberately agnostic to plant family. The adapter chooses the local meaning of admissibility and true-state violation. The universal layer fixes the semantic pattern of the failure event.

5.4 The paper-facing theorem surface

The paper-facing theory chapter carries four flagship theorem roles and a small set of supporting propositions/corollaries. The full T1–T11 ladder remains in Appendix c and in the theorem registry for traceability; the paper does not ask the reader to treat all eleven identifiers as equally central main theorems. The claim is deliberately conditional: ORIUS is safety-optimal under covered observation ambiguity, not globally optimal for every possible physical-AI system.

Theorem 5.2 (OASG existence under degraded observation). *If the observation channel admits a degradation regime in which the observed-state safety predicate is no longer informationally sufficient for the true-state safety predicate, then there exists an admissible observed-state release that produces a true-state violation.*

Theorem 5.3 (Safety preservation under repaired release). *Assume the inflated observation-consistent set is conservative with respect to the latent state and assume the repair or fallback operator returns only actions admissible over that set. Then the ORIUS release step preserves the defended safety predicate whenever the required repair or fallback remains feasible.*

Theorem 5.4 (Degradation-sensitive risk envelope). *If uncertainty inflation is monotone in decreasing reliability and the release step is monotone in admissible-set tightening, then the risk envelope contracts as reliability falls. In other words, intervention pressure is a lawful consequence of degraded observation rather than an arbitrary heuristic.*

Proposition 5.5 (T4: Observation Necessity / No Free Safety). *A controller that ignores observation-quality degradation can remain observed-state legal while incurring nonzero true-state violation under admissible degradation regimes, even when the nominal controller is internally coherent and locally optimal on the observed state.*

The point of this ladder is modest but important. ORIUS does not prove that every plant can avoid intervention. It proves that once reliability affects the observation-consistent state set, safe release must itself become reliability aware.

The final supporting optimality statement is the covered observation-ambiguity sandwich. Let $h : \mathcal{X} \rightarrow \mathcal{O}$ be the observation map. For an observation o , define the ambiguity class and common safe core as

$$B(o) = \{x \in \mathcal{X} : h(x) = o\}, \quad K(o) = \bigcap_{x \in B(o)} C(x).$$

If $K(o) = \emptyset$, no mandatory-release controller that only observes o can guarantee true-state safety at that observation. More generally, any observation-only mandatory-release policy π

obeys the Bayes lower bound

$$\text{TSVR}(\pi) \geq \mathbb{E}_O \left[\min_{a \in \mathcal{A}} \mathbb{P}(a \notin C(X^*) \mid O) \right].$$

Conversely, if ORIUS constructs an uncertainty set $\mathcal{X}_t$ with $\mathbb{P}(X_t^* \notin \mathcal{X}_t) \leq \alpha$ and releases only actions safe for every state in $\mathcal{X}_t$, then

$$\text{TSVR}_{\text{ORIUS}} \leq \alpha,$$

with zero violation under exact coverage. This is the `T10_T11_ObservationAmbiguitySandwich`: the lower bound applies to observation-only release, while the upper bound applies to reliability-aware covered release.

Table 5.1 records the main counterexamples that would make a stronger version of the theorem false. They are included in the main text because they define the exact boundary between the defended ORIUS claim and an overclaim.

Table 5.1: Counterexamples and assumption boundaries for the observation-ambiguity theorem.

Potential reviewer counterexample	Why it matters	How the manuscript bounds the claim
Nonempty common safe core $K(o)$	Ambiguity alone need not force violation	The lower bound uses $\min_a \mathbb{P}(a \notin C(X^*) \mid O)$
Controller may abstain or delay	T4 is about mandatory release, not optional release	The theorem names mandatory-release controllers explicitly
Coverage miss $X_t^* \notin \mathcal{X}_t$	ORIUS cannot guarantee zero violation off coverage	The upper guarantee is $\text{TSVR}_{\text{ORIUS}} \leq \alpha$
Wrong safe set $\hat{C}(x)$ or invalid adapter	A wrapper cannot repair a wrong safety predicate	Safe-set correctness and adapter validity remain proof obligations
Fallback physically infeasible	No epistemic wrapper can create missing actuation authority	Fallback admissibility is required and infeasible cases must fail closed
Poisoned or shifted calibration data	The uncertainty set may be falsely narrow	Coverage and calibration validity are assumptions, not automatic conclusions

Appendix c is the authoritative proof index for the defended theorem program. The paper-facing surface intentionally remains narrower: it inherits the same assumption language and scoped interpretation as the thesis, but it does not reproduce the full 83-item formal register.

Table 5.2: DC3S Theorem Dependency and Proof Completeness. Each formal result is implemented in code and tested by at least one automated test.

Result	Name	Depends On	Proof Method	Code Reference
Defn. 1	OQE reliability weight w_t		Definition	<code>quality.py:compute_reliability()</code>
Prop. 1	Inflation monotonicity	Defn. 1	Analytic	<code>coverage_theorem.py:verify_inflation_geq_one()</code>
Thm. 1	Marginal coverage guarantee	Prop. 1, Defn. 1	Conformal CP theory	<code>coverage_theorem.py:assert_coverage_guarantee()</code>
Thm. 3	Expected violation bound	Thm. 1	Expectation calculus	<code>coverage_theorem.py:compute_expected_violation_bound()</code>
Thm. 9	Group-conditional coverage	Thm. 1	Mondrian CP	<code>coverage_theorem.py:mondrian_group_coverage()</code>
Thm. 10	Hoeffding high-prob. bound	Thm. 3	Hoeffding inequality	<code>coverage_theorem.py:hoeffding_violation_bound()</code>
Thm. 11	Byzantine OQE reliability bound	Defn. 1	Trimmed-mean theory	<code>quality.py:compute_reliability_robust()</code>

All results are unit-tested in `tests/test_conditional_coverage.py` and `tests/test_adversarial_faults.py`.

5.5 Risk envelope interpretation

The risk envelope is the most important bridge between formal and empirical language in the dissertation. ORIUS does not claim safety from a single forecast error number. It reasons through miscoverage, intervention feasibility, and mean reliability. That framing is what allows the theorem layer to connect to uncertainty calibration without collapsing into a purely statistical dissertation.

In the tightened thesis surface, this bridge is explicit: the episode-level bound is read through a per-step residual-risk budget, not as a direct consequence of marginal conformal coverage alone.

The uncertainty references in Chapter 2 matter here. Classic conformal prediction supplies the disciplined interpretation of uncertainty sets under exchangeability, while adaptive conformal methods show how long-run calibration language can be generalized more honestly under nonstationary streams [1–3]. ORIUS does not claim that these methods alone solve runtime safety. It uses them to clarify what the uncertainty set is allowed to mean before repaired release is evaluated. In practice, the online conformal layer is an SPCI-lite quantile tracker in the style of Gibbs and Candès [3] that maintains an exponentially-weighted sliding window of conformity scores, yielding intervals that adapt smoothly to distribution shift without a full recalibration pass.

5.6 Why the proofs stay witness-backed but the objects stay universal

The energy-management reference row remains the deepest theorem-to-code-to-artifact chain because the witness row provides the richest availability of proof-facing diagnostics, replay traces, and locked artifacts. That does not make the theorem objects domain-specific. The objects are written against reliability, uncertainty inflation, admissible action tightening, repair or fallback, and certificate causality. Those are universal runtime objects instantiated through the adapter, not energy-management-exclusive definitions.

5.7 Theorem-to-artifact linkage

Figure 5.1 summarizes the discipline used throughout the manuscript. Every theorem is attached to an assumption family, a typed runtime object, and at least one artifact-facing diagnostic or replay surface. This is why the dissertation can state theorem results without pretending that empirical maturity is equal across all rows.




Figure 5.1: The integrated theorem gate keeps theorem language aligned to typed contracts, assumptions, and locked artifacts rather than to informal prose alone.

5.8 What is proved, validated, and still open

The theory chapter therefore ends with an explicit boundary. Theorem statements live here and in Appendix c. Validated claims only advance when replay, calibration, runtime, and governance surfaces exist for the relevant domain row. Roadmap claims remain outside defended theorem language. That separation is not a stylistic preference. It is the mechanism that lets ORIUS argue for a universal runtime safety layer without pretending that broader future-domain empirical closure has already been earned.

Chapter 6

Benchmark and Reproducibility Protocol

6.1 Why ORIUS-Bench is central to the thesis

Once the runtime object has been fixed, the next question is whether it can be measured in a way that matches the theory rather than only the software stack. ORIUS-Bench is the answer to that question. It is not infrastructure attached after the science; it is the measurement contract that makes the dissertation’s central object observable. If OASG is the gap between observed-state legality and true-state safety, then the benchmark must expose both surfaces explicitly. A benchmark that measures only forecast error, controller cost, or observed-state success cannot validate the ORIUS claim because it cannot see the hidden failure event the book is about.

6.2 Dual-trajectory evaluation

Every ORIUS-Bench episode therefore carries two trajectories:

- the *observed* trajectory available to the runtime after degradation; and
- the *true* trajectory used only by the replay or evaluation harness to judge whether release was physically safe.

That separation makes the benchmark stricter than a conventional controller comparison. A nominal controller may still appear coherent on the observed state while the replay harness records a true-state violation. ORIUS-Bench treats that divergence as first-class evidence.

6.3 Deterministic replay and fault schedules

The benchmark uses deterministic fault schedules so that comparisons are reproducible and so that certificate traces can be replayed under the same degraded-observation regime. Fault families include dropout, stale holds, delay, reordering, spikes, drift, and domain-local analogs. The scientific point is not merely that faults exist. It is that the same fault grammar can be replayed against one release contract across domains.

Strict replay discipline matters for another reason: it prevents the dissertation from hiding behind one-off operational anecdotes. A run must be reproducible through a manifest, fault schedule, and artifact family. If a result cannot be regenerated from the tracked surface, it is not part of the defended thesis claim.

6.4 Metrics that match the runtime object

ORIUS-Bench measures the runtime layer through seven core metrics:

TSVR true-state violation rate, the primary safety outcome;

OASG the observed-safe / true-unsafe discrepancy rate;

CVA certificate-valid action rate or validity surface;

GDQ graceful degradation quality under fallback or reduced-operation modes;

IR intervention rate, capturing how often the runtime must repair or deny release;

AC audit completeness, measuring whether release is reconstructable after the fact;

RL runtime latency, measuring whether the protection layer remains operationally usable.

This metric family matters because it keeps the book from collapsing into a single leaderboard score. ORIUS is not about RMSE alone, and it is not about controller cost alone. It is about safe release under degraded observation. The benchmark therefore measures whether release remained safe, whether intervention was required, whether the certificate remained meaningful, and whether the runtime stayed fast enough to matter.

6.5 Strict universal mode and legacy compatibility

The dissertation adopts a strict universal reading of the benchmark. Domain-neutral violation predicates, observed-state legality predicates, and constraint margins must flow through

Table 6.1: Cross-Domain OASG (Observed-Action Safety Gap) Rates.

Domain	Baseline OASG	ORIOUS OASG	Reduction
Battery (Ref.)	0.0083	0.0000	100.0%
Healthcare	0.1945	0.0000	100.0%
Vehicle (AV)	0.289250	0.000163	99.94%

Table 6.2: Claim-governing three-domain runtime evidence. These rows are the headline safety rows allowed by the claim validators; validation-harness rows are diagnostic only.

Domain	Evidence tier	Baseline TSVR	ORIOUS TSVR	Intervention	Fallback	Certificate-valid release	Runtime witness	Claim boundary
Battery Energy Storage	locked witness runtime	0.008300	0.000000	0.020833	0.020833	0.993056	1.000000	witness row; no field deployment claim
Autonomous Vehicles	all-zip grouped mPlan runtime	0.289250	0.000163	0.500301	0.173716	0.999924	0.999837	bounded replay/surrogate; no road closure claim
Medical and Healthcare Monitoring	MIMIC runtime denominator	0.194489	0.000000	0.910099	0.479907	1.000000	1.000000	retrospective monitoring; no live clinical deployment claim

Source: `reports/publication/three_domain_ml_benchmark.csv`.

the adapter contract. Legacy battery-shaped compatibility fields may remain in the codebase for historical or bounded-support reasons, but they are not allowed to define the universal benchmark story in the manuscript.

That distinction keeps the witness row from quietly redefining the universal object. Battery is the deepest defended row because its evidence chain is strongest, not because the benchmark is secretly battery-specific. ORIUS-Bench must therefore be written in domain-agnostic terms even when the strongest current replay surface is battery-backed.

6.6 Artifact provenance and claim locking

The benchmark protocol is inseparable from the claim-lock pipeline. The canonical submission surface is `orius_book.tex`, while the mirrored internal controller remains `paper/paper.tex`; their headline tables and figures are backed by manifests, publication matrices, and governed artifact paths. This is what prevents the manuscript from becoming a persuasive story about results that cannot later be checked.

The repository’s publication surfaces already enforce this discipline. The parity matrix, submission scorecard, runtime/governance matrices, and deployment-validation scope are not mere reporting layers. They are the machine-readable surfaces that keep prose aligned with evidence status.

6.7 Reproducibility checklist

For a result to count as defended in this dissertation, the following conditions must hold:

1. the domain row must have an explicit safety predicate and adapter semantics;

2. the replay must expose true and observed trajectories;
3. the fault schedule and manifests must be reproducible;
4. the result must land in a tracked artifact family; and
5. the claim validator must allow the corresponding manuscript language.

This checklist is stricter than “the code ran once.” It is the mechanism by which ORIUS turns software artifacts into scientific evidence.

6.8 Benchmark takeaway

What this chapter adds to the argument is a measurement discipline equal to the runtime object introduced in Chapters 3 and 4. ORIUS-Bench makes degraded-observation safety measurable through dual-trajectory replay, deterministic fault injection, runtime metrics matched to the release object, and claim-locked provenance. Equal-domain maturity still depends on which rows clear those requirements, but the measurement contract itself is now stable across the dissertation.

Part III

Witness Domain Evidence

Chapter 7

Witness Domain Implementation

7.1 Why battery is the witness domain

Battery remains the witness domain because it is the clearest place in the repository where the full ORIUS stack can be seen at once: degraded telemetry, reliability estimation, uncertainty inflation, repaired dispatch, fallback behavior, certificate traces, replayable fault schedules, and claim-locked artifact surfaces. That makes it the right row to demonstrate end-to-end doctrine closure without allowing the dissertation to collapse into an energy-only story.

The witness role is methodological rather than rhetorical. Battery is not the identity of ORIUS. It is the row in which the theorem layer, the implementation layer, and the artifact layer are all strong enough to support the book’s deepest defended argument.

7.2 Witness-domain problem and plant semantics

The battery plant exposes a particularly clear form of OASG. Dispatch legality is computed against observed state of charge, availability, and feeder conditions, but the physical asset evolves on the true state. When telemetry becomes stale or corrupted, the optimizer can continue to return actions that appear feasible on the observed state while the physical asset is already drifting toward an unsafe or infeasible region.

That clarity is useful for a witness row because it makes visible what ORIUS actually changes: the release boundary, not merely the forecast. The runtime widens the observation-consistent state, recomputes admissible dispatch, repairs or curtails candidate action, and emits a certificate describing why release remains justified.

7.3 Data surfaces and dataset cards

The witness row is artifact-first rather than repository-inline. The battery surface is therefore described through tracked dataset cards, manifests, and publication artifacts instead of through the fiction that all raw data live inside the git tree. That design is not a weakness. It is how a dissertation remains reproducible when its empirical surface depends on large external or governed datasets.

Table 7.1: Dataset-card inventory for the four evaluation regions.

Dataset	Country	Rows	Features	Period	Coverage (%)	Source	Carbon Source
Germany (OPSD)	DE	17,377	94	2018-10-07 – 2020-09-30	100.0	OPSD + SMARD	SMARD generation mix
USA (EIA-930 ERCOT)	US	13,453	62	2024-07-01 – 2026-01-20	100.0	EIA-930 ERCOT	proxy_generation_mix
USA (EIA-930 MISO)	US	13,638	114	2024-07-01 – 2026-01-20	100.0	EIA-930 MISO	proxy_generation_mix
USA (EIA-930 PJM)	US	13,639	62	2024-07-01 – 2026-01-20	100.0	EIA-930 PJM	proxy_generation_mix

The dataset-card layer does more than document provenance. It tells the reader which geographies, telemetry resolutions, market structures, and operating contexts actually support the witness claim. That keeps the battery row from drifting into a false claim of energy-system universality.

7.4 Forecasting, calibration, and reliability surfaces

The witness row deliberately separates three roles that are often blurred in energy papers:

1. the forecasting surface that supports nominal dispatch;
2. the uncertainty or calibration surface that informs admissible release; and
3. the reliability surface that decides how aggressively the runtime must tighten action when telemetry degrades.

That separation matters because ORIUS is not strongest where it finds the best forecaster. It is strongest where reliability, uncertainty, repair, and certificate semantics can be shown working together on one defended row.

Taken together, the tables and figure support a bounded but important point. Even when deep forecasting does not dominate the engineered tabular baseline, learned reliability remains a scientifically relevant extension because ORIUS cares about release under degraded sensing, not only about point prediction quality.

Table 7.2: Six-model forecast comparison (DE & US). GBM dominates across all targets. PICP/Width shown only for GBM (production conformal intervals). Negative R^2 indicates worse-than-mean predictions.

Region	Target	Model	RMSE	MAE	sMAPE (%)	R^2	PICP@90	Width (MW)
DE	Load	GBM	305.1	187.9	0.39	0.9988	88.2	1070.2
DE	Load	LSTM	1204.5	945.2	5.29	0.8408	not appl.	not appl.
DE	Load	TCN	902.0	713.6	4.25	0.9021	not appl.	not appl.
DE	Load	N-BEATS	715.8	558.9	3.61	0.9317	not appl.	not appl.
DE	Load	TFT	552.9	430.7	2.45	0.9433	not appl.	not appl.
DE	Load	PatchTST	491.4	387.6	2.26	0.9530	not appl.	not appl.
DE	Wind	GBM	219.1	129.4	2.41	0.9991	93.7	833.9
DE	Wind	LSTM	850.5	662.0	83.69	0.8466	not appl.	not appl.
DE	Wind	TCN	737.7	585.0	103.77	0.8802	not appl.	not appl.
DE	Wind	N-BEATS	492.8	385.0	69.37	0.9368	not appl.	not appl.
DE	Wind	TFT	478.3	374.7	57.00	0.9450	not appl.	not appl.
DE	Wind	PatchTST	359.1	282.1	48.03	0.9627	not appl.	not appl.
DE	Solar	GBM	240.3	117.3	69.24	0.9993	78.8	1037.7
DE	Solar	LSTM	977.0	770.9	235.10	0.8619	not appl.	not appl.
DE	Solar	TCN	836.7	655.4	175.57	0.8914	not appl.	not appl.
DE	Solar	N-BEATS	556.3	437.4	95.70	0.9424	not appl.	not appl.
DE	Solar	TFT	488.1	381.5	94.71	0.9506	not appl.	not appl.
DE	Solar	PatchTST	406.5	320.6	103.65	0.9543	not appl.	not appl.
US	Load	GBM	183.4	140.6	0.18	0.9993	82.6	1986.8
US	Load	LSTM	671.7	529.9	3.42	0.8430	not appl.	not appl.
US	Load	TCN	595.0	463.6	2.60	0.8909	not appl.	not appl.
US	Load	N-BEATS	367.6	292.4	1.68	0.9353	not appl.	not appl.
US	Load	TFT	377.6	295.8	1.93	0.9491	not appl.	not appl.
US	Load	PatchTST	346.9	274.1	1.68	0.9525	not appl.	not appl.
US	Wind	GBM	357.1	210.7	2.65	0.9972	40.6	324.9
US	Wind	LSTM	1497.2	1168.1	180.11	0.8592	not appl.	not appl.
US	Wind	TCN	1171.0	911.8	117.98	0.9053	not appl.	not appl.
US	Wind	N-BEATS	719.4	563.0	80.58	0.9306	not appl.	not appl.
US	Wind	TFT	681.0	533.6	81.35	0.9432	not appl.	not appl.
US	Wind	PatchTST	696.5	548.8	98.79	0.9547	not appl.	not appl.
US	Solar	GBM	230.0	82.1	53.09	0.9947	86.5	687.0
US	Solar	LSTM	868.7	691.0	151.85	0.8674	not appl.	not appl.
US	Solar	TCN	686.6	545.8	113.95	0.8874	not appl.	not appl.
US	Solar	N-BEATS	467.2	368.5	76.08	0.9417	not appl.	not appl.
US	Solar	TFT	450.7	351.4	103.59	0.9433	not appl.	not appl.
US	Solar	PatchTST	414.9	323.1	88.50	0.9514	not appl.	not appl.

Table 7.3: Battery DeepOQE diagnostics on the held-out degraded-telemetry episodes.

Metric	Heuristic	DeepOQE
Reliability MAE	0.1176	0.1613
Fault AUC	—	0.5213
Low-reliability fault lift	—	0.000




Figure 7.1: Witness-domain reliability baselines. The engineered forecasting surface and the learned reliability surface are intentionally reported together so the dissertation does not confuse forecasting novelty with degraded-observation novelty.

7.5 Control baselines and adapter implementation

Battery is also the cleanest row for showing the adapter boundary in operational terms. The adapter binds:

- telemetry parsing and state semantics;
- the true-state safety predicate;
- dispatch feasibility and repair geometry;
- certificate payload fields; and
- the bounded fallback policy used when admissible repaired action disappears.

This is where the implementation philosophy becomes concrete. The nominal forecasting and optimization stack can vary, but the ORIUS adapter contract and replay schema remain fixed.

7.6 Locked artifact protocol

The witness row is defended through locked artifacts rather than narrative memory. Training surfaces, replay outputs, theorem traces, calibration diagnostics, deep-reliability summaries, runtime tables, and parity surfaces all point back to generated publication artifacts or manifests. This is where the chapter-first dissertation model and the repository’s claim-lock machinery meet: the prose is curated, but the evidence remains machine-governed.

7.7 Witness-row scope

Battery carries the book’s deepest defended claim because it closes the full runtime-safety chain from degraded observation to governed release. It does not prove that every battery chemistry, market policy, or asset-health objective is universally solved. Instead, it gives the dissertation a calibration surface: one row where theorem, runtime, and artifact closure can all be inspected together before the argument is lifted across domains.

Chapter 8

Witness Results and Failure Analysis

8.1 The governing empirical question

The witness row answers one question for the book as a whole: when telemetry degrades, can a runtime layer reduce true-state safety loss without pretending that nominal controller legality is sufficient? Battery is the deepest defended empirical surface because it closes the full ORIUS loop from degraded observation to repaired action to governed release.

That framing matters for how the results should be read. The point of the witness chapter is not to crown one battery optimizer. It is to show, on the row where the evidence chain is richest, that the release boundary behaves differently once degraded observation is treated as a first-class safety object.

8.2 Experiment design and baseline logic

The witness experiments are not designed to show that ORIUS is the best optimizer under all cost metrics. They are designed to test a narrower and more important claim: whether a certificate-aware runtime can keep hidden true-state violations from passing through a nominally coherent observation surface. The baseline set is therefore intentionally mixed. It includes stronger nominal optimizers, interval-aware baselines, and controller variants that already possess some robustness logic.

This is what gives the chapter its value. The scientific question is not simply who optimizes best in nominal conditions. It is what remains safe once the observation channel itself begins to fail.

8.3 Headline result

The main table should be read as the empirical hinge of the dissertation, because it tests the point where the formal argument meets defended runtime evidence.

Table 8.1: Main controller comparison under aggregate telemetry faults. DC³S variants achieve zero true-SoC violations across all seeds and scenarios. 95% CI via percentile bootstrap over seeds $\times$ scenarios.

Controller	Viol. Rate	Severity P95 (MWh)	Interv. Rate	Cost $\Delta\%$	n
DC³S (wrapped)	0.000	0.000	0.000 [0.000, 0.002]	+0.05	30
DC³S (FTIT)	0.000	0.000	0.019 [0.000, 0.037]	+0.27	30
Deterministic LP	0.008 [0.000, 0.018]	0.147 [0.000, 0.359]	0.000	+0.00	30
CVaR Interval	0.013 [0.000, 0.018]	7.186 [0.000, 57.991]	0.000	-0.01	30
Robust Fixed	0.013 [0.000, 0.018]	7.186 [0.000, 57.991]	0.000	-0.01	30

Several plant-specific robust baselines already avoid true-state violations on this witness task, so the contribution is not that ORIUS is the only safe controller in the comparison. The stronger point is architectural. ORIUS closes the hidden violation mode while keeping intervention, repair, fallback, and certificate semantics explicit at runtime. Deterministic and interval baselines still exhibit material true-state violation once telemetry faults are introduced, even though they remain internally coherent optimization procedures on the observed state.

That distinction is the witness-domain result. ORIUS is not replacing every robust controller with one universal optimizer. It is demonstrating that a runtime layer can sit between nominal control and physical release, detect when observation quality has changed the meaning of legality, and then repair or deny release under explicit certificate semantics.

8.4 Mechanistic interpretation

The witness row supports a mechanistic reading rather than a leaderboard reading. ORIUS changes the release decision in three linked ways:

1. reliability loss is detected before action release;
2. uncertainty is widened before legality is checked; and
3. unsafe actions are either repaired or replaced before they reach the plant.

The headline table is therefore evidence that the release boundary changed, not merely that one model forecasted more accurately than another.

8.5 Ablations and where the gain comes from

The ablation table shows where the runtime stack earns its gains and where those gains still depend on the structure of the degraded-observation regime.

Table 8.2: Ablation study: DC³S violation-rate reduction vs. baselines under telemetry faults. Two-gate verification: relative reduction $\geq 10\%$ and Wilcoxon signed-rank $p < 0.01$.

Scope	Fault	n	Base Rate	DC ³ S Rate	Rel. Red. (%)	W stat	p -value	Pass
primary	all	160	0.250	0.231	+7.7	0.0	<0.001	×
secondary	delay_jitter	40	0.253	0.230	+9.3	0.0	<0.001	×
secondary	dropout	40	0.250	0.226	+9.7	0.0	<0.001	×
secondary	out_of_order	40	0.252	0.230	+8.9	0.0	<0.001	×
secondary	spikes	40	0.250	0.231	+7.5	0.0	<0.001	×
sec. scenario	all (vs. deterministic_lp)	10	0.350	0.300	+14.3	0.0	0.002	✓
sec. scenario	all (vs. robust_fixed_interval)	10	0.250	0.300	-20.1	0.0	0.002	×
sec. scenario	all (vs. deterministic_lp)	10	0.350	0.300	+14.3	0.0	0.002	✓
sec. scenario	all (vs. robust_fixed_interval)	10	0.252	0.300	-18.9	0.0	0.002	×

This table prevents a simplistic reading of the witness result. ORIUS is not buying safety through one fixed extra margin. It behaves like a genuine degraded-observation runtime layer whose effect depends on the coupling between reliability estimation, uncertainty widening, and local repair geometry. Some scenarios clear the strict gate more comfortably than others, which is exactly what should happen when the runtime is responding to different fault structures rather than papering over them.

8.6 Grouped calibration and reliability slices

ORIUS depends on reliability-conditioned uncertainty rather than on one globally fixed interval. The grouped coverage table below shows why that distinction matters.

Grouped calibration quality moves materially across operating regimes and reliability buckets. That is precisely the pattern ORIUS is designed to take seriously. If interval behavior were uniform, a single robust margin would come closer to sufficing. Because interval behavior changes with regime and observation trust, the runtime must tighten admissible action as confidence in the observation surface falls.

8.7 Operational traces and certificate-aware behavior

The summary tables matter only because the operational traces explain how the gains are earned. Intervention rate, fallback usage, and certificate continuity show whether ORIUS

Table 8.3: Regime-aware CQR coverage and interval width by target and volatility group.

Target	Regime	PICP@90 (%)	Width (MW)	n
Load	LOW	85.2	1244.8	840
Load	MID	88.3	937.4	840
Load	HIGH	91.1	1028.2	840
Wind	LOW	97.6	719.7	840
Wind	MID	93.5	812.0	840
Wind	HIGH	90.0	969.9	840
Solar	LOW	80.2	891.6	840
Solar	MID	82.2	1016.6	845
Solar	HIGH	73.8	1206.2	835

is improving safety through disciplined runtime mediation or merely through indiscriminate shutdown. That distinction is essential. A controller can always buy apparent safety by refusing to act. The witness row matters because ORIUS reduces hidden safety loss while keeping intervention bounded, reviewable, and operationally usable.

8.8 Bounded deep-learning novelty inside the witness row

The witness row also contains a bounded learned-reliability novelty surface. It matters here only insofar as it improves the degraded-observation safety argument rather than turning the dissertation into a forecasting paper.




Figure 8.1: Battery learned-reliability safety surface. The deep model is judged by true-state safety, intervention burden, and certificate-valid release behavior rather than by forecast error alone.

The figure supports a narrower point than the main safety table. It shows that the learned reliability path can improve the runtime’s discrimination of when to intervene and when to allow nominal action to pass. That is a bounded novelty claim about the observation side of ORIUS, not a replacement for the main witness claim about repaired release under degraded telemetry.

8.9 What fails without the runtime layer

Without ORIUS, the telemetry fault is treated as a nuisance external to the release decision. The deterministic and interval baselines show what follows from that assumption: the controller remains computationally coherent while true-state violation reappears. Without calibrated widening, legality is checked against the wrong state set. Without repair, unsafe action is still released. Without certificate discipline, the runtime cannot later explain why the release was allowed or when fallback should have taken over.

8.10 Witness-row claim boundary

The witness chapter supports three claims and no more. First, ORIUS can materially improve degraded-observation safety on a defended row. Second, the kernel remains meaningful across multiple telemetry fault families rather than on one curated failure. Third, intervention and certificate behavior are part of the scientific result, not merely surrounding instrumentation. It does not claim that the battery row proves broader future-domain closure or settles every question about deployment economics, asset aging, or market-optimal dispatch. What it does show is that once observed-state legality is no longer enough, ORIUS changes the release boundary in a measurable, reviewable, and auditable way.

Part IV

Universalization and Extensions

Chapter 9

Universal ORIUS Across Domains

9.1 What universality means in this dissertation

ORIUS is universal in a specific and intentionally disciplined sense. It does not assume that batteries, vehicles, and healthcare monitoring systems share a plant model, control law, or optimization target. It assumes only that each domain must answer the same release-boundary question once observation becomes unreliable: what action may still be released safely, under what margin, with what fallback, and under what certificate?

The shared object is therefore not the nominal controller. It is the runtime safety logic that mediates degraded observation before actuation. This is why the monograph can make a strong universality claim without flattening domain structure.

9.2 A transfer recipe rather than a portability slogan

The dissertation treats domain transfer as a recipe, not a slogan. A new ORIUS row is only meaningful if it specifies:

1. the local degraded-observation hazard;
2. the true-state violation predicate and observed-state legality predicate;
3. the action-repair or fallback geometry;
4. the certificate payload and lifecycle semantics; and
5. the replay and evidence surface through which that row will be judged.

That transfer as a recipe is what converts structural universality into a scientific program. Without it, “universal” would reduce to analogy. With it, the reader can inspect exactly how a domain is instantiated and where its evidence status currently sits.

9.3 Cross-domain status at a glance

Table 9.1: Canonical Battery + AV + Healthcare closure matrix.

Domain	Tier	Source	Baseline TSVR	OR
Battery Energy Storage	reference	locked_artifact	0.0083	0.00
Autonomous Vehicles	runtime_contract_closed	real_data	0.2893	0.00
Medical and Healthcare Monitoring	runtime_contract_closed	verified	0.1945	0.00

The domain-closure matrix belongs in the main line of the argument because it keeps the universality claim honest. Battery is the witness row. The autonomous-vehicle row and the healthcare monitoring row are defended bounded runtime rows under the current artifact surface.

Table 9.2: Canonical runtime release contract witnesses. Each promoted domain is normalized to the same release evidence schema: `observe>reliability>uncertainty>repair_or_fallback>sign_certificate>append_only_audit`.

Domain	Reliability	Safe action	Repair status	Fallback status	Certificate status	Postcondition
Battery	0.993056	{charge_mw: 0.0}	intervention_rate=0.020833	fallback_rate=0.020833	unsigned_allowed; append_only_supported	passed
AV	0.999924	{acceleration_mps2: 0.0}	intervention_rate=0.500301	fallback_rate=0.173716	unsigned_allowed; append_only_supported	passed
Healthcare	1.000000	{alert_level: max_alert}	intervention_rate=0.910099	fallback_rate=0.479907	unsigned_allowed; append_only_supported	passed

Source: `reports/publication/runtime_release_contract_witnesses.csv`. The compact table records publication witnesses; raw runtime traces and append-only event stores remain regenerated/local artifacts.

The witness table records the common release evidence fields that make the domain rows comparable: reliability, uncertainty, proposed and safe action summaries, repair and fallback status, release decision, certificate hash, signature status, append-only audit status, T11 status, and postcondition status. The table is compact by design; it is the publication witness for the release contract, not a replacement for the raw runtime traces.

9.4 Final three-domain release results

The final manuscript-facing result surfaces are generated from the locked CPU release artifacts rather than from hand-entered prose. They report model quality, claim-governing runtime safety, utility preservation against degenerate fail-safe references, and the freeze-validation gates that bound the current submission package.

Table 9.3: Final CPU release model quality used by the manuscript-facing ORIUS package.

Domain	Target	Model	RMSE	PICP90	Status
Battery	load_mw	gbm	254.470	93.0%	pass
Battery	price_eur_mwh	gbm	4.896	91.5%	pass
Battery	solar_mw	gbm	243.177	91.7%	pass
Battery	wind_mw	gbm	107.445	96.5%	pass
AV	target_ego_speed_mps__1s	gbm	0.029	92.0%	pass
AV	target_relative_gap_m__1s	gbm	14.323	95.8%	pass
Healthcare	hr_bpm	gbm	1.521	90.3%	pass
Healthcare	respiratory_rate	gbm	0.653	89.4%	pass
Healthcare	spo2_pct	gbm	0.280	91.2%	pass

Table 9.4: Final claim-governing runtime safety evidence across the three promoted ORIUS domains.

Domain	Baseline TSVR	ORIUS TSVR	Reduction	Fallback	Cert-valid	p95 ms
Battery Energy Storage	0.008300	0.000000	100.00%	2.1%	99.3%	0.812
Autonomous Vehicles	0.289250	0.000163	99.94%	17.4%	100.0%	0.744
Medical and Healthcare Monitoring	0.194489	0.000000	100.00%	48.0%	100.0%	0.701

Table 9.5: Utility-preserving safety scorecard. ORIUS is compared with a degenerate fail-safe reference rather than only with unsafe persistence.

Domain	Fail-safe reference	Excess TSVR	Utility gain	Utility delta	Gate
Battery Energy Storage	immediate_shutdown	0.000000	nonzero/zero	10.100000	True
Autonomous Vehicles	always_brake	0.000000	3.985808	5869.636141	True
Medical and Healthcare Monitoring	always_alert	0.000000	nonzero/zero	142767.000000	True

Table 9.6: Final freeze and validation gates for the locked CPU release package.

Gate	Value	Pass
Split training	pass	yes
Model-quality release gate	13 release models, 0 blockers	yes
Downstream freeze validators	15 validators	yes
Runtime and stress gates	three domains, 12 stress families	yes
Release manifest	1307 hashed artifacts	yes

Table 9.4 is the governing safety table for the three-domain result: Battery and Healthcare close with zero ORIUS TSVR on the promoted runtime denominator, while AV is epsilon-closed on the bounded nuPlan replay/surrogate denominator. Table 9.5 answers the most important conservatism objection: ORIUS is compared against immediate shutdown, always-brake, or always-alert references and must preserve useful work without adding material excess TSVR over those fail-safe policies.




Figure 9.1: Final claim-governing runtime TSVR across Battery, AV, and Healthcare.

9.5 Domain synthesis

The point of the synthesis is not to suggest that every row is interchangeable. It is to show how the same release-boundary grammar survives contact with different plants, action spaces, and safety objects while each row keeps its own evidence ceiling.

Battery energy storage. The battery hazard is hidden constraint violation under degraded feeder, load, or state telemetry. ORIUS instantiates the runtime with tight state margins, explicit repair, and auditable release over a heavily instrumented control surface. Its universality contribution is to show that the full theorem-to-code-to-artifact chain can be closed on a real physical control problem.

Autonomous vehicles. The autonomous-vehicle hazard is unsafe longitudinal release under stale or degraded scene understanding. ORIUS instantiates the runtime through bounded




Figure 9.2: Final release-model PICP90 values for the manuscript-facing model-quality gate.




Figure 9.3: Useful work preserved by ORIUS over each domain’s degenerate fail-safe reference.

time-to-collision and predictive-entry-barrier logic, with explicit brake-oriented fallback and certificate-aware release checks. Its universality contribution is to show that the same release semantics remain meaningful in a fast-horizon, continuous-action mobility setting. The current claim-governing AV denominator is the all-zip grouped nuPlan replay/surrogate row with 1,531,104 runtime steps; it is not a road-deployment or full autonomous-driving closure claim.

Healthcare monitoring. The healthcare hazard is clinically unsafe release under partially observed physiological state, especially when alerts and interventions are triggered from noisy or degraded streams. ORIUS instantiates the runtime through physiological margin checks, alert fallback, and certificate-aware release semantics. Its universality contribution is to show that the same runtime grammar survives in a setting where the protected object is patient-facing monitoring behavior rather than mechanical actuation.




Figure 9.4: Cross-domain ORIUS validation surface. One runtime grammar carries across the active three-domain program, while each row remains limited by its own defended evidence surface.

9.6 Architecture universality versus promoted-row closure

Two conclusions follow from the domain synthesis. First, what transfers is not a controller family but a release-boundary grammar: reliability loss, uncertainty inflation, action repair,

fallback semantics, and certificate-aware release. Second, what remains local is the geometry of the safety object itself: a battery margin is not a vehicle headway constraint, and a vehicle headway constraint is not a physiological alert threshold.

That is why the book distinguishes architecture universality from promoted-row closure. Architecture universality is already defended by the shared adapter contract, kernel, benchmark schema, and governance surfaces. Promoted-row closure is an empirical achievement that requires each active row to clear the same replay, calibration, runtime, and governance gates.

9.7 Default proof mode for the dissertation

The default proof mode of the dissertation is therefore: *universal in architecture, tiered in evidence*. Battery calibrates the deepest witness depth. Healthcare and the bounded autonomous-vehicle runtime row show that the same contract is not confined to one plant family.

9.8 Boundary of the universality claim

This chapter is the place where the thesis must be most explicit. The universal claim is a claim about a reusable runtime grammar, not a claim that every conceivable physical-AI row has already been closed. The domain-closure matrix makes that difference visible in one artifact surface. In that exact sense, ORIUS already establishes a reusable runtime safety grammar across the active program while keeping the evidence boundary explicit.

Chapter 10

Temporal Certificates and Graceful Degradation

10.1 Certificate horizon as a runtime object

ORIUS certificates are temporal objects, not one-step labels. A certificate has a validity horizon, an expiration condition, and a bounded interpretation under worsening observation. This matters because degraded telemetry can remain acceptable for one step and become unacceptable only a short time later. The temporal question is therefore not “safe or unsafe forever,” but “how long does the runtime have before the current certificate no longer justifies release?”

T5 (Certificate Validity Horizon) formalises the constructive horizon definition, and T6 (Certificate Expiration Bound) supplies the supporting confidence-aware expiration proposition. Both are registered in the `THEOREM_REGISTER` (in `theoretical_guarantees.py`) and enforced at every runtime step: the API router computes the validity horizon τ_t and the expiry lower bound, appends them to the uncertainty metadata, and fails the guarantee check (`guarantee_passed = False`) whenever $\tau_t < 1$.

10.2 Why temporal safety belongs inside runtime assurance

This temporal view is one of the places where ORIUS aligns most clearly with classical runtime-assurance intuition. Simplex-style architectures distinguish between high- performance nominal control and a safer supervisory authority that must take over when the nominal contract is no longer justified [4, 5]. ORIUS extends that idea by attaching horizon

semantics directly to the certificate object. The runtime is not only deciding whether to intervene. It is also deciding how long the present release claim is allowed to remain live before fallback becomes mandatory.

10.3 Expiration, half-life, and safe-hold

The current evidence supports a conservative temporal reading. Horizon and half-life are most strongly defended as bounded expiration and safe-hold semantics under blackout-style degradation. That is the right level of claim: certificate time is measurable and governable, but it is not yet promoted into a universal law of every plant and sensor regime.




Figure 10.1: Blackout half-life surface used to motivate certificate expiration, bounded safe-hold, and temporal intervention discipline.

The half-life framing matters because it gives the runtime a language for degradation that is stronger than an unstructured timeout. A certificate can decay in justification even when the nominal policy remains capable of producing an action. ORIUS uses that decay to decide when safe-hold, curtailment, or reduced-operation behavior should replace ordinary release.

10.4 Severe-blackout protocol

The severe-blackout setting is the cleanest empirical test of temporal semantics because the runtime cannot rely on fresh observation arriving in time to rescue nominal control. The question becomes: can the runtime preserve a governed reduced-operation mode long enough to avoid unsafe release while still remaining operationally legible? That is why the blackout protocol belongs in the main body. It is not a corner-case stress test. It is the clearest demonstration that certificates are temporal objects with behavioral consequences.

10.5 Graceful degradation policies

When no full-quality control action remains justified, ORIUS benchmarks degradation policies explicitly rather than treating them as informal emergency behavior. This is the runtime-assurance connection in practice: the system needs a governed way to do less, safer, and more explainably when observation quality collapses.




Figure 10.2: Four-policy graceful degradation comparison. The dissertation uses this family to treat fallback quality as a benchmarked layer instead of an unexamined emergency default.

Benchmarking degradation policies matters scientifically because a runtime could always buy safety by immediately collapsing to a maximally conservative shutdown or no-op. ORIUS asks a more demanding question: what bounded policy still preserves safety while keeping the system operationally interpretable and certificate-governed?

10.6 Temporal extension

This chapter extends the same runtime object introduced earlier rather than adding a separate theory of fallback. Certificate time and graceful degradation become measurable and

governable parts of the release object itself. The defended claim remains bounded: time-to-expiration, safe-hold, and fallback quality now belong to the scientific evaluation surface of ORIUS, but no universal optimal degradation policy is claimed here.

Chapter 11

Compositional Safety

11.1 Why local safety does not automatically compose

One of the easiest mistakes in safety argumentation is to assume that individually safe controllers remain jointly safe when they share resources or constraints. ORIUS treats composition as a bounded extension precisely because this implication fails in general. Local certificates can remain valid while shared budgets, timing windows, or coupled plant dynamics still create system-level failure.

11.2 A simple counterexample

Consider two locally admissible controllers that each satisfy their own state and action constraints. If both consume a shared feeder budget, corridor occupancy, or actuator authority reserve, their simultaneous release can still violate the system-level constraint even though neither local controller is individually unsafe. This is the compositional version of OASG: local legality can diverge from system-level safety when the shared constraint is not represented in the local release object.

11.3 Shared constraints as the right composition surface

Composition is therefore framed around shared constraints rather than around an immediate claim of universal multi-agent closure. Shared power budgets, corridor occupancy, common actuator authority, or coupled timing envelopes are the right first-class objects because they say where local certificates stop being sufficient. ORIUS can then lift the certificate object

into a negotiated setting without pretending that single-agent repair theorems automatically solve the multi-agent case.

Table 11.1: Multi-Agent Fleet Coordination: joint safety violations and useful work under shared transformer capacity constraint (limit = 80 MW, two batteries at 60 MW each = 120 MW proposed). Independent protocol violates; ORIUS-coordinated eliminates violations.

Protocol	Joint Violations	Local Violations	Useful Work (MWh)	Margin C
Independent (non-compositional)	24	0	1920	
Centralized (ORIUS-coordinated)	0	0	1920	
Distributed Negotiation	0	0	1920	

11.4 Multi-agent certificates and negotiation

The composition story in ORIUS is therefore narrow but meaningful. A local certificate says that an action is safe given one agent’s uncertainty set and fallback semantics. A compositional certificate must additionally encode the shared constraint on which multiple agents depend. This turns negotiation and fairness from policy afterthoughts into runtime objects. If the shared constraint cannot be satisfied jointly, the runtime must decide which agents may continue, which must degrade, and which fallbacks remain legitimate.

11.5 Why composition remains bounded in the dissertation

Composition matters scientifically because it shows that the ORIUS grammar does not end at single-agent control. At the same time, it remains outside the strongest universal claim. The present dissertation does not defend a composition theorem beyond the active three-domain lane. It shows that shared constraints can be represented through the same runtime object and that local certificates alone are not sufficient once coupling enters the system.

11.6 Compositional takeaway

What this chapter contributes is a bounded extension, not a new center of gravity for the thesis. ORIUS can lift degraded-observation safety into a shared-constraint setting by making system-level resources explicit in the release object. It does not claim that local safety

composes automatically, and it does not claim that every current domain row already clears a defended multi-agent promotion discipline.

Chapter 12

CertOS Runtime Assurance

12.1 Why CertOS belongs in the main argument

CertOS is the part of ORIUS that turns repaired action into a governed runtime object. Without lifecycle state, fallback rules, and an audit chain, certificates would be inert metadata. With CertOS, certificate validity becomes part of the release decision itself: which action was allowed, under which reliability state, for which horizon, and under which fallback if the certificate expires or becomes invalid?

This chapter therefore does not read as a deployment appendix. It is the point at which the ORIUS kernel acquires operational semantics. `DomainAdapter` defines the local safety predicate and fallback object that CertOS must protect, and the universal benchmark schema defines the replay traces from which certificate validity, latency, and lifecycle behavior are evaluated.

12.2 Lifecycle operations

CertOS organizes runtime release through a small but strict lifecycle:

- issue a certificate only after the repaired or fallback action is shown admissible;
- validate a certificate before each protected release;
- renew a certificate when fresh evidence preserves the release claim;
- expire or revoke a certificate when the underlying assumptions no longer hold; and
- force explicit fallback when no valid certificate remains.

These operations make ORIUS operationally legible. The runtime is not merely deciding whether an action is numerically acceptable; it is managing the validity of the claim under which release occurs.

12.3 Runtime invariants

The lifecycle logic is narrow and strong. A protected action is never released without a valid certificate or an explicit fallback. Certificate state must be renewable, expirable, revocable, and auditable. Runtime traces must preserve enough information to reconstruct why release was allowed and when the governing assumptions ceased to hold. These are not administrative conveniences. They are the invariants that keep repaired action tied to a reviewable safety argument.




Figure 12.1: CertOS lifecycle and runtime-governance surface across ORIUS domains. The figure shows how certificate issue, validation, expiry, fallback, and audit continuity are exposed as part of runtime safety rather than as post hoc reporting.

12.4 Audit chain and post hoc verification

The audit chain matters because ORIUS is not only a forward-looking safety layer. It is also a mechanism for reconstructing the justification of past release decisions. Once an unsafe event, intervention, or fallback occurs, the system must be able to explain what the reliability state was, which repair or fallback was chosen, what certificate was valid, and why the runtime believed the action remained admissible. This is why audit completeness is a first-class benchmark metric rather than a logging implementation detail.

12.5 Governance alignment, not compliance

CertOS aligns the runtime evidence surface with disciplined safety-argument practice without claiming compliance with external standards. In that limited and precise sense, the governance chapter is legible to readers who care about claim-based safety arguments, lifecycle discipline, and trustworthy AI governance [9–12].

The distinction matters. UL 4600, the NIST AI RMF, ISO 26262, and IEC 61508 are used here as evidence-structure and lifecycle analogs only. They motivate why a runtime safety layer should make claims, assumptions, evidence, lifecycle transitions, and audit continuity explicit. They do not turn the dissertation into a compliance claim, and the manuscript must never be read that way.

12.6 Universality and evidence-tiered maturity

CertOS is universal at the level of release semantics. Every defended ORIUS row needs a certificate object, fallback policy, and audit trace. What is not yet universal is equal deployment maturity across all future domains. Battery provides the deepest defended lifecycle surface. The autonomous-vehicle and healthcare rows show defended bounded runtime surfaces under current artifacts. CertOS therefore supplies one common governance object without pretending that all future deployment surfaces are equally mature today.

12.7 Governance boundary

This chapter is one of the places where the non-claim boundary has to remain explicit. ORIUS provides a policy-driven runtime governance layer in which repaired action, fallback, and certificate validity are coupled, auditable, and recoverable after the fact. That claim is scientifically important because a runtime safety layer without release semantics cannot say what action was justified once degraded observation has already entered the control loop. It does not claim formal certification compliance, and it does not claim equal deployment maturity across every current domain row.

Part V

Submission Boundary

Chapter 13

Governance, Reproducibility, Limitations, and Conclusion

13.1 What is proved

The dissertation proves a bounded but durable theory claim. Under explicit assumptions, degraded observation can create a gap between observed-state legality and true-state safety; reliability-conditioned uncertainty inflation changes the admissible release set; and repair-or-fallback discipline preserves safety whenever the inflated set and lifecycle semantics remain feasible. Those results live in the theorem chapter and proof appendix and should not be overstated beyond their assumptions.

13.2 What is empirically validated

The empirical surface is strongest at the witness and defended bounded rows. Battery remains the deepest theorem-to-code-to-artifact chain. Healthcare monitoring and the autonomous-vehicle row support defended bounded interpretations under the current replay and governance surfaces.

Those limits matter because ORIUS is not presented here as a universal optimal controller or as a finished deployment program for every physical-AI setting. It is presented as a runtime safety layer whose strongest current evidence is deep in one witness domain and bounded across two peer rows.

13.3 What is out of scope

The dissertation does not claim:

- empirical closure beyond the active Battery + AV + Healthcare program;
- universal statistical calibration under arbitrary shift or adversarial calibration-data poisoning;
- sub-millisecond inner-loop safety filters (DC_SS acts as an outer-loop supervisor at 5–10 Hz, constrained by runtime optimization latency);
- closed-loop reactive physics for Autonomous Vehicles (the AV evaluation uses open-loop trace surrogate replay to isolate the OASG hazard, bounding it short of a live ROS2/CARLA integration);
- hardware control-authority preservation (ORIUS assumes actuation mechanisms remain physically viable even when observation degrades);
- production deployment readiness in every plant family; or
- formal certification compliance under UL 4600, NIST AI RMF, ISO 26262, or IEC 61508.

The standards language used in the governance chapter is an evidence-structure alignment and lifecycle-analogy argument only. The deployment language is deliberately bounded by the domain-closure matrix, runtime/governance surfaces, and explicit non-claim statements. For the healthcare row, this boundary follows the public critical-care dataset lineage of MIMIC-IV and eICU [47–49] and the reporting/applicability discipline of TRIPOD+AI, PROBAST+AI, CONSORT-AI, DECIDE-AI, and STARD-AI [50–54]. Those sources support transparent retrospective reporting and future clinical-evaluation design; they do not convert the current ORIUS healthcare row into live clinical deployment.

Table 13.1: Validation boundary for the current ORIUS evidence package.

Evidence tier	Current status	Claim allowed	Claim not allowed
Battery runtime	completed locked witness runtime	theorem-grade witness plus bounded software/HIL-style evidence	physical field deployment or certified grid operation
nuPlan AV replay	completed all-zip grouped runtime replay/surrogate	bounded AV runtime-denominator evidence	road deployment or full autonomous-driving closure
CARLA closed loop	not completed in this PDF	future controllable stress-test tier	completed simulator closure
Healthcare MIMIC runtime	completed retrospective monitoring row	bounded alert-release monitoring evidence	live clinical deployment or clinical decision-support approval
eICU/source holdout	not staged/completed in this PDF	future external source-holdout tier	completed multi-site clinical validation
Physical battery HIL/bench	not completed in this PDF	future physical validation tier	physical HIL, bench, or field validation

13.4 Reproducibility discipline

Every headline claim stays tied to tracked figures, tables, manifests, or claim matrices so that the defended surface remains auditable after the fact. The purpose of that discipline is simple: a universal safety thesis should be inspectable at the level of artifacts, not only persuasive at the level of prose.

That is why the dissertation remains chapter-first in prose but artifact-first in evidence. The prose chapters explain the argument. The claim validator, publication matrices, release manifests, and replay tables control what the prose is allowed to say. In that exact sense, the dissertation is universal in architecture, tiered in evidence.

13.5 Threats to validity

The largest threats to validity are asymmetry and interpretation. Asymmetry arises because the witness row is much deeper than the other promoted rows. Interpretation risk arises because readers may confuse structural universality with broader empirical closure. For example, while DC_S mathematically proves Feasible Fallback (Theorem 7), if a vehicle’s physical brakes are severed (an actuation failure), no epistemic observation wrapper can save it. Furthermore, the Conformal Coverage guarantees (Theorem 3) implicitly assume the historical calibration dataset X_{calib} has not been systemically poisoned. If calibration data inherently contains degraded telemetry, the resulting probability intervals will silently fail to bound the constraint geometry.

The dissertation addresses the first risk by keeping the witness role explicit and by preserving the domain-closure matrix in the main narrative. It addresses the interpretation risk by strictly scoping the actuation authority boundaries, declaring the AV trace-replay limitations, and separating proved, validated, and roadmap claims at the major boundary chapters.

13.6 Conclusion

The argument of this dissertation can now be stated plainly. Degraded observation is not a peripheral nuisance attached to otherwise reliable autonomy; it is a release-boundary safety problem in its own right. Once observed-state legality and true-state legality can diverge, safe actuation requires more than nominal intelligence. It requires explicit reliability estimation, conservative state widening, repaired action or bounded fallback, and certificate-aware release.

ORIOUS makes that claim operational through three connected surfaces. A `DomainAdapter` exposes the local safety predicate and fallback object. The universal benchmark schema exposes one replay contract for truth, observation, intervention, fallback, certificate validity, and latency. CertOS supplies the lifecycle semantics that determine when repaired action may still be released. Together, those surfaces turn degraded observation from a vague systems concern into a measurable and governable runtime object.

That is the durable contribution of the thesis. ORIOUS does not ask every domain to share a controller, plant model, or optimization objective. It asks every domain to make the release boundary explicit. In doing so, it turns true-state violation, intervention burden, fallback usage, certificate validity, and runtime latency into scientific objects that can be compared across domains without pretending the domains themselves are identical.

The final thesis claim is therefore ambitious and disciplined at the same time. ORIOUS already supports a reusable runtime safety layer for physical AI under degraded observation. Its active repo-tracked program is intentionally limited to Battery, AV, and Healthcare. In that exact sense, structural universality is defended now while the empirical boundary stays explicit. That boundary is not a rhetorical retreat. It is part of what makes the argument credible.

Part VI

Archival Provenance Index

Chapter 14

Archival Source Register

Earlier drafts of this repository compiled a polished thirteen-chapter submission spine, a battery-era chapter archive, a merged historical narrative, longform extension packets, and standalone proof PDFs into one PDF. That made the document complete as an archive but repetitive as a monograph: the introduction, related work, theorem ladder, universality boundary, proof surface, and conclusion appeared multiple times in different historical voices.

This senior-review controller keeps the main body as one continuous monograph. Historical sources remain tracked and inspectable, but they are no longer compiled inline as duplicate narrative chapters. When archive material is needed for provenance or review traceability, the source path below is the canonical location.

14.1 Tracked archive locations

- `chapters/`: battery-era chapter archive and earlier theorem development.
- `chapters_merged/`: merged historical narrative draft, including older broader-domain validation text that is not part of the current defended claim surface.
- `paper/longform/`: longform extension packets for conditional coverage, Byzantine robustness, committee response, portability, and deeper theoretical guarantees.
- `appendices/app_c_flagship_proofs.pdf` and `appendices/app_c_all_theorems.pdf`: standalone proof packets.
- `appendices/`: legacy appendix packets, artifact registers, theorem audits, release manifests, and reproducibility records.

14.2 Reader-facing rule

The defended scientific flow is the preceding main body plus the curated system appendices that follow. Archive sources are evidence history, not additional main-argument chapters. This prevents older wording from silently restating or overstating the current Battery + AV + Healthcare claim boundary.

Part VII

System Appendices

Appendix a

Notation Sheet

This appendix collects every symbol used in the ORIUS battery-safety framework. Symbols are grouped by role: telemetry and state variables, quality and drift indicators, uncertainty and action objects, battery plant parameters, metrics, and tuning constants.

a.1 Telemetry and State Variables

Table a.1: Telemetry and state symbols.

Symbol	Domain	Description
z_t	$\mathbb{R}^n$	Clean (ground-truth) telemetry vector at step t .
$\tilde{z}_t$	$\mathbb{R}^n$	Degraded telemetry vector after dropout, delay, or drift.
o_t	$\mathbb{R}^n$	Observed state passed to the controller.
x_t	$\mathbb{R}^n$	True (latent) physical state; not directly accessible.
$\hat{x}_t$	$\mathbb{R}^n$	State estimate reconstructed from $\tilde{z}_t$.
ξ_t	$\mathbb{R}^n$	Observation error: $\xi_t = \hat{x}_t - x_t$.
ϵ_t	$\mathbb{R}$	One-step model (dynamics) prediction error.

a.2 Quality, Drift, and Sensitivity Indicators

Table a.2: Quality and detector symbols.

Symbol	Domain	Description
w_t	$[0, 1]$	Observation-quality (reliability) score. Aggregates freshness, consistency, and drift evidence.
d_t	$\{0, 1\}$	Drift-detected flag from the Page–Hinkley detector.
s_t	$\mathbb{R}_+$	Sensitivity / staleness indicator; large s_t implies stale data.

a.3 Uncertainty and Action Objects

Table a.3: Uncertainty-set and action symbols.

Symbol	Domain	Description
$\hat{C}_t$	interval	Base conformal prediction interval at coverage $1 - \alpha$.
$\hat{C}_t^{\text{RAC}}$	interval	Reliability-adaptive conformal interval after RUI inflation.
U_t	set	Uncertainty set for robust dispatch: $U_t = [\hat{x}_t - m_t, \hat{x}_t + m_t]$.
m_t	$\mathbb{R}_+$	Half-width (margin) of the uncertainty set: $m_t = q_t / (w_t + \epsilon)$.
a_t^*	$\mathbb{R}^m$	Candidate action proposed by the base controller.
a_t^{safe}	$\mathbb{R}^m$	Repaired action after DC3S safety projection.
$\mathcal{A}_t$	set	Tightened safe-action set under reliability-aware constraints.
$\phi(\cdot)$	$\{0, 1\}$	Safety predicate: $\phi(x) = 1$ iff x satisfies all constraints.

a.4 Battery Plant Parameters

Table a.4: Battery plant parameters.

Symbol	Domain	Description
SOC_t	$[0, E_{\max}]$	State of charge at step t (MWh).
$E_{\max}$	$\mathbb{R}_+$	Rated energy capacity (MWh).
$P_{\max}$	$\mathbb{R}_+$	Maximum charge/discharge power (MW).
$R_{\max}$	$\mathbb{R}_+$	Maximum ramp rate (MW/step).
η_c	$(0, 1]$	Charging efficiency.
η_d	$(0, 1]$	Discharging efficiency.
Δt	$\mathbb{R}_+$	Dispatch interval length (hours).

a.5 Metrics

Table a.5: Evaluation metrics.

Symbol	Domain	Description
TSVR	$[0, 1]$	True-SOC Violation Rate: fraction of steps with $\text{SOC}_t \notin [\text{SOC}_{\min}, \text{SOC}_{\max}]$.
IR	$[0, 1]$	Intervention Rate: fraction of steps where DC3S repairs the action.
OASG	$[0, 1]$	Observed-safe / Actually-unsafe State Gap rate.
PICP	$[0, 1]$	Prediction-Interval Coverage Probability.

a.6 Tuning Constants and Greek Letters

Table a.6: Tuning constants and thresholds.

Symbol	Domain	Description
α	$(0, 1)$	Conformal miscoverage level (target: 0.10).
κ_r	$\mathbb{R}_+$	Reliability decay rate in the OQE score.
κ_d	$\mathbb{R}_+$	Drift penalty coefficient.
κ_s	$\mathbb{R}_+$	Staleness penalty coefficient.

$w_{\min}$	$(0, 1)$	Quality-score floor to prevent division by zero.
q_t	$\mathbb{R}_+$	Conformal quantile at step t .
ϵ	$\mathbb{R}_+$	Small constant in margin denominator: $m_t = q_t / (w_t + \epsilon)$.

a.7 Set Notation and Operators

Table a.7: Set notation and mathematical operators.

Symbol	Description
$\mathbb{R}, \mathbb{R}_+$	Real numbers; non-negative reals.
$\mathcal{U}$	Power- and ramp-feasible action set (before safety tightening).
$\mathbb{P}[\cdot]$	Probability measure.
$\mathbb{E}[\cdot]$	Expectation operator.
$\arg \min, \arg \max$	Argument of the minimum / maximum.
$\mathbf{1}[\cdot]$	Indicator function.
$ \cdot $	Absolute value or set cardinality, depending on context.
$\ \cdot\ $	Euclidean norm unless otherwise stated.

Appendix b

Assumption Register

This appendix records the assumptions that delimit the dissertation’s strongest claims. The purpose is not to hide caveats in the back matter. It is to make the theorem and evidence boundary explicit enough that every promoted claim can be traced to the assumptions it needs, the diagnostics that check those assumptions, and the fallback interpretation used when the assumptions become only partially satisfied.

Table b.1: ORIUS dissertation assumption register.

ID	Name	Meaning	Diagnostics / evidence	Where it matters
A1	Bounded observation error surface	Reliability-conditioned uncertainty produces a finite observation-consistent state set.	Reliability summaries, calibration diagnostics, and bounded widening logic in the theorem gate.	Chapters 3, 5, and Appendix h.
A2	Feasible safe repair or fallback	When the admissible set collapses, the run-time can still choose a certified fallback action or safe-hold policy.	Repair traces, fallback tables, and CertOS lifecycle surfaces.	Chapters 4, 10, and 12.

ID	Name	Meaning	Diagnostics / evidence	Where it matters
A3	Causal certificate issuance	Certificates depend only on information available at runtime up to the current step.	Certificate payload inspection, audit chain reconstruction, and lifecycle verification.	Chapters 4, 12, and Appendix e.
A4	Dual-trajectory replay fidelity	The benchmark exposes both observed and true safety surfaces strongly enough to measure OASG.	Replay manifests, deterministic fault schedules, and domain closure artifacts.	Chapters 6, 8, and 9.
A5	Adapter faithfulness	Each domain adapter encodes the relevant safety predicate, fallback semantics, and certificate payload honestly for its plant.	Domain closure matrix, protocol-card review surfaces, and non-claim boundaries for weaker rows.	Chapters 4, 9, and 12.
A6	Evidence-tier discipline	Architectural universality and empirical parity are not allowed to collapse into one claim.	Parity matrix, submission score-card, gap matrix, and claim validator.	Chapters 1, 9, and 13.

The dissertation’s formal claims should always be read together with this table. When the evidence surface is weaker than the architectural surface, the text is intentionally written as a bounded semantic or governance statement rather than as a broader future-domain validation claim.

Appendix c

Proofs

This appendix records the proof-facing structure of the dissertation. It is not intended to replace the theorem chapter; it is intended to make the proof obligations inspectable and to show which claims are formal, which are empirical, and which rely on shared typed runtime objects.

c.1 Proof map

Table c.1: Theorem-to-proof map for the ORIUS dissertation.

Result	Main idea	Depends on	Artifact-facing link
Theorem 5.2	If observed legality is not informationally sufficient for true-state safety, then there exists a degraded-observation regime that admits an observed-safe but true-unsafe release.	A1, A4, A5	Dual-trajectory replay and the OASG metric family.
Theorem 5.3	If inflation is conservative and the repair or fallback operator respects the tightened set, the released action preserves the defended safety predicate.	A1, A2, A3, A5	Repair traces, fallback logs, and certificate-valid release surfaces.

Result	Main idea	Depends on	Artifact-facing link
Theorem 5.4	As reliability falls, uncertainty inflation and admissible-set tightening contract the release envelope lawfully rather than heuristically.	A1, A3, A6	Calibration diagnostics, theorem gate figure, and runtime intervention traces.
Proposition 5.5	Quality-ignorant release can remain observed-safe while incurring true-state failure.	A1, A4, A5	Witness-row violation tables and fault-family replay.
T10_T11_ObservationAmbiguitySandwich	ObservationAmbiguitySandwich release has a Bayes ambiguity-risk lower bound, while ORIUS has an α -covered upper bound under correct uncertainty-set coverage.	A1–A6 plus coverage, safe-set correctness, repair membership, fallback admissibility, and adapter validity	src/orius/universal_theory/observation_ambiguity_active_theorem_audit

c.2 Proof sketches

OASG existence. The proof proceeds by separating observed-state legality from true-state legality. Once the adapter exposes a true-state violation predicate and the replay harness exposes an observed state that is not informationally sufficient for that predicate, one can construct a degradation regime in which observed legality survives while true-state safety fails. The scientific role of the theorem is not to characterize every such regime exhaustively. It is to show that the gap is structural rather than anecdotal.

Safety preservation under repair. The second theorem uses the conservativeness of the inflated observation-consistent set. If the repair operator only emits actions admissible over that set, or the fallback operator replaces release when such action does not exist, then the released action remains safe for the latent state whenever the inflation really upper-bounds that state. This is the exact point at which fallback becomes part of the proof surface rather than an implementation afterthought.

Degradation-sensitive risk envelope. The third theorem formalizes the intuition that reduced trust must shrink release freedom. Monotone reliability loss implies monotone widening of the uncertainty set; that widened set then implies monotone tightening of admissible release. The theorem does not say intervention is free. It says intervention pressure is a lawful consequence of degraded observation.

No-free-safety proposition. The proposition follows by contradiction against any controller that checks legality only on observed state while ignoring observation degradation. If the degraded observation regime is admissible yet insufficiently informative about the latent state, then there exists a release that remains locally coherent while still violating the true safety predicate.

Covered observation-ambiguity sandwich. Let $L(a, o) = \mathbb{P}(a \notin C(X^*) \mid O = o)$ be the conditional true-state violation risk of releasing action a after observing o . Any observation-only mandatory-release controller π selects a single action $\pi(o)$ on the entire ambiguity class $B(o) = \{x : h(x) = o\}$. Therefore, for each observation value o ,

$$L(\pi(o), o) \geq \min_{a \in \mathcal{A}} L(a, o).$$

Taking expectation over O gives

$$\text{TSVR}(\pi) = \mathbb{E}_O[L(\pi(O), O)] \geq \mathbb{E}_O\left[\min_{a \in \mathcal{A}} \mathbb{P}(a \notin C(X^*) \mid O)\right].$$

This proves the Bayes lower bound. If, for some o , the common safe core $K(o) = \cap_{x \in B(o)} C(x)$ is empty and the conditional law of X^* puts positive mass on states that exclude each candidate action, then the conditional minimum is strictly positive. If $K(o)$ is nonempty, the lower bound correctly allows zero conditional risk; this is why the theorem is not the false claim that all ambiguity forces violation.

For the ORIUS upper bound, suppose the released action u_t satisfies $u_t \in C(x)$ for every $x \in \mathcal{X}_t$ and $\mathbb{P}(X_t^* \notin \mathcal{X}_t) \leq \alpha$. Split the violation event:

$$\mathbb{P}(u_t \notin C(X_t^*)) = \mathbb{P}(u_t \notin C(X_t^*), X_t^* \in \mathcal{X}_t) + \mathbb{P}(u_t \notin C(X_t^*), X_t^* \notin \mathcal{X}_t).$$

The first term is zero by set-wise release safety. The second term is bounded by $\mathbb{P}(X_t^* \notin \mathcal{X}_t) \leq \alpha$. Hence $\text{TSVR}_{\text{ORIUS}} \leq \alpha$, and exact coverage gives zero one-step true-state violation probability.

c.3 Interpretation

These proofs are deliberately modest in scope. They do not prove broader future-domain closure and they do not replace the need for replay, calibration, runtime, and governance artifacts. Their role is to anchor the manuscript’s strongest claims to explicit assumptions and typed objects so that later empirical chapters can be read as validation of the same runtime contract rather than as isolated case studies.

Appendix d

Artifact and Claim Index

This appendix records the canonical bridge from dissertation claims to governed repository surfaces. The prose chapters are curated, but every headline statement is expected to land on a theorem object, publication matrix, replay artifact, or explicit non-claim surface.

d.1 Claim classes and governing surfaces

Table d.1: Claim classes and their governing artifact surfaces.

Claim class	Primary surface	What governs it
Proved	Typed universal objects, assumption register, proof appendix	Theorems must point to assumptions, typed runtime objects, and the theorem-to-artifact bridge.
Validated	Publication matrices, replay tables, domain closure rows	Claims may advance only when the corresponding replay, calibration, runtime, and governance artifacts exist.
Governance	CertOS lifecycle surfaces, runtime/governance matrices	Certificate lifecycle, fallback semantics, and audit continuity remain first-class evidence, not narrative-only claims.
Roadmap	Explicit non-claims, parity blockers, promotion obligations	Future work is useful only when it stays outside headline proof or validation language until artifacts close the gap.

d.2 Headline claim families

Table d.2: Headline claim families and their canonical evidence surfaces.

Claim family	Chapter anchor	Canonical evidence surfaces
Degraded-observation hazard	Chapters 1–3	Typed adapter semantics, dual-trajectory benchmark contract, and OASG-facing replay logic.
Theorem ladder	Chapter 5	Assumption register, proof appendix, theorem dependency table, integrated theorem gate figure.
Witness-domain validation	Chapters 7–8	Main results, ablations, grouped coverage tables, learned-reliability surfaces, runtime traces.
Cross-domain universality	Chapter 9	Parity matrix, domain closure matrix, cross-domain validation figure, runtime/-governance matrices.
Temporal and governance semantics	Chapters 10 and 12	Blackout half-life figure, graceful degradation figure, governance lifecycle matrix, runtime budget matrix.
Exact non-claims and parity blockers	Chapter 13	Submission scorecard, deployment validation scope, 93-plus gap matrix, parity matrix.

d.3 Canonical manuscript-facing evidence

The following generated surfaces define the defended manuscript state at any given time:

- `reports/publication/orius_equal_domain_parity_matrix.csv`
- `reports/publication/orius_submission_scorecard.csv`
- `reports/publication/orius_governance_lifecycle_matrix.csv`
- `reports/publication/orius_runtime_budget_matrix.csv`
- `reports/publication/orius_deployment_validation_scope.csv`

- `reports/publication/orius_93plus_gap_matrix.csv`
- `reports/publication/final_training_quality_for_paper.csv`
- `reports/publication/final_runtime_safety_for_paper.csv`
- `reports/publication/utility_preserving_safety_scorecard.csv`
- `reports/publication/final_freeze_validation_for_paper.csv`
- `reports/publication/final_paper_results_summary.json`

d.4 Interpretation

The canonical readiness summary lives in Chapter 9 because parity promotion is part of the main scientific narrative, not an appendix-only afterthought. This appendix records the governing surfaces that support those statements and keeps exact non-claims visible so the dissertation cannot silently promote a bounded row into a stronger evidence class.

Appendix e

Safety-Case Bundle Index

This appendix maps the dissertation to a safety-case style evidence bundle without making any compliance claim. The goal is simple: a reviewer should be able to see which hazard, which argument, which evidence packet, and which non-claim boundary belong together. The appendix is therefore an evidence-structure aid, not a sector-specific certification claim [9–12].

Table e.1: Safety-case bundle index for the ORIUS dissertation.

Bundle		Dissertation role	Primary evidence surfaces
Hazard	definition	Defines degraded observation and OASG as the governing safety problem.	Chapters 1–3, universal claim matrix, assumption register.
Theorem bundle		States what is proved under explicit assumptions.	Chapter 5, Appendix c, theorem dependency table, integrated theorem gate figure.
Benchmark bundle		Defines replay, dual trajectories, metrics, and fault families.	Chapter 6, Appendix h, domain closure matrix, submission scorecard.
Witness bundle		Shows the deepest theorem-to-artifact row.	Chapters 7–8, battery tables and figures, release manifest.
Cross-domain bundle		Shows structural universality and bounded transfer.	Chapter 9, submission scorecard, runtime/governance matrices.

Bundle	Dissertation role	Primary evidence surfaces
Governance bundle	Shows certificate lifecycle, runtime budgets, fallback, and audit continuity.	Chapter 12, governance lifecycle matrix, runtime budget matrix, review-facing governance derivatives.
Boundary bundle	States what remains gated, roadmap, or non-promotional.	Chapter 13, deployment validation scope, gap matrix, exact non-claim prose.

e.1 How to read the bundle index

The bundle index is intentionally orthogonal to the chapter order. A chapter can contribute to more than one bundle, and a bundle can span multiple generated artifacts. That structure mirrors real safety-argument practice more closely than a chapter-by-chapter checklist: one claim family may rely simultaneously on a theorem statement, a replay trace, a runtime lifecycle table, and an explicit non-claim boundary.

e.2 Interpretive boundary

The dissertation uses this bundle index only as an evidence-structure aid. It aligns the manuscript with disciplined safety-argument practice and lifecycle-oriented risk governance without implying certification or sector-specific regulatory compliance. Its role is to keep the argument inspectable, not to relabel the dissertation as a compliance submission.

Appendix f

Extended Battery Result Tables

This appendix presents the full result tables that support the main-text witness-domain summaries in Chapters 7 and 8. Every number traces to a locked CSV in `reports/publication/`.

f.1 Extended Controller Comparison

Table f.1 reports every metric for all five controllers under the default fault profile (`drift_combo`, 20 % dropout, DE region).

Table f.1: Extended controller comparison (DE, `drift_combo`, 20 % dropout). Source: `dc3s_main_table.csv`.

Controller	TSVR (%)	IR (%)	PICP (%)	Sev. p_{95}	Revenue	Δ Rev. (%)	Latency (ms)
Det. LP	3.9	0.0	—	0.42	1.000	0.0	1.2
Robust Fixed	0.8	12.1	—	0.11	0.931	−6.9	1.4
CVaR	1.2	8.4	—	0.18	0.952	−4.8	2.1
DC3S	0.0	14.7	91.3	0.00	0.948	−5.2	3.8
DC3S + FTIT	0.0	13.2	92.1	0.00	0.956	−4.4	4.1

Revenue is normalised to the Det. LP baseline. Severity p_{95} is the 95th-percentile true-SOC violation magnitude (MWh, normalised by $E_{\max}$). Latency is median per-step wall-clock time.

f.2 Metric Definitions

For completeness, the extended metrics reported above are:

TSVR

True-SOC Violation Rate (Appendix a).

IR Intervention Rate: fraction of steps where the shield modifies the candidate action.

PICP

Prediction-Interval Coverage Probability of the conformal interval on held-out data.

Sev. p_{95}

95th percentile of the violation severity distribution; zero when TSVR = 0.

Revenue

Normalised arbitrage revenue over the evaluation horizon.

Latency

Median per-step wall-clock time including OQE, uncertainty inflation, shield, and certificate emission.

f.3 Regional Impact Breakdown

Table f.2 disaggregates the DC3S results by region (DE = Germany, US = CAISO/ERCOT composite).

Table f.2: Regional impact breakdown for DC3S + FTIT. Source: `dc3s_main_table.csv`, `48h_trace_final_de.csv`, `48h_trace_final_us.csv`.

Region	TSVR (%)	IR (%)	PICP (%)	Sev. p_{95}	Δ Rev. (%)	Latency (ms)
DE (Germany)	0.0	13.2	92.1	0.00	−4.4	4.1
US (CAISO/ERCOT)	0.0	15.8	90.7	0.00	−5.1	4.3

f.4 Fault-Profile Sensitivity

Table f.3 shows TSVR and IR across fault families at the 20% dropout level.

f.5 Artifact Provenance

All figures referenced in the main text trace to locked publication-report artifacts:

- `fig_true_soc_violation_vs_dropout.png` — TSVR vs. dropout rate.
- `fig_true_soc_severity_p95_vs_dropout.png` — severity surface.

Table f.3: DC3S + FTIT under individual fault families (DE, 20 % dropout). Source: `fault_performance_table.csv`.

Fault family	TSVR (%)	IR (%)	PICP (%)	Sev. p_{95}
dropout	0.0	11.4	93.2	0.00
delay_jitter	0.0	10.1	92.8	0.00
spike	0.0	16.3	90.4	0.00
stale	0.0	12.7	91.6	0.00
drift	0.0	14.9	91.1	0.00
reorder	0.0	9.8	93.0	0.00
drift_combo	0.0	13.2	92.1	0.00

- `fig_calibration_tradeoff.png` — coverage-width Pareto front.
- `fig_transfer_coverage.png` — transfer coverage.
- `fig_coverage_width.png` — coverage vs. interval width.
- `fig_48h_trace_final_de.png` — 48 h operational trace (DE).
- `fig_48h_trace_final_us.png` — 48 h operational trace (US).

Appendix g

Reliability-Group Coverage Audits

This appendix documents the stratified coverage audit that verifies the conformal guarantee holds *within* each reliability bin, not only in aggregate. The audit is a key defence against the scenario where aggregate PICP hides systematic under-coverage for low-reliability episodes.

g.1 Audit Logic

The reliability-bin stratification proceeds as follows:

1. Partition all evaluation steps into K bins by their OQE reliability score w_t . Default: $K = 5$ with edges $\{0.0, 0.2, 0.4, 0.6, 0.8, 1.0\}$.
2. Within each bin k , compute the empirical coverage:

$$\widehat{\text{PICP}}_k = \frac{1}{|B_k|} \sum_{t \in B_k} \mathbf{1}[x_t \in \hat{C}_t^{\text{RAC}}],$$

where $B_k = \{t : w_t \in [\ell_k, u_k)\}$.

3. Flag bin k as **under-covered** if $\widehat{\text{PICP}}_k < 1 - \alpha - \delta_{\text{audit}}$, where $\delta_{\text{audit}} = 0.05$ is the audit tolerance.
4. The audit **passes** if no bin is flagged.

The tolerance δ_{audit} accounts for finite-sample fluctuation in small bins and is calibrated via a Clopper–Pearson binomial interval at the 95 % level.

g.2 Illustrative Audit Table

Table g.1 shows per-bin coverage for DC3S + FTIT under `drift_combo` at 20% dropout (DE region).

Table g.1: Coverage by reliability bin (DE, `drift_combo`, 20% dropout). Source: `cqr_group_coverage.csv`.

Reliability bin	n steps	$\widehat{\text{PICP}}$ (%)	Avg. width	Flag
[0.0, 0.2)	184	94.6	0.38	—
[0.2, 0.4)	271	93.0	0.29	—
[0.4, 0.6)	402	92.3	0.22	—
[0.6, 0.8)	518	91.7	0.17	—
[0.8, 1.0]	793	91.1	0.12	—
All	2168	92.1	0.19	Pass

Key observations:

- Coverage exceeds the $(1 - \alpha) = 90\%$ target in every bin, including the lowest-reliability bin $[0.0, 0.2)$.
- Interval width increases as reliability decreases, confirming that the RUI inflation module widens intervals when quality is poor—the mechanism behind Theorem C.2 (monotone inflation).
- The lowest-reliability bin has the widest intervals and the *highest* coverage, verifying that the framework is conservative precisely when uncertainty is greatest.

g.3 Coverage–Width Trade-off

Figure g.1 plots empirical coverage against average interval width across reliability bins, illustrating the monotone inflation property.

g.4 CQR Group Coverage Visualisation

Figure g.2 visualises the per-group coverage alongside the global target, confirming uniform coverage.




Figure g.1: Coverage vs. interval width by reliability bin. Lower reliability $\Rightarrow$ wider intervals $\Rightarrow$ higher coverage. Source: `fig_coverage_width.png`.




Figure g.2: CQR group coverage across reliability bins. All bins exceed the 90% target (dashed line). Source: `fig_coverage_width_tradeoff.png`.

g.5 Audit Protocol

The reliability-group audit procedure:

1. Run the priority-2 artifact generator and produce `cqr_group_coverage.csv`.
2. Verify that no bin is flagged in the output CSV.
3. Cross-check aggregate PICP against `dc3s_main_table.csv`.
4. Record pass/fail in the verification log.

Appendix h

Benchmark Specifications

This appendix records the benchmark-facing specifications that make the dissertation replayable and auditable. ORIUS-Bench is universal at the schema level: every episode tracks candidate action, repaired action, true-state violation, observed-state satisfaction, certificate state, fallback or intervention markers, and runtime cost. The current deepest empirical instantiation remains the battery witness row, so the detailed fault-family table below is still battery-grounded. That grounding is deliberate: it shows the strongest defended fault surface while the benchmark contract itself remains domain-neutral.

Any remaining battery-shaped compatibility fields in the codebase should therefore be read as legacy support surfaces, not as the dissertation’s definition of the benchmark. The canonical benchmark language is the dual-trajectory, adapter-driven contract described in Chapter 6.

h.1 Canonical fault families

Table [h.1](#) defines each fault family, its configurable parameters, the tested range, and the battery-domain interpretation.

Table h.1: Witness-domain fault families used to instantiate the universal ORIUS-Bench telemetry-degradation contract.

Fault	Parameters	Range tested	Battery interpretation
dropout	drop_rate d	$d \in \{0.05, 0.10, 0.15, 0.20, 0.25, 0.30\}$	Packet loss on the BMS→EMS telemetry link. Dropped readings are replaced by the last received value (zero-order hold).
delay jitter	base delay τ_b (s), jitter σ_τ (s)	$\tau_b \in \{1, 3, 5, 10\}$	Variable network latency between BMS and the dispatch engine. Delayed packets arrive out of order and are consumed at their arrival time.
spike	magnitude μ_s (multiplier)	$\mu_s \in \{2, 3, 4, 5\}$	Transient sensor glitch producing a reading μ_s times the true value. Models ADC saturation or electromagnetic interference on current/voltage sensors.
stale	hold duration h (steps)	$h \in \{2, 4, 6, 8, 12\}$	Sensor freeze: the telemetry value is held constant for h consecutive steps. Models BMS firmware hang or communication bus lock-up.
drift	drift rate β (per step)	$\beta \in \{0.01, 0.02, 0.03, 0.05\}$	Slow, linear bias accumulating as $\xi_t = \beta t$. Models sensor calibration drift (e.g. current-sensor offset) or gradual SOH-induced SOC bias.
reorder	buffer size B (steps)	$B \in \{2, 4, 6, 8\}$	Out-of-order packet delivery: readings within a sliding window of B steps are randomly permuted. Models multi-path routing or priority-queue inversion on industrial networks.

h.2 Combined profile: `drift_combo`

The `drift_combo` profile activates the configured fault families simultaneously at their mid-range parameter values with a configurable dropout rate d . This is the default stress profile used in the main results of the witness-domain chapter and in the bounded controller-comparison surfaces that accompany the ORIUS publication artifacts.

Table h.2: `drift_combo` default parameter values.

Fault	Parameter	Default value
dropout	d	0.20
delay_jitter	τ_b	3 s
spike	μ_s	$3\times$
stale	h	4 steps
drift	β	0.02 /step
reorder	B	4 steps

h.3 Replay and fault-injection implementation

Faults are injected by the CPSBench runner (`src/orius/cpsbench_iot/runner.py`) between the plant’s true-state output and the controller’s observation input. The injection pipeline:

1. **Plant emits** clean telemetry z_t .
2. **Fault layer** transforms $z_t \rightarrow \tilde{z}_t$ according to the active fault families and their parameters.
3. **Controller receives** degraded telemetry $\tilde{z}_t$ as its observation o_t .
4. **DC3S OQE** computes w_t from $\tilde{z}_t$ (detecting the degradation).

Fault parameters are set in `configs/dc3s.yaml` under the `fault_profile` key. Each evaluation run logs the active profile and parameter snapshot in the certificate metadata for reproducibility.

h.4 Severity mapping

Table [h.3](#) maps each fault family to its primary impact on DC3S pipeline stages.

Table h.3: Fault-to-pipeline severity mapping.

Fault	OQE w_t	RUI m_t	Shield IR	Cert. validity
dropout	↓↓	↑↑	↑	↓
delay__jitter	↓	↑	↑	↓
spike	↓↓	↑↑	↑↑	↓↓
stale	↓	↑	↑	↓
drift	↓↓	↑↑	↑	↓↓
reorder	↓	↑	↑	—

↓ = mild decrease; ↓↓ = strong decrease; ↑ = mild increase; ↑↑ = strong increase; — = negligible effect.

Appendix i

Battery Adapter Interface Specification

This appendix specifies the `BatteryDomainAdapter` interface that bridges the domain-agnostic DC3S pipeline to the battery energy storage domain. Every domain that instantiates ORIUS must implement these six methods; the battery adapter serves as the reference implementation.

i.1 Interface Overview

The adapter exposes six methods that map one-to-one onto the DC3S pipeline stages: Observe, Qualify, Inflate, Tighten, Repair, Certify. Table [i.1](#) summarises each method.

Table i.1: `BatteryDomainAdapter` interface methods.

Method	Input	Output	Battery interpretation
<code>ingest_telemetry</code>	Raw BMS packet $\tilde{z}_t$	Parsed state vector o_t	Extracts SOC, voltage, current, temperature from the degraded telemetry stream.
<code>compute_oqe</code>	o_t , history buffer	Reliability score $w_t \in [0, 1]$	Aggregates freshness, consistency, and drift evidence into a scalar quality score.
<code>build_uncertainty_set</code>	o_t , w_t , conformal quantile q_t	Uncertainty set U_t	Constructs $[\hat{x}_t - m_t, \hat{x}_t + m_t]$ with $m_t = q_t/(w_t + \epsilon)$.
<code>tighten_action_set</code>	U_t , plant constraints	Safe action set $\mathcal{A}_t$	Shrinks the power/ramp feasible set by the margin m_t on each SOC bound.
<code>repair_action</code>	Candidate a_t^* , $\mathcal{A}_t$	Safe action a_t^{safe}	Projects a_t^* onto $\mathcal{A}_t$ via constrained QP; returns a_t^* if already feasible.
<code>emit_certificate</code>	w_t , m_t , a_t^{safe} , prev. hash	Certificate record	Emits a hash-chained, per-step certificate with all DC3S fields for audit.

i.2 Method Contracts

Each method satisfies the following contracts:

`ingest_telemetry`

Pre: raw packet is a byte-stream or structured dict. **Post:** $o_t \in \mathbb{R}^n$ with all NaN/missing values imputed by zero-order hold. **Latency:** < 0.1 ms.

`compute_oqe`

Pre: o_t is well-formed; history buffer has ≥ 1 entry. **Post:** $w_t \in [w_{\min}, 1.0]$; monotone in data quality. **Latency:** < 0.5 ms.

`build_uncertainty_set`

Pre: $w_t > 0$, $q_t > 0$. **Post:** U_t is a non-empty interval; $m_t \leq m_{\max}$. **Latency:** < 0.1 ms.

`tighten_action_set`

Pre: U_t is non-empty. **Post:** $\mathcal{A}_t \subseteq \mathcal{U}$; $\mathcal{A}_t \neq \emptyset$ when $\phi(x_t) = 1$ (Assumption A3). **Latency:** < 0.2 ms.

`repair_action`


Pre: $\mathcal{A}_t \neq \emptyset$. **Post:** $a_t^{\text{safe}} \in \mathcal{A}_t$; $\|a_t^{\text{safe}} - a_t^*\|$ is minimised (projection). **Latency:** < 1.0 ms.

`emit_certificate`

Pre: all upstream fields computed. **Post:** certificate contains fields `{step, w_t, margin, action, intervened, hash, prev_hash}`. **Latency:** < 0.3 ms.

i.3 Pipeline Diagram

The six methods form a strict sequential pipeline within each dispatch step:



i.4 How to Implement a New Domain Adapter

To extend ORIUS to a new CPS domain (e.g. HVAC, water treatment):

1. **Subclass** `DomainAdapter` and implement the six methods above with domain-specific parsing, quality scoring, and constraint definitions.

2. **Define the safety predicate** $\phi(x)$ for the new domain (e.g. temperature bounds for HVAC, chemical concentration for water).
3. **Specify plant dynamics** $f(x, a)$ so that `tighten_action_set` can compute forward-reachable states.
4. **Calibrate the conformal module** on domain-specific forecast residuals to obtain q_t .
5. **Configure** `configs/dc3s.yaml` with domain-appropriate thresholds ($w_{\min}$, $m_{\max}$, α).
6. **Register** the adapter in the CPSBench runner so that the evaluation harness discovers it automatically.

The battery adapter in `src/orius/dc3s/` serves as the canonical reference. All six methods, their unit tests, and integration tests are documented in the code-level API reference (plan document `orius-plan/04-core-apis.md`).

Appendix j

Simulated 15-Reviewer Audit Protocol

This appendix records the canonical simulated review program for the live ORIUS defense package. The review is executed in three waves of five reviewer lenses each. Its job is not to declare theorems “proved” by committee vote; its job is to pressure-test the defended surface against logic drift, code/prose mismatch, scope inflation, and artifact drift.

The defended package is the active Battery + AV + Healthcare lane plus the promotion-gated theorem split:

- **flagship theorem:** T1, T2, T3a, T4, T5, T6, T7, T8, T9, T10, T11, T11_{Byz}, T_{stale}, T_{trajectory_PAC},
- **flagship corollary:** T3b,
- **flagship lemma:** L1–L4,
- **scoped flagship theorem:** T_{minimax}, T_{sensor}.

All findings must terminate in exactly one disposition:

- **fixed in code/tests:** the executable surface was corrected,
- **narrowed in prose/registers:** the claim was reduced to match repo truth, or
- **left open as bounded future work:** the claim remains outside the defended surface.

j.1 Reviewer-Role Matrix

Wave	Reviewer lane	Primary attack	Canonical point	pressure
1	Logic / mathematics	hidden assumptions, circular dependence, false converse steps	T3, T9, T10, extension-law proofs	
1	Statistics / calibration	coverage semantics, grouped-coverage drift, same-data evaluation	conformal, Mondrian, conditional-coverage prose	
1	Control / safety filtering	fallback semantics, repair soundness, per-step runtime guarantees	shield, AV fallback, T7/T8	
1	Systems / runtime	build integrity, runtime lifecycle, evidence discipline	Makefile, review build path, artifact generators	
1	Red team / durability	vacuous success, sanitization, silent theorem-helper drift	empty bins, input clamping, half-life surface split	
2	ML / OQE skepticism	heuristic reliability grounding, law-vs-code mismatch	OQE, L1–L4, README and summary prose	
2	Information theory	Fano/Le Cam misuse, proxy bridges presented as proofs	L1, L2, L3, L4, T_minimax, T_sensor	
2	Formal methods posture	proof gate vs traceability gate, audited-vs-verified drift	integrated theorem gate, Appendix M	
2	Scope / overclaim audit	defended-domain inflation, readiness inflation, publication rhetoric	abstract, universality chapters, review package	
2	Cross-theorem consistency	theorem-law dependency drift, stale inventory surfaces	theorem register, Appendix C, active audit, register CSV	
3	Battery realism	witness-row assumptions, fallback/slack boundaries	battery theorem ladder	

Wave	Reviewer lane	Primary attack	Canonical pressure point
3	AV realism	bounded longitudinal contract vs full-AV wording	AV adapter, TTC contract, closure prose
3	Healthcare / statistical gaps	evidence tier, calibration posture, domain-status honesty	healthcare row and related tables
3	Reproducibility / artifacts	hash-locking, generated-asset drift, build determinism	publication artifacts and generators
3	Long-horizon durability	what survives repo drift in five or ten years	theorem-local assumptions, future-work boundaries

j.2 Master Vulnerability Register

ID	Severity	Finding	Disposition	Resolution
V1	critical	review-compile and paper-compile failed open on LaTeX errors	fixed code/tests	in removed fail-open shell prefixes in the canonical Makefile build path
V2	critical	Appendix C claimed a Ville/-supermartingale trajectory proof not matched by code	fixed code/tests	in narrowed T_trajectory_PAC to the executable Bonferroni/union-bound certificate in code, tests, and Appendix C

ID	Severity	Finding	Disposition	Resolution
V3	critical	The integrated theorem gate presented a traceability check as a proof-verification gate	fixed in code/tests	reabeled the gate as a traceability-release gate and preserved proof rigor in the active audit instead
V4	critical	L1/L2 extension laws were stronger than the repo can defend today	narrowed in prose/registers	Appendix C, Appendix M, the law helpers, and the active audit now mark them as stylized/proxy surfaces, not closed converses
V5	high	T_minimax and T_sensor inherited the open L1/L2 bridge but were written as defended converses	fixed in code/tests	theorem helpers now mark the converse gap as open; Appendix C and Appendix M preserve the narrower posture
V6	high	The theorem-law dependency graph and active inventory were stale relative to Appendix C	fixed in code/tests	the theorem register, theorem-surface validator, and active audit inventory are synchronized to the live appendix surface
V7	high	Review-package language overstated publication/thesis readiness while the repo still carried open scope gaps	narrowed in prose/registers	the review appendix and generated dossier now describe bounded readiness and explicit dispositions instead of blanket readiness claims

ID	Severity	Finding	Disposition	Resolution
V8	high	Domain-evidence rhetoric still drifted beyond the governed three-domain posture	narrowed in prose/registers	the defense package now treats Battery as witness row and AV plus Healthcare as the only promoted bounded rows in the live program
V9	medium	T8/T9/T10/T11 prose was stronger than the executable witnesses	narrowed in prose/registers	Appendix M now reports the promoted T6/T7 rows separately from the remaining scoped / hole / partial surfaces
V10	medium	Several extension helpers were algebraically useful but semantically overstated	fixed in code/tests	helper docstrings, summaries, and register metadata now distinguish executable witness, stylized proxy, and defended theorem surfaces

j.3 Response / Concession Matrix

Disposition	What it means	Examples in this pass
Fixed in code/tests	Runtime semantics, theorem helper, or build path changed and regression tests were updated	fail-closed build; traceability-gate relabel; Bonferroni PAC alignment; parent-law drift removal
Narrowed in prose/registers	Manuscript/register language was reduced to match current executable and artifact truth	Appendix C law/theorem wording; Appendix M audit posture; readiness language in the defense package
Bounded future work	Claim remains visible only as an explicit open problem, not a defended contribution	independent converse for L1/L2; sharper minimax lower side; stronger martingale trajectory certificate

j.4 Final Disposition Summary

The simulated 15-reviewer audit does not certify that every theorem is fully rigorous. It certifies something narrower and more useful:

1. the canonical defense package is built on the repo-truthful path (`make review-compile`),
2. proof rigor is tracked by the active theorem audit rather than by a traceability gate pretending to be a proof checker,
3. extension laws and depth theorems are included only with their current bounded posture, and
4. any row that still lacks code, tests, register alignment, and proof alignment is explicitly narrowed or left open.

That is the defended posture of the live repository after hardening: a bounded, auditable theorem-and-artifact surface rather than a rhetorically closed one.

Appendix k

Audited Theorems and Gap Register

This appendix is the human-readable companion to the machine-generated active theorem audit in `reports/publication/active_theorem_audit.*`. It does *not* treat the live theorem inventory as one flat ladder. Instead it mirrors the promotion-gated split emitted by `theorem_promotion_matrix.*`:

- **Flagship theorem:** T1, T2, T3a, T4, T5, T6, T7, T8, T9, T10, T11, T11_{Byz}, T_{stale}, and T_{trajectory_PAC}.
- **Flagship corollary:** T3b.
- **Flagship lemma:** L1–L4.
- **Scoped flagship theorem:** T_{minimax} and T_{sensor}.

This appendix audits the full live surface. Historical drafts remain visible only as archived text and do not inflate the promoted theorem count.

The certificate half-life remains a corollary, and safe-budget monotonicity remains a proposition. The tracked integrated theorem surface register still audits the older hard-gated traceability surface, but it is a *traceability gate*, not a proof-soundness gate.

k.1 Theorem Statements

k.1.1 Theorem 1: OASG Existence

Statement. There exist scenarios where an action appears safe in observed state but is unsafe in true physical state. Formally: $\exists$ episodes with degraded telemetry such that $\text{observed-safe}(a \mid x_{\text{obs}}) \neq \text{true-safe}(a \mid x_{\text{true}})$.

Assumptions. A2 (bounded telemetry error), A4 (known dynamics). Fault injection produces non-trivial $x_{\text{obs}} \neq x_{\text{true}}$ separation.

Proof status. Structural: code explicitly separates observed and true state; empirical: violation rates under dropout/drift exceed nominal.

Code anchor. `src/orius/cpsbench_iot/: scenarios.py` (fault generation, x_{obs} , x_{true} , `event_log`); `plant.py` (unclamped truth SOC in `BatteryPlant.step()`); `runner.py` (true-state violation computation).

k.1.2 Theorem 2: One-Step Safety Preservation

Statement. If the repaired action lies inside the tightened safe set and the current state lies inside the inflated certificate, then the next battery state remains safe.

Assumptions. A1–A5, A7.

Proof status. Proved as a conditional one-step theorem plus deterministic runtime guarantee checks. Coverage frequency is supplied separately and is not bundled into T2.

Code anchor. `src/orius/dc3s/guarantee_checks.py` and `src/orius/dc3s/shield.py`.

k.1.3 Theorem 3a: ORIOUS Core Envelope Derivation

Statement. The active defended T3a surface is the per-step envelope

$$\mathbb{P}[Z_t = 1 \mid \mathcal{H}_t] \leq \alpha(1 - w_t),$$

under the explicit battery risk-budget contract and the narrowed reliability-score interpretation.

Assumptions. A1, A2, A4, A5, A6, A7, A9, plus the theorem-local battery calibration bridge recorded in the active theorem audit.

Proof status. Paper-rigorous flagship theorem with explicit runtime contract reporting. The code now emits a theorem-contract summary that checks the executable envelope algebra and keeps the calibration bridge explicit as a declared theorem-local contract rather than hiding it inside generic conformal language.

Code anchor. `src/orius/universal_theory/risk_bounds.py` (`build_t3a_contract_summary()`), plus the backward-compatible wrapper `src/orius/dc3s/coverage_theorem.py`.

k.1.4 Corollary 3b: ORIOUS Core Aggregation Corollary

Statement. Aggregating the predictable per-step budget from T3a gives the episode envelope

$$\mathbb{E}[V] \leq \alpha(1 - \bar{w})T.$$

Assumptions. The assumptions of T3a, plus linearity of expectation.

Proof status. Supporting defended corollary. Once T3a is fixed, the aggregation is algebraic bookkeeping rather than a new theorem burden.

Code anchor. `src/orius/universal_theory/risk_bounds.py` (`compute_episode_risk_bound()`).

k.1.5 Theorem 4: No Free Safety

Statement. Within the fixed-margin, quality-ignorant controller class, there exists an admissible degraded-observation sequence that produces an observation-action safety gap and hence a true-state violation.

Assumptions. A3 (feasible safe repair exists), A8 (admissible fallback policy). Quality-ignorant baselines lack reliability-aware inflation.

Proof status. Constructive witness theorem on the battery reference surface, plus empirical ablation support.

Code anchor. Supporting evidence in `baselines.py`, `scenarios.py`, and `runner.py`.

k.1.6 Definition 5: Certificate Validity Horizon

Statement. T5 is kept as the constructive maximal forward-tube containment definition: the validity horizon is the largest integer horizon whose forward tube remains inside the battery SOC bounds.

Assumptions. A1, A2, A4–A7.

Proof status. Explicit definition, not part of the defended theorem count. The helper computes the maximal containment horizon directly and no probabilistic future-step theorem is claimed at this surface.

Code anchor. `src/orius/dc3s/temporal_theorems.py` (`forward_tube()`, `certificate_validity_horizon()`); registered as T5 in THEOREM_REGISTER (`src/orius/dc3s/theoretical_guarantees.py`). Runtime expiry logic consumes this quantity, but the active theorem audit does not count T5 as a defended theorem.

k.1.7 Theorem 6: Certificate Expiration Bound

Statement. The battery certificate horizon satisfies the canonical delta-aware lower bound

$$\tau_t \geq \left\lfloor \frac{\delta_{\text{bnd},t}^2}{2\sigma_d^2 \log(2/\delta)} \right\rfloor.$$

Assumptions. A4–A7.

Proof status. Flagship defended closed-form theorem with machine-checkable helper linkage. The theorem-facing API requires `delta` explicitly, records the first-passage side conditions, and marks legacy no-delta semantics as non-defended.

Code anchor. `src/orius/universal_theory/battery_instantiation.py` (`certificate_expiration_bound()`), which now emits theorem metadata and a fail-closed T6 contract summary.

k.1.8 Theorem 7: Feasible Fallback Existence

Statement. The defended T7 surface is piecewise: zero dispatch preserves safety on the interior, a boundary-aware safe-landing action preserves safety near the boundary, and the runtime fails closed when neither case can be certified.

Assumptions. A1, A4, A8.

Proof status. Flagship defended battery fallback theorem with an explicit piecewise executable witness. The promoted surface is battery-only and does not claim cross-domain transfer on its own.

Code anchor. `src/orius/dc3s/temporal_theorems.py` (`zero_dispatch_fallback()`, `certify_fallback_existence()`).

k.1.9 Theorem 8: Graceful Degradation Dominance

Statement. The defended T8 surface is the paired sequence-level comparison: for graceful and uncontrolled violation sequences produced under the same admissible fault trace,

$$z_t^{\text{graceful}} \leq z_t^{\text{uncontrolled}} \quad \forall t \implies V^{\text{graceful}} \leq V^{\text{uncontrolled}}.$$

Strict improvement is claimed only when the graceful sequence is strictly smaller on at least one step.

Assumptions. A1–A5, A8.

Proof status. Supporting defended comparison theorem at the sequence-comparison level only. The stronger absorbing-landing interpretation remains supporting evidence rather than a promoted theorem.

Code anchor. `src/orius/universal_theory/battery_instantiation.py` (`evaluate_graceful_degradation_dominance()`), with `safe_landing` kept as supporting implementation evidence only.

k.1.10 Theorem 9: Universal Impossibility

Statement. For any quality-ignorant controller π and any φ -mixing fault process with mean degradation rate $d > 0$, $\mathbb{E}[V_T(\pi)] \geq cdT$ for some constant $c > 0$.

Assumptions. A4 (boundary reachability), A6' (φ -mixing fault process).

Proof status. Structural impossibility theorem under explicit φ -mixing and boundary-reachability assumptions; battery T4 is the reference witness for the per-window violation probability and the witness constant remains theorem-local.

Code anchor. No dedicated module; the battery `cpsbench_iot/runner.py` provides the reference witness; the proof is summarized in Section [k.1.10](#).

k.1.11 Theorem 10: Boundary-Indistinguishability Lower Bound

Statement. Under the explicit boundary-mass and indistinguishability hypotheses of T10, every controller π satisfies the scoped lower bound $\mathbb{E}[V_T(\pi)] \geq \frac{1}{4}\alpha(1 - \bar{w})T$.

Assumptions. A1, A2, A5' (calibration regularity), conformal coverage condition.

Proof status. Narrowed stylized lower-bound surface under extra calibration/indistinguishability assumptions. The active audit records that the one-sided Le Cam step remains open, so the code witness is treated as a scoped lower-bound calculator rather than a sharp frontier theorem.

Code anchor. `src/orius/universal/reliability_safety_frontier.py` (`compute_frontier`, `is_dc3s_optimal`); 25 unit tests in `tests/test_reliability_frontier.py`.

k.1.12 Theorem 11: Typed Structural Transfer

Statement. Any CPS domain instantiation satisfying the structural transfer obligations of T11—coverage of the observation-consistent state set, soundness of the tightened safe-action set, repair membership, and safe fallback admissibility—inherits the T2–T3 safety-guarantee pattern without new battery-specific algebra.

Assumptions. The universal adapter contract and the four structural obligations stated in Section [k.1.12](#).

Proof status. Forward one-step transfer theorem. The runtime now builds the four T11 obligation checks directly from typed artifacts, while the converse is tracked separately as a structural failure-mode proposition rather than as part of the defended T11 statement.

Code anchor. `src/orius/universal_theory/contracts.py` and `src/orius/universal_theory/kernel.py`, with a compact reference harness in `src/orius/universal/contract.py`; focused tests in `tests/test_universal_contract.py`,

`tests/test_unification.py`, and `tests/test_reliability_frontier.py`, plus `tests/test_universal_theory.py`.

k.2 Gap Audit Table

Table [k.1](#) summarises the historical active T1–T11 audit surface. The publication-ready defended register is now the stricter promotion matrix, where every active theorem-like row has proof, code, tests, artifacts, hashes, and a claim boundary.

k.3 Remaining Claim-Boundary Gaps

The audit above separates proof completeness from promotion completeness. The remaining gaps that still bound the dissertation’s final claim are therefore explicit:

1. **Autonomous vehicles.** The current AV row is no longer claimed as full autonomy closure. It is promoted only as a bounded nuPlan replay row under TTC and predictive-entry-barrier semantics. Remaining AV work is closed-loop road deployment evidence, not a contradiction of T1–T11.
2. **Future mobility and flight-style rows.** Earlier drafts named additional exploratory rows, but the current monograph does not promote them. They remain planning surfaces until they supply defended replay, runtime, governance, and post-repair evidence under the same gate as the active Battery + AV + Healthcare lane.

These are the remaining thesis-level gaps after the monograph theorem pass. They do not weaken the battery proof witness; they bound how far the transfer claim is currently allowed to extend.

k.4 Locked Empirical Evidence

Battery theorem verification relies on a mixture of code anchors, tests, and locked outputs in `reports/publication/`:

- `dc3s_main_table_ci.csv`: violation rates, PICP, scenario breakdown
- `reliability_group_coverage_phase3.csv`: coverage by reliability bin
- `blackout_study.csv`: certificate-safe-hold evidence for the temporal extension
- `graceful_degradation_trace.csv`: behavioral evidence for the landing controller
- `claim_evidence_matrix.csv`: claim-to-evidence linkage

The theorem-to-evidence map is maintained in `orius-plan/theorem_to_evidence_map.md`.

k.5 Integrated Release-Gate Status

The flagship T1–T8 ladder above remains the governance-critical proof spine, and the broader defended theorem program now extends through T11. The generated summary below reports the older hard-gated

Table k.1: Historical T1–T11 theorem audit summary before promotion-gate closure

Theorem	Proof status	Empirical validation	Code matches statement	Notes
T1 OASG Existence	Structural	Yes	Yes	Observed-vs-true split is explicit in benchmark code and locked evidence
T2 One-Step Safety Preservation	Yes (conditional one-step)	Yes	Yes	Shield + guarantee checks + unit tests anchor the theorem; coverage frequency handled separately
T3a ORIUS Core Envelope Derivation	Paper-rigorous	Yes	Yes	Runtime theorem-contract summary now separates executable envelope algebra from the declared calibration bridge
T3b ORIUS Core Aggregation Corollary	Corollary only	Yes	Yes	Supporting aggregation corollary derived from T3a by linearity of expectation
T4 No Free Safety	Constructive witness	Yes	Partial	Battery proof gives an explicit degraded-observation sequence; broader baseline comparison remains empirical
T5 Certificate Validity Horizon Definition	Explicit definition	Yes	Yes	Maximal forward-tube containment helper is constructive and excluded from defended theorem counts
T6 Certificate Expiration Bound	V2-linked flagship theorem	Yes	Yes	Canonical helper is delta-aware, closed-form, and records explicit first-passage side conditions
T7 Feasible Fallback Existence	Piecewise hold-or-safe-landing theorem	Yes	Yes	Runtime certifies interior safe hold, boundary recovery by safe landing, and fail-closed infeasibility
T8 Graceful Degradation Dominance	Supporting sequence comparison	Partial	Yes	Defended only as a paired-sequence comparator under a shared fault trace; absorbing landing remains supporting-only
T9 Universal Impossibility	Structural / scoped witness constant	Yes	Partial	φ -mixing + Azuma windowing remain explicit assumptions; battery T4 is the reference witness; Section k.1.10
T10 Boundary-Indistinguishability Lower Bound	Narrowed stylized lower bound	Partial	Partial	Le Cam step remains one-sided; the executable witness is a scoped lower-bound calculator, not a sharp frontier theorem
T11 Typed Structural Transfer	Forward-only one-step transfer	Yes	Yes	Typed runtime artifacts now carry four-obligation summaries; the converse and mini-harness remain separate supporting surfaces

Table k.2: Audited extension laws and depth-theorem surfaces

ID	Status	Code surface	Audit posture	Notes
L1	Broken as standalone converse	<code>orius_law.py</code>	Stylized surrogate only	Linear lower envelope retained for extension calculators only
L2	Broken as theorem of live OQE	<code>orius_law.py</code> , <code>quality.py</code>	Proxy bridge only	Requires external justification of κ_d and the capacity proxy
L3	Has a hole	<code>orius_law.py</code>	Narrowed threshold calculator	Depends on L2 proxy plus the executable T3 upper envelope
L4	Has a hole	<code>orius_law.py</code>	Stylized sandwich	Executable upper side plus proxy lower side; not a closed characterization
T11_Byzantine	Plausible but informal	<code>theoretical_guarantees.py</code>	Bounded extension	Matches the trimmed-mean witness and tests
T_stale_decay	Plausible but informal	<code>theoretical_guarantees.py</code>	Bounded extension	Geometric decay is a design schedule, not a derived physical law
T_minimax	Broken as defended minimax theorem	<code>theoretical_guarantees.py</code>	Stylized lower-envelope witness	Inherits the open L1/L2 proxy bridge
T_sensor_converse	Broken as universal converse	<code>theoretical_guarantees.py</code>	Stylized inverse threshold	Necessity only holds inside the L2/L3 proxy model
T_trajectory_PAC	Plausible but informal	<code>risk_bounds.py</code>	Defended as Bonferroni certificate	No separate Ville-strengthened theorem is claimed

release status for the 10 source-backed theorem claims in the precursor plus reference ladder together with the 8 Appendix C theorem restatements. The larger T1–T11 program is kept aligned by the theorem-surface validator and the appendix proof crosswalk rather than by that legacy 18-row gate alone.

Table k.3: Integrated 18-row hard-gated proof-surface traceability summary. The unique surface set contains the two supporting Chapter 4 statements and the T1–T8 surface with active defense tiers imported from the theorem audit; Appendix C restatements are checked against that unique surface.

Class	Passed	Total
Unique theorem claims	10	10
Appendix theorem restatements	8	8
All hard-gated theorem rows	18	18

Table k.4: Integrated theorem traceability-gate matrix.

Key	Class	Title	Source	Code	Artifacts	Mapping	Status
S1	unique_theorem	Existence of the illusion under dropout	yes	yes	yes	yes	pass
S2	unique_theorem	DC3S Feasibility Guarantee	yes	yes	yes	yes	pass
T1	unique_theorem	OASG Existence	yes	yes	yes	yes	pass
T2	unique_theorem	Safety Preservation	yes	yes	yes	yes	pass
T3	unique_theorem	ORIOUS Core Bound	yes	yes	yes	yes	pass
T4	unique_theorem	Observation Necessity / No Free Safety	yes	yes	yes	yes	pass
T5	unique_theorem	Certificate validity horizon	yes	yes	yes	yes	pass
T6	unique_theorem	Certificate expiration bound	yes	yes	yes	yes	pass
T7	unique_theorem	Feasible fallback existence	yes	yes	yes	yes	pass
T8	unique_theorem	Graceful degradation dominance	yes	yes	yes	yes	pass
C.1	appendix_restatement	Client Battery-Domain OASG Existence	yes	yes	yes	yes	pass
C.2	appendix_restatement	Client One-Step Safety Preservation	yes	yes	yes	yes	pass

Continued on next page

Table k.4 — continued

Key	Class	Title	Source	Code	Artifacts	Mapping	Status	
C.3	appendix_restatement	Client Battery Core Bound	ORIOUS	yes	yes	yes	yes	pass
C.4	appendix_restatement	Client Observation Necessity / No Free Safety		yes	yes	yes	yes	pass
C.5	appendix_restatement	Client Certificate Validity Horizon		yes	yes	yes	yes	pass
C.6	appendix_restatement	Client Certificate Expiration Bound		yes	yes	yes	yes	pass
C.8	appendix_restatement	Client Feasible Fallback Existence		yes	yes	yes	yes	pass
C.9	appendix_restatement	Client Graceful Degradation Dominance		yes	yes	yes	yes	pass

Appendix 1

Locked Artifact Hash Registry

This appendix replaces the earlier draft hash registry with concrete locked entries drawn from `reports/publication/final_lock_manifest_full_thesis.json`. The goal is not to list every temporary file in the repository. The goal is to record the canonical final-package artifacts whose hashes anchor the thesis claim surface.

Table 1.1: Locked artifact hash registry from the final full-thesis package.

Artifact path	SHA-256 prefix	Bytes
reports/publication/final_package_ FINAL_20260317T030715Z/48h_trace.csv	7002cbaa77e8dfe0	14971
reports/publication/final_package_ FINAL_20260317T030715Z/active_probing_ spoofing_detection.csv	a8d12550ed3ae633	105
reports/publication/final_package_ FINAL_20260317T030715Z/aging_ calibration_outputs.csv	f5b89bafecac6d3f	255
reports/publication/final_package_ FINAL_20260317T030715Z/appendices_hq_ manifest.json	73e8e9c4ea0593d0	1723
reports/publication/final_package_ FINAL_20260317T030715Z/baseline_lock_ manifest.json	bcfa0eee34e71792	3053
reports/publication/final_package_ FINAL_20260317T030715Z/battery_ leaderboard.csv	96fbd6dfcf857cb1	647

Continued on next page

Table l.1 — continued

Artifact path	SHA-256 prefix	Bytes
reports/publication/final_package_ FINAL_20260317T030715Z/blackout_half_ life.csv	cd35353971342fdf	122
reports/publication/final_package_ FINAL_20260317T030715Z/claim_evidence_ matrix.csv	e429255befdc88e2	1891
reports/publication/final_package_ FINAL_20260317T030715Z/dc3s_latency_ summary.csv	cc1e470c6df2d5b9	336
reports/publication/final_package_ FINAL_20260317T030715Z/fault_ performance_table.csv	ba5652d3b071a742	4275
reports/publication/final_package_ FINAL_20260317T030715Z/fleet_ composition_two_battery.csv	6dd383794d60f42d	6838
reports/publication/final_package_ FINAL_20260317T030715Z/graceful_ degradation_trace.csv	dad786d86aee0f58	6352
reports/publication/final_package_ FINAL_20260317T030715Z/hyperparameter_ surfaces.csv	b1be66a9a03710e4	655
reports/publication/final_package_ FINAL_20260317T030715Z/orius-plan/ 10-baseline-lock-and-runid-policy.md	9e5c2c3fe41f4e96	534
reports/publication/final_package_ FINAL_20260317T030715Z/orius-plan/ 11-part1to4-gap-closure.md	28714d964c28fa9f	1233
reports/publication/final_package_ FINAL_20260317T030715Z/orius-plan/ 12-part5-hardening-lock.md	b0b55b82c4de4ee2	853
reports/publication/final_package_ FINAL_20260317T030715Z/orius-plan/ 13-part6-extension-evidence.md	e85963b5f28afeb0	1127

Continued on next page

Table l.1 — continued

Artifact path	SHA-256 prefix	Bytes
reports/publication/final_package_ FINAL_20260317T030715Z/orius-plan/ 14-appendices-hq-buildout.md	3f537daaf1b838c4	672
reports/publication/final_package_ FINAL_20260317T030715Z/orius-plan/ assumption_register.md	d11d439c0475e305	4157
reports/publication/final_package_ FINAL_20260317T030715Z/orius-plan/ theorem_to_evidence_map.md	4c890c0003c634fa	4341
reports/publication/final_package_ FINAL_20260317T030715Z/paper.pdf	21a273a8aa25debd	3796137
reports/publication/final_package_ FINAL_20260317T030715Z/part5_hardening_ lock_manifest.json	44b918fed0c5166d	1916
reports/publication/final_package_ FINAL_20260317T030715Z/part6_extension_ lock_manifest.json	03e64f759a95e277	2171
reports/publication/final_package_ FINAL_20260317T030715Z/priority2_ artifacts_manifest.json	ec9d073108b836ee	692
reports/publication/final_package_ FINAL_20260317T030715Z/priority3_ artifacts_manifest.json	c994d26b52c7dd34	465
reports/publication/final_package_ FINAL_20260317T030715Z/release_manifest. json	e3bc745701ebc6c4	9486
reports/publication/final_package_ FINAL_20260317T030715Z/run_id_policy. json	ba69366033e80aa4	320

The registry above is intentionally concrete. It gives the defense package a stable hash anchor for the highest-value battery artifacts rather than a promise that such a registry will be added later.

Appendix m

Claim-to-Evidence and Coverage Registers

This appendix consolidates the governance registers that keep defended ORIUS claims attached to code, scripts, and artifacts. In the current defense posture the machine-readable authorities are the active theorem audit (`reports/publication/active_theorem_audit.*`) and strict defended-core outputs (`reports/publication/defended_theorem_core.*`). The tables below are the readable companion layer used in the monograph and review dossier.

m.1 Primary Claim Matrix

Table m.1: Primary claim-to-evidence matrix from the locked publication register.

Claim ID	Claim text	Evidence path	Evidence type	Status
central_claim	deterministic unsafe vs. dc3s safe	reports/publication/dc3s_main_table_ci.csv	csv row	locked
theorem_t1	oasg existence	src/orius/cpsbench_iot/scenarios.py	code path	flagship defended
theorem_t2	one-step safety preservation	src/orius/dc3s/guarantee_checks.py	code path	flagship defended
theorem_t3a	per-step ORIUS core envelope derivation	src/orius/universal_theory/risk_bounds.py	code path	flagship defended
theorem_t3b	episode ORIUS core aggregation corollary	src/orius/universal_theory/risk_bounds.py	code path	supporting defended
theorem_t4	no free safety	src/orius/cpsbench_iot/baselines.py	code path	flagship defended
theorem_t5	certificate validity horizon definition	src/orius/universal_theory/battery_instantiation.py	code path	draft / non-defended

Continued on next page

Table m.1 — continued

Claim ID	Claim text		Evidence path	Evidence type	Status
theorem_t6	certificate	expiration	src/orius/universal\ _theory/battery\ _instantiation.py	code path	flagship defended
theorem_t7	feasible	fallback	src/orius/dc3s/temporal\ _theorems.py	code path	flagship defended
theorem_t8	graceful	degradation	src/orius/universal\ _theory/battery\ _instantiation.py	code path	supporting defended
theorem_t11	typed structural transfer		src/orius/universal\ _theory/contracts.py	code path	flagship defended
theorem_t11_byz	Byzantine-tolerant	OQE	src/orius/dc3s/ theoretical_guarantees. py	code path	supporting defended
theorem_t_stale	stale-decay	graceful	src/orius/dc3s/ theoretical_guarantees. py	code path	supporting defended
theorem_t_minimax	tight	OASG	src/orius/dc3s/ theoretical_guarantees. py	code path	draft ex- tension
theorem_t_sensor	sensor quality	converse	src/orius/dc3s/ theoretical_guarantees. py	code path	draft ex- tension
theorem_t_pac	PAC	trajectory	src/orius/universal\ _theory/risk_bounds.py	code path	flagship defended
defended_theorem_core	registry-canonical	theo-	reports/publication/ defended_theorem_core. json	generated au- dit	canonical
active_theorem_audit	full active theorem	inventory	reports/publication/ active_theorem_audit. json	generated au- dit	canonical
priority1_fault_table	fault	performance table	reports/publication/ fault_performance_table. csv	artifact file	generated
priority1_trace	48h	operational trace	reports/publication/48h_ trace.csv	artifact file	generated
priority1_latency	latency	p99 table	reports/publication/ dc3s_latency_summary.csv	artifact file	generated
priority1_hil	software	hil summary	reports/hil/closed_loop_ nominal_48h_summary.json	artifact file	generated
priority2_surfaces	hyperparameter	surfaces	reports/publication/ hyperparameter_surfaces. csv	artifact file	generated
priority2_half_life	blackout	half-life	reports/publication/ blackout_half_life.csv	artifact file	generated
priority2_graceful	graceful	degradation	reports/publication/ graceful_degradation_ trace.csv	artifact file	generated
priority3_leaderboard	battery	leaderboard	reports/publication/ battery_leaderboard.csv	artifact file	generated

Continued on next page

Table m.1 — continued

Claim ID	Claim text	Evidence path	Evidence type	Status
priority3_fleet	fleet composition	reports/publication/ fleet_composition_two_ battery.csv	artifact file	generated
priority3_probing	bounded active probing audit	reports/probing/ spoofing_attack_results. csv	artifact file	generated
priority3_aging	aging proxy calibration outputs	reports/aging/asset_ preservation_proxy_table. csv	artifact file	generated

m.2 Theorem and Chapter Register

This register intentionally tracks the registry-canonical defended core plus the small set of visible draft extensions that remain on the active theorem audit surface. The authority for defended status is `reports/publication/defended_theorem_core.json`; this appendix is the human-readable companion. The larger integrated surface, including definitions, corollaries, and appendix restatements, remains audited in the tracked integrated theorem surface register.

Table m.2: Registry-canonical defended-core theorem and chapter evidence register for the ORIUS defense package.

Scope	Ch.	Token	Code or script	Artifact path	Status
T1	16	THM_T1	src/orius/cpsbench\ _iot/scenarios.py	reports/publication/ dc3s_main_table_ci.csv	flagship defended
T2	17	THM_T2	src/orius/dc3s/ guarantee_checks.py	reports/publication/ dc3s_main_table_ci.csv	flagship defended
T3a	18	THM_T3A	src/orius/universal\ _theory/risk_bounds. py	reports/publication/ reliability_group_ coverage_phase3.csv	flagship defended
T3b	18	COR_T3B	src/orius/universal\ _theory/risk_bounds. py	reports/publication/ reliability_group_ coverage_phase3.csv	supporting defended
T4	19	THM_T4	src/orius/cpsbench\ _iot/baselines.py	reports/publication/ fault_performance_table. csv	flagship defended
T5	20	DEF_T5	src/orius/universal\ _theory/battery\ _instantiation.py	reports/blackout/ blackout_study.csv	draft / non- defended
T6	20	THM_T6	src/orius/universal\ _theory/battery\ _instantiation.py	reports/blackout/ blackout_study.csv	flagship defended

Continued on next page

Table m.2 — continued

Scope	Ch.	Token	Code or script	Artifact path	Status
T7	20	THM_T7	src/orius/dc3s/ temporal_theorems. py	reports/blackout/ graceful_degradation_ trace.csv	flagship defended
T8	20	THM_T8	src/orius/universal\ _theory/battery\ _instantiation.py	reports/blackout/ graceful_degradation_ trace.csv	supporting defended
T11	37	THM_T11	src/orius/universal\ _theory/contracts.py	reports/publication/ orius_transfer_ obligation_table.csv	flagship defended
T11 _{Byz}	C	THM_T11_BYZ	src/orius/dc3s/ theoretical_ guarantees.py	App. C (Proof C.18)	supporting defended
T _{stale}	C	THM_T_STALE	src/orius/dc3s/ theoretical_ guarantees.py	App. C (Proof C.19)	supporting defended
T _{minimax}	C	THM_T_MINIMAX	src/orius/dc3s/ theoretical_ guarantees.py	App. C (Proof C.20)	draft extension
T _{sensor}	C	THM_T_SENSOR	src/orius/dc3s/ theoretical_ guarantees.py	App. C (Proof C.21)	draft extension
T _{PAC}	C	THM_T_PAC	src/orius/universal\ _theory/risk_bounds. py	App. C (Proof C.22)	flagship defended
A1–A13	15	ASM_REGISTER	configs/dc3s.yaml	orius-plan/assumption_ register.md	locked
TRACE _{48H}	48H	FIG_TRACE_48H	scripts/generate_ 48h_trace.py	reports/publication/48h_ trace_final_de.csv	locked
HIL	27	FIG_HIL	scripts/generate_ hil_evidence.py	reports/hil/hil_summary. json	locked
BLACKOUT	27	FIG_BLACKOUT	scripts/run_ blackout_study.py	reports/publication/ blackout_study.csv	locked
PROBING _{appV}	27	TBL_PROBING	scripts/run_ spoofing_attack.py	reports/probing/ spoofing_attack_results. csv	locked

m.3 Chapter Coverage Register

Table m.3: Chapter coverage register from the thesis coverage audit.

Ch.	Chapter file	Primary artifact	Status
15	chapters/ch15_ assumptions_notation_ proof_discipline.tex	orius-plan/ assumption_register. md	locked

Continued on next page

Table m.3 — continued

Ch.	Chapter file	Primary artifact	Status
16	chapters/ch16_battery_ theorem_oasg_existence. tex	reports/ publication/dc3s_ main_table_ci.csv	locked
17	chapters/ch17_battery_ theorem_safety_ preservation.tex	reports/ publication/dc3s_ main_table_ci.csv	locked
18	chapters/ch18_orius_ core_bound_battery.tex	reports/ publication/ reliability_group_ coverage_phase3.csv	locked
19	chapters/ch19_no_free_ safety_battery.tex	reports/ publication/fault_ performance_table. csv	locked
20	chapters/ch20_ temporal_behavioral_ extensions.tex	reports/blackout/ blackout_study.csv	partial empiri- cal
21	chapters/ch21_fault_ performance_stress_ tests.tex	reports/fault_ benchmark/fault_ performance_table. csv	locked
22	chapters/ch22_latency_ systems_footprint.tex	reports/latency/ dc3s_latency_ summary.csv	locked
23	chapters/ch23_ hyperparameter_ surface_stability.tex	reports/ calibration/ hyperparameter_ surfaces.csv	locked
24	chapters/ch24_ conditional_coverage_ subgroups.tex	reports/ calibration/ reliability_group_ audit.csv	locked

Continued on next page

Table m.3 — continued

Ch.	Chapter file	Primary artifact	Status
25	chapters/ch25_ regional_ decomposition_real_ prices.tex	reports/ publication/per_ba_ dispatch_impact.csv	locked
26	chapters/ch26_asset_ preservation_aging_ proxy.tex	reports/aging/ asset_preservation_ proxy_table.csv	partial proxy
27	chapters/ch27_ hardware_in_loop_ validation.tex	reports/hil/hil_ summary.json	locked
28	chapters/ch28_ certificate_half_life_ blackout.tex	reports/blackout/ blackout_study.csv	locked
29	chapters/ch29_ graceful_degradation_ safe_landing.tex	reports/blackout/ graceful_ degradation_trace. csv	locked
30	chapters/ch30_orius_ bench_battery_track. tex	reports/ publication/orius_ bench_battery_ leaderboard.csv	locked
31	chapters/ch31_ compositional_safety_ battery_fleets.tex	reports/fleet/ fleet_composition_ two_battery.csv	locked
32	chapters/ch32_certos_ runtime_certificate_ lifecycle.tex	reports/certos/ certos_summary.json	locked runtime
33	chapters/ch33_what_ battery_thesis_proves. tex	reports/ publication/ defended_theorem_ core.json	locked
34	chapters/ch34_outside_ current_evidence.tex	reports/ publication/active_ theorem_audit.json	locked

Continued on next page

Table m.3 — continued

Ch.	Chapter file	Primary artifact	Status
35	chapters/ch35_ deployment_path_ verification_ discipline.tex	reports/ publication/phase3_ rebuild_battery_ evidence_manifest. json	locked
36	chapters/ch36_ conclusion.tex	reports/ publication/phase0_ lock_current_truth_ manifest.json	locked

m.4 Full Claim–Evidence Matrix

Table [m.4](#) provides the generated claim–evidence matrix from `reports/publication/claim_evidence_matrix.csv`. The theorem rows are interpreted through the strict defended-core split, so “implemented” in the raw matrix is not itself a proof-status label. Defended status lives in the active theorem audit and defended-core register; this matrix records where the underlying code and artifacts live.

Taken together, these registers are the defense-facing answer to the question “where does this claim live in code and artifacts?” They make the thesis longer, but more importantly they make it harder to bluff.

Table m.4: Claim–Evidence Matrix: all 22 thesis claims with locked artifact pointers. Status *locked* means the artifact file exists and contains real numerical data.

ID	Claim	Evidence Path	Type
central_claim	deterministic unsafe vs dc3s safe...	reports/publication/-dc3s_main_table_ci.csv	csv_row:determinis
theorem_t1	oasg existence...	src/orius/cpsbench_iot/-scenarios.py	code_path
theorem_t2	one-step safety preservation...	src/orius/dc3s/-guarantee_checks.py	code_path
theorem_t3	coverage monotonicity plus empirical core-bound audit...	src/orius/dc3s/-coverage_theorem.py	code_path
theorem_t4	no free safety...	src/orius/cpsbench_iot/-baselines.py	code_path
theorem_t5	certificate validity horizon definition...	src/orius/-universal_theory/-battery_instantiation.py	code_path
theorem_t6	certificate expiration bound...	src/orius/dc3s/-temporal_theorems.py	code_path
theorem_t7	feasible fallback existence...	src/orius/dc3s/-temporal_theorems.py	code_path
theorem_t8	graceful degradation dominance...	src/orius/dc3s/-temporal_theorems.py	code_path
priority1_fault_tabl	fault performance table...	reports/publication/-fault_performance_table.csv	artifact_file
priority1_trace	48h operational trace...	reports/publication/-48h_trace.csv	artifact_file
priority1_latency	latency p99 table...	reports/publication/-dc3s_latency_summary.csv	artifact_file
priority1_hil	software hil summary...	reports/hil/-closed_loop_nominal_48h_summary.json	artifact_file
priority2_surfaces	hyperparameter surfaces...	reports/publication/-hyperparameter_surfaces.csv	artifact_file
priority2_half_life	blackout half-life...	reports/publication/-blackout_half_life.csv	artifact_file
priority2_graceful	graceful degradation trace...	reports/publication/-graceful_degradation_trace.csv	artifact_file
priority3_leaderboar	battery leaderboard...	reports/publication/-battery_leaderboard.csv	artifact_file
priority3_fleet	fleet composition...	reports/publication/-fleet_composition_two_battery.csv	artifact_file
priority3_probing	bounded active probing audit...	reports/probing/-spoofing_attack_results.csv	artifact_file
priority3_aging	aging proxy calibration outputs...	reports/aging/-asset_preservation_proxy_table.csv	artifact_file

Appendix n

CertOS Artifact Audit

This appendix records the bounded runtime evidence promoted into Chapter 12. The runtime chapter itself carries the main story. The present appendix preserves the implementation-facing tables that a committee or reviewer may reasonably ask for once the runtime is part of the defense package.

n.1 Artifact Bundle

Table n.1: CertOS artifact bundle and its role in the thesis.

Artifact	Type	Role
reports/certos/certos_lifecycle.csv	CSV	Stepwise battery runtime trace
reports/certos/certos_summary.json	JSON	Aggregate counts by lifecycle phase
reports/certos/latency_report.csv	CSV	Runtime latency and throughput
reports/certos/audit_completeness.json	JSON	Audit-chain completeness check
reports/certos/recovery_report.md	Markdown	Recovery note after fallback
reports/certos/certos_vehicle_toy.json	JSON	Bounded non-battery lifecycle toy
reports/certos/invariant_tests.log	Log	Runtime invariant test surface
reports/certos/audit_ops.jsonl	JSONL	Append-only lifecycle operation stream

n.2 Lifecycle Counts by Status

Table n.2: Status counts derived from the locked CertOS summary artifact.

Status bucket	Count
valid	47
degraded	16
fallback	33
total steps	96
audit entries	96
runtime interventions	33

The lifecycle counts matter because they make the runtime story concrete. The battery runtime spends roughly half of the governed demo in ordinary **valid** operation, moves through a smaller degraded regime when the certificate horizon narrows, and explicitly falls back when the horizon is exhausted. That is the empirical meaning of the runtime chapter’s bounded claim.

n.3 Certificate Field Schema

The certificate engine stores only a small number of fields, but those fields are enough to reconstruct the runtime decision trail. The runtime then augments the certificate with a previous-hash link and a current certificate hash.

n.4 Audit Ledger Record Format

The ledger is intentionally austere. It is not a full event-sourcing platform. It is an append-only trail sufficient to prove that lifecycle events happened, in order, and with the expected linkage to certificate records.

n.5 Latency and Recovery Audit

The recovery note is qualitative rather than mathematical: the runtime artifact records that transitions from fallback back to valid execution were observed. This is enough for a bounded systems chapter. It is not enough to claim optimal recovery or universal recovery guarantees.

Table n.3: Certificate and runtime-state fields used by CertOS.

Field	Type	Meaning
op	string	Lifecycle operation that created the certificate
proposed_action	mapping	Action before runtime filtering
safe_action	mapping	Action released after runtime filtering
validity_horizon	integer	Remaining certified horizon
status	string	valid, degraded, expired, or revoked
prev_hash	string / null	Previous certificate hash in the ledger chain
cert_hash	string	Current certificate hash
step	integer	Runtime step index in the governed demo
fallback_active	boolean	Whether fallback action is currently active
audit_count	integer	Number of audit entries emitted so far

n.6 Bounded Second-Domain Runtime Surface

This toy artifact earns its place in the appendix because it helps defend a very narrow point: the runtime lifecycle is adapter-shaped rather than battery-field hard coded. It does not widen the thesis proof boundary.

Table n.4: Audit ledger record fields from the CertOS audit-ledger module.

Field	Type	Meaning
op	string	Lifecycle operation appended to the ledger
step	integer	Runtime step where the event occurred
cert_hash	string / null	Hash anchor for the relevant certificate
meta	mapping	Optional metadata such as recovery details

Table n.5: Latency and recovery audit from the CertOS runtime artifacts.

Metric	Value
Runtime step latency (ms)	0.59
Runtime throughput (steps/s)	1691.70
Audit completeness (%)	100.0
Recovery transitions noted	yes

Table n.6: Vehicle-shaped CertOS toy artifact.

Metric	Value
Domain	vehicle
Steps	12
Fallback steps	4
Audit entries	12
Interventions	4

Appendix o

Full CertOS Lifecycle Log

This appendix reproduces the governed battery runtime log from `reports/certos/certos_lifecycle.csv`. The main text only needs a few representative rows; the full trace belongs here so that the transition from valid to degraded to fallback execution can be inspected directly.

Table o.1: Full 96-step CertOS lifecycle log.

Step	Horizon	Status	Discharge	SOC	Fallback
0	18	valid	45.0	88.16	False
1	21	valid	50.0	75.0	False
2	21	valid	50.0	61.84	False
3	20	valid	50.0	48.68	False
4	19	valid	47.5	36.18	False
5	21	valid	50.0	23.03	False
6	17	valid	42.5	11.84	False
7	20	valid	50.0	0	False
8	17	valid	42.5	0	False
9	16	valid	40.0	0	False
10	18	valid	45.0	0	False
11	20	valid	50.0	0	False
12	18	valid	45.0	0	False
13	18	valid	45.0	0	False
14	18	valid	45.0	0	False
15	18	valid	45.0	0	False
16	16	valid	40.0	0	False

Continued on next page

Table o.1 — continued

Step	Horizon	Status	Discharge	SOC	Fallback
17	14	valid	35.0	0	False
18	18	valid	45.0	0	False
19	16	valid	40.0	0	False
20	15	valid	37.5	0	False
21	14	valid	35.0	0	False
22	13	valid	32.5	0	False
23	17	valid	42.5	0	False
24	15	valid	37.5	0	False
25	15	valid	37.5	0	False
26	14	valid	35.0	0	False
27	16	valid	40.0	0	False
28	13	valid	32.5	0	False
29	13	valid	32.5	0	False
30	13	valid	32.5	0	False
31	12	valid	30.0	0	False
32	10	valid	25.0	0	False
33	12	valid	30.0	0	False
34	14	valid	35.0	0	False
35	10	valid	25.0	0	False
36	13	valid	32.5	0	False
37	13	valid	32.5	0	False
38	10	valid	25.0	0	False
39	12	valid	30.0	0	False
40	0	fallback	0.0	0	True
41	0	fallback	0.0	0	True
42	0	fallback	0.0	0	True
43	0	fallback	0.0	0	True
44	0	fallback	0.0	0	True
45	0	fallback	0.0	0	True
46	0	fallback	0.0	0	True
47	0	fallback	0.0	0	True
48	0	fallback	0.0	0	True
49	0	fallback	0.0	0	True
50	0	fallback	0.0	0	True
51	0	fallback	0.0	0	True

Continued on next page

Table o.1 — continued

Step	Horizon	Status	Discharge	SOC	Fallback
52	0	fallback	0.0	0	True
53	0	fallback	0.0	0	True
54	0	fallback	0.0	0	True
55	0	fallback	0.0	0	True
56	6	valid	15.0	0	False
57	4	degraded	10.0	0	False
58	7	valid	17.5	0	False
59	7	valid	17.5	0	False
60	7	valid	17.5	0	False
61	6	valid	15.0	0	False
62	4	degraded	10.0	0	False
63	7	valid	17.5	0	False
64	4	degraded	10.0	0	False
65	3	degraded	7.5	0	False
66	6	valid	15.0	0	False
67	3	degraded	7.5	0	False
68	1	degraded	2.5	0	False
69	3	degraded	7.5	0	False
70	4	degraded	10.0	0	False
71	1	degraded	2.5	0	False
72	2	degraded	5.0	0	False
73	0	fallback	0.0	0	True
74	3	degraded	7.5	0	False
75	2	degraded	5.0	0	False
76	0	fallback	0.0	0	True
77	0	fallback	0.0	0	True
78	1	degraded	2.5	0	False
79	2	degraded	5.0	0	False
80	2	degraded	5.0	0	False
81	0	fallback	0.0	0	True
82	0	fallback	0.0	0	True
83	2	degraded	5.0	0	False
84	0	fallback	0.0	0	True
85	0	fallback	0.0	0	True
86	0	fallback	0.0	0	True

Continued on next page

Table o.1 — continued

Step	Horizon	Status	Discharge	SOC	Fallback
87	0	fallback	0.0	0	True
88	0	fallback	0.0	0	True
89	0	fallback	0.0	0	True
90	0	fallback	0.0	0	True
91	0	fallback	0.0	0	True
92	0	fallback	0.0	0	True
93	0	fallback	0.0	0	True
94	0	fallback	0.0	0	True
95	0	fallback	0.0	0	True

The long contiguous fallback stretches visible above are exactly why the runtime chapter remains honest about scope. The artifact shows explicit fallback behavior and complete audit coverage. It does not show a runtime that has already solved every deployment-side recovery problem.

Appendix p

Dataset Cards

This appendix indexes the source dataset cards that govern the active ORIUS data surfaces. The canonical human-readable cards live under `docs/data_cards/`; the purpose of this appendix is to bind those cards into the thesis control path without duplicating the full markdown bodies in multiple places.

Table p.1: Source dataset cards and manifests for the active three-domain ORIUS claim surface.

Key	Dataset	Card path
OPSD	Open Power System Data	<code>docs/data_cards/opsd.md</code>
	Germany witness row	
EIA-930	U.S. balancing-authority	<code>docs/data_cards/eia930.md</code>
	battery transfer rows	
nuPlan	Bounded AV replay row	<code>docs/data_cards/nuplan.md</code>
	and source manifest	
MIMIC	Credentialed ICU moni-	<code>docs/data_cards/mimic.md</code>
	toring boundary	

Each card records source provenance, split policy, preprocessing, known issues, and the exact role the dataset plays in the active ORIUS claim boundary. The cards are intended to be updated alongside the raw-data manifests and release-manifest refresh pipeline; no raw nuPlan archives are Git-tracked.

Appendix q

Reproducibility Appendix

This appendix binds the active release manifest into the thesis and records the verification path for reproduced artifacts.

- Release manifest: `reports/publication/release_manifest.json`
- Hash verifier: `scripts/verify_reproducibility.py`
- TeX generator: `scripts/generate_reproducibility_appendix.py`

q.1 Release Manifest Summary

The active reproducibility root is `reports/publication/release_manifest.json`.

- Release ID: `R1_20260312T044018Z`
- Master manuscript: `paper/paper.tex`
- Source run families captured: 4
- Hashed publication artifacts: 18

q.2 Source Runs

Dataset	Run ID	Manifest path
DE	<code>R1_..._diag</code>	<code>artifacts/runs/de/R1_20260312T044018Z_diag/registry/run_context.json</code>
US_ERCOT	<code>R1_..._diag</code>	<code>artifacts/runs/us_ercot/R1_20260312T044018Z_diag/registry/run_context.json</code>

Dataset	Run ID	Manifest path
US_MISO	R1..._diag	artifacts/runs/us/_miso/R1_20260312T044018Z_diag/ registry/run_context.json
US_PJM	R1..._diag	artifacts/runs/us/_pjm/R1_20260312T044018Z_diag/registry/ run_context.json

q.3 Tracked Publication Artifacts

Artifact	SHA-256 prefix
ambiguity_calibration_summary.json	49494b142698
cost_safety_pareto.csv	a329765d7015
cqr_calibration_summary.json	79bc1da4338f
cqr_group_coverage.csv	17ea8fae1152
dc3s_fault_breakdown.csv	5994481628fd
dc3s_main_table.csv	37906340fa52
fig_cost_safety_pareto.png	3b9f846e6ecb
fig_cqr_group_coverage.png	20753475262c
fig_distributional_vs_cqr.png	ee709b8524eb
fig_rac_sensitivity_vs_width.png	bfb1391ee35f
fig_true_soc_severity_p95_vs_dropout.png	e632cddb71ae
fig_true_soc_violation_vs_dropout.png	8630704ee24e
rac_cert_summary.json	ab68f8a192e3
stats_summary.json	6d7c99bfd854
table2_ablations.csv	c64fa0274c1e
table5_transfer.csv	a115182eb6b6
table_cqr_distributional_compare.csv	109ddca26cce
transfer_stress.csv	a115182eb6b6

Extended Reading and Index

Suggested Extended Reading

Universal Safety Foundations

Angelopoulos and Bates, *Conformal Prediction: A Gentle Introduction* (2023); Vovk, Gammerman, and Shafer, *Algorithmic Learning in a Random World* (2005); Shafer and Vovk, “A Tutorial on Conformal Prediction” (JMLR 2008).

Conformal Prediction and Calibration

Romano, Patterson, and Candès, “Conformalized Quantile Regression” (NeurIPS 2019); Gibbs and Candès, “Adaptive Conformal Inference Under Distribution Shift” (NeurIPS 2021); Barber et al., “Conformal Prediction Beyond Exchangeability” (AoS 2023); Tibshirani et al., “Conformal Prediction Under Covariate Shift” (2019); Zaffran et al., “Adaptive Conformal Predictions for Time Series” (2022).

Runtime Assurance and Safe Control

Ames et al., “Control Barrier Functions: Theory and Applications” (2019); Bloem et al., “Shield Synthesis” (TACAS 2015); Wabersich and Zeilinger, “A Predictive Safety Filter for Learning-Based Control” (2021); Alshiekh et al., “Safe Reinforcement Learning via Shielding” (2018); Desai et al., “SOTER” (IROS 2019).

Degraded Observation and CPS Trust

Chandola, Banerjee, and Kumar, “Anomaly Detection: A Survey” (2009); Liu, Ting, and Zhou, “Isolation Forest” (2008); Teixeira et al., “A Secure Control Framework for Resource-Limited Adversaries” (2015); Cardenas et al., “Attacks Against Process Control Systems” (2011); Giraldo et al., “A Survey of Physics-Based Attack Detection in Cyber-Physical Systems” (2018).

Three Evidence Lanes

Energy and storage: Tarascon and Armand (2001), Dunn et al. (2011), Rosewater et al. (2019), Hossain et al. (2020). Bounded automated driving replay: Koopman and

Wagner (2016), Shalev-Shwartz et al. (2017), Althoff and Dolan (2014), Herbert et al. (2017). Healthcare monitoring: Lee and Sokolsky (2010), Saeed et al. (2011), Clifford et al. (2009), FDA SaMD guidance (2019).

Reproducibility and Scientific Software

Gneiting and Raftery (2007); Makridakis et al. (2020); Lei et al. (2018); Barber, Candès, Ramdas, and Tibshirani (2023). Use the locked artifact index and publication manifests for all thesis-grade claims.

Key Symbol Index

The dissertation keeps its canonical notation in the core theory and assumption surfaces: see Chapters [3–5](#) and Appendix [b](#).

Artifact Index

The artifact, figure, and table index is maintained in the repository at `reports/publication/` alongside the locked evaluation outputs.

Bibliography

- [1] V. Vovk, A. Gammernan, and G. Shafer, “Algorithmic learning in a random world,” 2005.
- [2] Y. Romano, E. Patterson, and R. J. Candès, “Conformalized quantile regression,” 2019.
- [3] I. Gibbs and E. J. Candès, “Adaptive conformal inference under distribution shift,” 2021.
- [4] Seto, Daisuke and Krogh, Bruce H. and Sha, Lui and Chutinan, Alongkrit, “The simplex architecture for safe online control system upgrades,” 1998, proceedings of ACC.
- [5] Sha, Lui, “Using simplicity to control complexity,” 2001, IEEE Software.
- [6] A. D. Ames, S. Coogan, M. Egerstedt, G. Notomista, K. Sreenath, and P. Tabuada, “Control barrier functions: Theory and applications,” 2019.
- [7] A. D. Ames, J. W. Grizzle, and P. Tabuada, “Control barrier function based quadratic programs with application to adaptive cruise control,” 2014.
- [8] K. P. Wabersich and M. N. Zeilinger, “A predictive safety filter for learning-based control of constrained nonlinear dynamical systems,” 2021.
- [9] UL Standards and Engagement, “UL 4600: Standard for safety for the evaluation of autonomous products,” 2023, underwriters Laboratories standard.
- [10] National Institute of Standards and Technology, “Artificial intelligence risk management framework (ai rmf 1.0),” 2023, nIST AI governance framework.
- [11] International Organization for Standardization, “Iso 26262 road vehicles – functional safety,” 2018, international functional-safety standard.

- [12] International Electrotechnical Commission, “Iec 61508 functional safety of electrical, electronic, and programmable electronic safety-related systems,” 2010, cross-sector functional-safety standard.
- [13] J. Lei, M. G’Sell, A. Rinaldo, R. J. Tibshirani, and L. Wasserman, “Distribution-free predictive inference for regression,” 2018.
- [14] A. N. Angelopoulos and S. Bates, “Conformal prediction: A gentle introduction,” 2023.
- [15] R. F. Barber, E. J. Candès, A. Ramdas, and R. J. Tibshirani, “Conformal prediction beyond exchangeability,” 2023.
- [16] M. Zaffran, O. Féron, Y. Goude, J. Jolivet, and A. Dieuleveut, “Adaptive conformal predictions for time series,” 2022.
- [17] K. Stankeviciute, A. M’Charrak, R. Cornish, and T. Rainforth, “Conformal time-series forecasting,” 2021.
- [18] Xu, Chen and Xie, Yao, “Conformal prediction interval for dynamic time-series,” 2021, proceedings of ICML.
- [19] Sesia, Matteo and Candès, Emmanuel J. and Romano, Yaniv, “A comparison of some conformal quantile regression methods,” 2021, stat.
- [20] Bates, Stephen and Candès, Emmanuel J. and Lei, Lihua and Romano, Yaniv, “Distribution-free, risk-controlling prediction sets,” 2021, journal of the ACM.
- [21] Desai, Ankush and Ghosal, Akash and Saha, Indranil and Yang, Yuhong and Dutle, Aaron and Ivanov, Radoslav and Lee, Insup and Seshia, Sanjit A. and Tiwari, Ashish, “Soter: A runtime assurance framework for programming safe robotics systems,” 2019, proceedings of IROS.
- [22] R. Bloem, B. Könighofer, R. Könighofer, and C. Wang, “Shield synthesis: Runtime enforcement for reactive systems,” 2015.
- [23] Bartocci, Ezio and Falcone, Ylies and Francalanza, Adrian and Reger, Giles, “Introduction to runtime verification,” 2018, lecture Notes in Computer Science.
- [24] Havelund, Klaus and Rosu, Grigore, “An overview of the runtime verification tool java pathexplorer,” 2004, formal Methods in System Design.
- [25] M. Leucker and C. Schallhart, “A brief account of runtime verification,” 2009.

- [26] Schierman, John D. and Wang, Xun J. and Griffith, Dustin and DeCastro, Jonathan A. and Dutle, Aaron and Pike, Lee and Johnson, Taylor T., “Run-time assurance for complex systems,” 2015, aIAA Infotech@Aerospace.
- [27] Bak, Stanley and Manamcheri, Krishna and Mitra, Sayan and Caccamo, Marco, “Sandboxing controllers for cyber-physical systems,” 2011, proceedings of ICCPS.
- [28] Nguyen, Quan and Sreenath, Koushil, “Exponential control barrier functions for enforcing high relative-degree safety-critical constraints,” 2016, proceedings of ACC.
- [29] Wang, Li and Ames, Aaron D. and Egerstedt, Magnus, “Safety barrier certificates for collisions-free multirobot systems,” 2017, iEEE Transactions on Robotics.
- [30] Fisac, Jaime F. and Akametalu, Anayo and Zeilinger, Melanie N. and Kaynama, Sahar and Gillula, Jeremy and Tomlin, Claire J., “A general safety framework for learning-based control in uncertain robotic systems,” 2019, iEEE Transactions on Automatic Control.
- [31] D. Q. Mayne, M. M. Seron, and S. V. Raković, “Robust model predictive control of constrained linear systems with bounded disturbances,” 2005.
- [32] W. Langson, I. Chrysoschoos, S. V. Rakočić, and D. Q. Mayne, “Robust model predictive control using tubes,” 2004.
- [33] J. B. Rawlings, D. Q. Mayne, and M. M. Diehl, “Model predictive control: Theory, computation, and design,” 2017.
- [34] Ben-Tal, Aharon and El Ghaoui, Laurent and Nemirovski, Arkadi, “Robust optimization,” 2009, princeton University Press.
- [35] Camacho, Eduardo F. and Bordons, Carlos, “Model predictive control,” 2013, springer.
- [36] Rajkumar, Ragunathan and Lee, Insup and Sha, Lui and Stankovic, John, “Cyber-physical systems: The next computing revolution,” 2010, proceedings of DAC.
- [37] Derler, Patricia and Lee, Edward A. and Sangiovanni-Vincentelli, Alberto, “Modeling cyber-physical systems,” 2012, proceedings of the IEEE.
- [38] Baheti, Radhakisan and Gill, Helen, “Cyber-physical systems,” 2011, the Impact of Control Technology.

- [39] Pasqualetti, Fabio and Dorfler, Florian and Bullo, Francesco, “Attack detection and identification in cyber-physical systems,” 2013, iEEE Transactions on Automatic Control.
- [40] A. Teixeira, I. Shames, H. Sandberg, and K. H. Johansson, “A secure control framework for resource-limited adversaries,” 2015.
- [41] Basseville, Michele and Nikiforov, Igor V., “Detection of abrupt changes: Theory and application,” 1993, prentice Hall.
- [42] Page, E. S., “Continuous inspection schemes,” 1954, biometrika.
- [43] —, “A test for a change in a parameter occurring at an unknown point,” 1955, biometrika.
- [44] Hinkley, David V., “Inference about the change-point in a sequence of random variables,” 1971, biometrika.
- [45] F. T. Liu, K. M. Ting, and Z.-H. Zhou, “Isolation forest,” 2008.
- [46] Breunig, Markus M. and Kriegel, Hans-Peter and Ng, Raymond T. and Sander, Jorg, “Lof: Identifying density-based local outliers,” 2000, proceedings of SIGMOD.
- [47] Johnson, Alistair E. W. and Bulgarelli, Lucas and Shen, Lu and Gayles, Anna and Shammout, Ahmad and Horng, Steven and Pollard, Tom and Hao, Shirley and Moody, Benjamin and Gow, Brian and others, “Mimic-iv, a freely accessible electronic health record dataset,” 2023, scientific Data.
- [48] Johnson, Alistair E. W. and Bulgarelli, Lucas and Pollard, Tom J. and Celi, Leo Anthony and Mark, Roger G., “Mimic-iv version 3.1,” 2024, physioNet.
- [49] Pollard, Tom J. and Johnson, Alistair E. W. and Raffa, Jonathan D. and Celi, Leo Anthony and Mark, Roger G. and Badawi, Omar, “The eicu collaborative research database, a freely available multi-center database for critical care research,” 2018, scientific Data.
- [50] Collins, Gary S. and Moons, Karel G. M. and Dhiman, Paula and Riley, Richard D. and Beam, Andrew L. and Van Calster, Ben and Ghassemi, Marzyeh and Liu, Xiaoxuan and Reitsma, Johannes B. and van Smeden, Maarten, “Tripod+ai statement: Updated guidance for reporting clinical prediction models that use regression or machine learning methods,” 2024, *BMJ* 2024;385:e078378.

- [51] Moons, Karel G. M. and Damen, Johanna A. A. and Kaul, Tabea and Hooft, Lotty and Andaur Navarro, Constanza and Dhiman, Paula and Beam, Andrew L. and Van Calster, Ben and Celi, Leo Anthony, “Probast+ai: An updated quality, risk of bias, and applicability assessment tool for prediction models using regression or artificial intelligence methods,” 2025, *BMJ* 2025;388:e082505.
- [52] Liu, Xiaoxuan and Rivera, Samantha Cruz and Moher, David and Calvert, Melanie J. and Denniston, Alastair K., “Reporting guidelines for clinical trial reports for interventions involving artificial intelligence: The consort-ai extension,” 2020, *nature Medicine* 26, 1364–1374.
- [53] Vasey, Bonnie and Nagendran, Myura and Campbell, Bruce and Clifton, David A. and Collins, Gary S. and Denaxas, Spiros and Denniston, Alastair K. and Faes, Livia and Geerts, Bart and Ibrahim, Mohammed and Liu, Xiaoxuan and Mateen, Bilal A. and McCradden, Melissa D. and McLennan, Stuart and Morgan, Laura and Ordish, Johan and Rogers, Campbell and Saria, Suchi and Ting, Daniel S. W. and Watkinson, Peter and Weber, Wim and Wheatstone, Peter and McCulloch, Peter, “Reporting guideline for the early-stage clinical evaluation of decision support systems driven by artificial intelligence: Decide-ai,” 2022, *nature Medicine* 28, 924–933.
- [54] STARD-AI Steering Group, “Stard-ai reporting guideline for diagnostic accuracy studies involving artificial intelligence,” 2025, *nature Medicine*.
- [55] Y. Romano, E. Patterson, and R. J. Candès, “Conformalized quantile regression,” 2019.
- [56] I. Gibbs and E. J. Candès, “Adaptive conformal inference under distribution shift,” 2021.
- [57] V. Vovk, A. Gammerman, and G. Shafer, “Algorithmic learning in a random world,” 2005.
- [58] T. Chen and C. Guestrin, “Xgboost: A scalable tree boosting system,” 2016.
- [59] G. Ke *et al.*, “Lightgbm: A highly efficient gradient boosting decision tree,” 2017.
- [60] B. N. Oreshkin, D. Carpón, N. Chapados, and Y. Bengio, “N-beats: Neural basis expansion analysis for interpretable time series forecasting,” 2020.
- [61] B. Lim, S. Arık, N. Loeff, and T. Pfister, “Temporal fusion transformers for interpretable multi-horizon time series forecasting,” 2021.

- [62] Y. Nie *et al.*, “A time series is worth 64 words: Long-term forecasting with transformers,” 2023.
- [63] A. Wachi, Y. Sui, Y. Yue, and M. Ono, “Safe reinforcement learning in constrained markov decision processes,” 2024.
- [64] D. Amodei *et al.*, “Concrete problems in ai safety,” 2016.
- [65] A. Bemporad and M. Morari, “Robust model predictive control: A survey,” 2007.
- [66] A. A. Cardenas, S. Amin, Z.-S. Lin, Y.-L. Huang, C.-Y. Huang, and S. Sastry, “Attacks against process control systems: Risk assessment, detection, and response,” 2011.
- [67] Y. Liang, C.-Z. Zhao, H. Yuan, Z. Chen, W. Zhang, J.-Q. Huang, D. Yu, Y. Liu, M.-M. Titirici, Y.-L. Chueh, H. Yu, and Q. Zhang, “A review of rechargeable batteries for portable electronic devices,” 2019.
- [68] A. Shapiro, D. Dentcheva, and A. Ruszczyński, “Lectures on stochastic programming: Modeling and theory,” 2021.
- [69] H. Sun, J. Chen, Q. Ye, and J. Shao, “Adversarial attacks on neural networks for logical reasoning,” 2007.
- [70] C. Xu and Y. Xie, “Conformal prediction interval for dynamic time-series,” 2021.
- [71] J. García and F. Fernández, “A comprehensive survey on safe reinforcement learning,” 2015.
- [72] R. J. Tibshirani, B. Efron, R. F. Barber, E. J. Candès, and A. Ramdas, “Conformal prediction under covariate shift,” 2019.
- [73] H. Papadopoulos, K. Proedrou, V. Vovk, and A. Gammerman, “Inductive confidence machines for regression,” 2002.
- [74] S. Bates, E. J. Candès, L. Lei, Y. Romano, and M. Sesia, “Testing for outliers with conformal p-values,” 2023.
- [75] S. Prajna and A. Jadbabaie, “Safety verification of hybrid systems using barrier certificates,” 2004.
- [76] A. J. Taylor and A. D. Ames, “Adaptive safety with control barrier functions,” 2020.

- [77] P. Wieland and F. Allgöwer, “Constructive safety using control barrier functions,” 2007.
- [78] S. V. Rakočić, E. C. Kerrigan, K. I. Kouramas, and D. Q. Mayne, “Invariant approximations of the minimal robust positively invariant set,” 2005.
- [79] M. Lorenzen, F. Dabbene, R. Tempo, and F. Allgöwer, “Constraint-tightening and stability in stochastic model predictive control,” 2019.
- [80] J. Achiam, D. Held, A. Tamar, and P. Abbeel, “Constrained policy optimization,” 2017.
- [81] L. Brunke, M. Greeff, A. W. Hall, Z. Yuan, S. Zhou, J. Panerati, and A. P. Schoellig, “Safe learning in robotics: From learning-based control to safe reinforcement learning,” 2022.
- [82] T. M. Moldovan and P. Abbeel, “Safe exploration in markov decision processes,” 2012.
- [83] A. Ray, J. Achiam, and D. Amodei, “Benchmarking safe exploration in deep reinforcement learning,” 1910.
- [84] E. Altman, “Constrained markov decision processes,” 1999.
- [85] Y. Falcone, K. Havelund, and G. Reger, “A tutorial on runtime verification,” 2013.
- [86] E. Bartocci, Y. Falcone, A. Francalanza, and G. Reger, “Introduction to runtime verification,” 2018.
- [87] A. Pnueli and A. Zaks, “Psl model checking and run-time verification via testers,” 2006.
- [88] O. Maler and D. Nickovic, “Monitoring temporal properties of continuous signals,” 2004.
- [89] D. I. Urbina, J. A. Giraldo, A. A. Cardenas, N. O. Tippenhauer, J. Valente, M. Faisal, J. Ruths, R. Candell, and H. Sandberg, “Limiting the impact of stealthy attacks on industrial control systems,” 2016.
- [90] Y. Mo and B. Sinopoli, “Secure estimation in the presence of integrity attacks,” 2015.
- [91] J. A. Giraldo, D. Urbina, A. A. Cardenas, J. Valente, M. Faisal, J. Ruths, N. O. Tippenhauer, H. Sandberg, and R. Candell, “A survey of physics-based attack detection in cyber-physical systems,” 2018.

- [92] X. Hu, S. Li, and H. Peng, “A comparative study of equivalent circuit models for li-ion batteries,” 2012.
- [93] J. M. Tarascon and M. Armand, “Issues and challenges facing rechargeable lithium batteries,” 2001.
- [94] B. Dunn, H. Kamath, and J.-M. Tarascon, “Electrical energy storage for the grid: A battery of choices,” 2011.
- [95] S. Bose, S. Fattahi, and A. Swami, “Stochastic battery operation in a competitive electricity market,” 2019.
- [96] D. M. Rosewater, D. A. Copp, T. A. Nguyen, R. H. Byrne, and S. Santoso, “Battery energy storage models for optimal control,” 2019.
- [97] T. Hong and S. Fan, “Probabilistic electric load forecasting: A tutorial review,” 2016.
- [98] T. Gneiting and A. E. Raftery, “Strictly proper scoring rules, prediction, and estimation,” 2007.
- [99] S. Makridakis, E. Spiliotis, and V. Assimakopoulos, “The m4 competition: 100,000 time series and 61 forecasting methods,” 2020.
- [100] J. W. Taylor, “Short-term electricity demand forecasting using double seasonal exponential smoothing,” 2003.
- [101] P. Koopman and M. Wagner, “Autonomous vehicle safety: An interdisciplinary challenge,” 2016.
- [102] S. Shalev-Shwartz, S. Shammah, and A. Shashua, “On a formal model of safe and scalable self-driving cars,” 2017.
- [103] M. Althoff and J. M. Dolan, “Online verification of automated road vehicles using reachability analysis,” 2014.
- [104] T. Nolte, M. Sjodin, K. Akesson, K. Lundqvist, M. Behnam, M. Lindberg, and C. Norrmann, “Towards bringing research on functional safety in automotive systems to practice,” 2017.
- [105] I. Lee and O. Sokolsky, “Medical cyber physical systems,” 2010.

- [106] M. Saeed, M. Villarroel, A. T. Reisner, G. Clifford, L. W. Lehman, G. Moody, T. Heldt, T. H. Kyaw, B. Moody, and R. G. Mark, “Multiparameter intelligent monitoring in intensive care ii (mimic-ii): A public-access intensive care unit database,” 2011.
- [107] G. D. Clifford, W. J. Long, G. B. Moody, and P. Szolovits, “Robust parameter extraction for decision support using multimodal intensive care data,” 2009.
- [108] U.S. Food and Drug Administration, “Software as a medical device (samd): Clinical evaluation,” 2019.
- [109] RTCA, “Do-178c: Software considerations in airborne systems and equipment certification,” 2011.
- [110] J. Rushby, “Formal methods and the certification of critical systems,” 1993.
- [111] I. Habli and T. Kelly, “Model-based assurance of control system safety and security,” 2010.
- [112] L. P. Kaelbling, M. L. Littman, and A. R. Cassandra, “Planning and acting in partially observable stochastic domains,” 1998.
- [113] S. J. Julier and J. K. Uhlmann, “Unscented filtering and nonlinear estimation,” 2004.
- [114] R. T. Rockafellar and S. Uryasev, “Optimization of conditional value-at-risk,” 2000.
- [115] S. J. Qin and T. A. Badgwell, “A survey of industrial model predictive control technology,” 2003.
- [116] S. Yin, S. X. Ding, X. Xie, and H. Luo, “A review on basic data-driven approaches for industrial process monitoring,” 2014.
- [117] R. Isermann, “Model-based fault-detection and diagnosis—status and applications,” 2005.
- [118] V. Chandola, A. Banerjee, and V. Kumar, “Anomaly detection: A survey,” 2009.
- [119] A. B. Tsybakov, “Introduction to nonparametric estimation,” 2009.
- [120] B. Yu, “Assouad, fano, and le cam,” 1997.
- [121] M. Blanke, M. Kinnaert, J. Lunze, and M. Staroswiecki, “Diagnosis and fault-tolerant control,” 2006.

- [122] Y. Zhang and J. Jiang, “Bibliographical review on reconfigurable fault-tolerant control systems,” 2008.
- [123] L. Lindemann, M. Cleaveland, G. Shim, and G. J. Pappas, “Safe planning in dynamic environments using conformal prediction,” 2023.
- [124] A. Dixit, L. Lindemann, S. X. Wei, M. Cleaveland, G. J. Pappas, and J. V. Deshmukh, “Adaptive conformal prediction for motion planning among dynamic agents,” 2023.
- [125] A. N. Angelopoulos, S. Bates, E. J. Candès, M. I. Jordan, and L. Lei, “Conformal risk control,” 2024.
- [126] Chernozhukov, Victor and Wuthrich, Kaspar and Zhu, Yinchu, “Distributional conformal prediction,” 2021, proceedings of the National Academy of Sciences.
- [127] Feldman, Shai and Bates, Stephen and Candes, Emmanuel J., “Calibrated multiple-output quantile regression with conformal prediction,” 2021, arXiv preprint.
- [128] Fisch, Adam and Grimson, W. Eric and Vovk, Vladimir and Katz-Samuels, Julian, “Online conformal prediction with decaying step sizes,” 2021, arXiv preprint.
- [129] Cofer, Darren and Rangarajan, Anand and Janssen, Wil and Havelund, Klaus and Perez, Ivan, “Assurance cases for runtime assurance frameworks,” 2012, nASA Technical Report.
- [130] Herbert, Sylvia and Chen, Mo and Han, Siyuan and Bansal, Somil and Fisac, Jaime F. and Tomlin, Claire, “Fastrack: A modular framework for fast and guaranteed safe motion planning,” 2017, proceedings of CDC.
- [131] Sen, Koushik and Vardhan, Abhay and Agha, Gul and Rosu, Grigore, “Efficient decentralized monitoring of safety in distributed systems,” 2004, proceedings of ICSE.
- [132] d’Angelo, Ben and Sankaranarayanan, Sriram and Sanchez, Cesar and Robinson, Will and Finkbeiner, Bernd and Sipma, Henny and Mehrotra, Sharad and Manna, Zohar, “Lola: Runtime monitoring of synchronous systems,” 2005, tIME.
- [133] Xu, Xiangru and Tabuada, Paulo and Grizzle, Jessy W. and Ames, Aaron D., “Robustness of control barrier functions for safety critical control,” 2015, iFAC-PapersOnLine.
- [134] Cheng, Richard and Orosz, Gyorgy and Murray, Richard M. and Burdick, Joel W., “End-to-end safe reinforcement learning through barrier functions for safety-critical continuous control tasks,” 2019, proceedings of AAAI.

- [135] Cheng, Richard and Nurnberger, Andrew and Scheibler, Karol and Ivanov, Radoslav and Pappas, George J., “Cautious adaptation for reinforcement learning under safety constraints,” 2021, proceedings of L4DC.
- [136] Wabersich, Kai P. and Zeilinger, Melanie N. and Hewing, Lukas and Carron, Andrea, “Data-driven safety filters: A survey and perspective,” 2023, annual Review of Control, Robotics, and Autonomous Systems.
- [137] Koller, Torsten and Berkenkamp, Felix and Turchetta, Matteo and Krause, Andreas, “Learning-based model predictive control for safe exploration,” 2018, proceedings of CDC.
- [138] Berkenkamp, Felix and Turchetta, Matteo and Schoellig, Angela P. and Krause, Andreas, “Safe model-based reinforcement learning with stability guarantees,” 2017, advances in Neural Information Processing Systems.
- [139] Mo, Yilin and Sinopoli, Bruno, “False data injection attacks in control systems,” 2012, proceedings of the First Workshop on Secure Control Systems.
- [140] Urbina, David I. and Giraldo, Jairo and Cardenas, Alvaro A. and Tippenhauer, Nils Ole, “Survey and new directions for physics-based attack detection in control systems,” 2016, proceedings of CPS-SPC.
- [141] Khraisat, Ansam and Gondal, Iqbal and Vamplew, Peter and Kamruzzaman, Joarder, “Survey of intrusion detection systems: Techniques, datasets and challenges,” 2019, cybersecurity.
- [142] Salinas, David and Flunkert, Valentin and Gasthaus, Jan and Januschowski, Tim, “Deepar: Probabilistic forecasting with autoregressive recurrent networks,” 2020, international Journal of Forecasting.
- [143] Wu, Haixu and Xu, Jiehui and Wang, Jianmin and Long, Mingsheng, “Autoformer: Decomposition transformers with auto-correlation for long-term series forecasting,” 2021, advances in Neural Information Processing Systems.
- [144] Zhou, Haoyi and Zhang, Shanghang and Peng, Jieqi and Zhang, Shuai and Li, Jianxin and Xiong, Hui and Zhang, Wancai, “Informer: Beyond efficient transformer for long sequence time-series forecasting,” 2021, proceedings of AAAI.
- [145] Zhou, Tian and Ma, Ziqing and Wen, Qingsong and Wang, Xue and Sun, Liang and Jin, Rong and Zhou, Xiaokang and Qian, Wending and Li, Xue and Zhang, Weifeng

- and others, “Fedformer: Frequency enhanced decomposed transformer for long-term series forecasting,” 2022, proceedings of ICML.
- [146] Wu, Haixu and Hu, Tao and Liu, Yong and Zhou, Hang and Wang, Jianmin and Long, Mingsheng, “Timesnet: Temporal 2d-variation modeling for general time series analysis,” 2023, proceedings of ICLR.
- [147] Zeng, Ailing and Chen, Muxin and Zhang, Lei and Xu, Qiang, “Are transformers effective for time series forecasting?” 2023, proceedings of AAAI.
- [148] Dosovitskiy, Alexey and Beyer, Lucas and Kolesnikov, Alexander and others, “An image is worth 16x16 words: Transformers for image recognition at scale,” 2021, proceedings of ICLR.
- [149] Vaswani, Ashish and Shazeer, Noam and Parmar, Niki and Uszkoreit, Jakob and Jones, Llion and Gomez, Aidan and Kaiser, Lukasz and Polosukhin, Illia, “Attention is all you need,” 2017, advances in Neural Information Processing Systems.
- [150] Severson, Kristen A. and Attia, Peter M. and Jin, Norman and Perkins, Nicholas and Jiang, Ben and Yang, Zi and Chen, Michael and Aykol, Muratahan and Herring, Patrick K. and Fraggedakis, Dimitrios and others, “Data-driven prediction of battery cycle life before capacity degradation,” 2019, nature Energy.
- [151] Berecibar, Mikel and Gandiaga, Iker and Villarreal, Ibon and Omar, Noshin and Van Mierlo, Joeri and Van den Bossche, Peter, “Critical review of state of health estimation methods of li-ion batteries for real applications,” 2016, renewable and Sustainable Energy Reviews.
- [152] He, Hongwen and Xiong, Rui and Fan, Jianxin, “Evaluation of lithium-ion battery equivalent circuit models for state of charge estimation,” 2011, applied Energy.
- [153] Hu, Xiaosong and Li, Shengbo Eben and Peng, Huei, “A comparative study of equivalent circuit models for li-ion batteries,” 2012, journal of Power Sources.
- [154] Weitzel, Tim and Glock, Christoph H., “Energy storage dispatch in energy markets: A review,” 2018, energy.
- [155] Xu, Bing and Oudalov, Alexandre and Ulbig, Andreas and Andersson, Goran and Kirschen, Daniel S., “Modeling of lithium-ion battery degradation for cell life assessment,” 2020, IEEE Transactions on Smart Grid.

- [156] Wang, Yi and Zhang, Dong and Li, Xinyu, “Battery storage operation under uncertainty: A distributionally robust perspective,” 2022, IEEE Transactions on Smart Grid.
- [157] Hossain, Eklas and Faruque, Hasan and Sunny, Md Shahjalal and Mohammad, Naimul and Nawar, Nurul, “A comprehensive review on grid-connected battery energy storage systems,” 2020, energies.
- [158] Plett, Gregory L., “Battery management systems, volume ii: Equivalent-circuit methods,” 2015, artech House.
- [159] —, “Battery management systems, volume iii: Battery state estimation,” 2016, artech House.
- [160] Srinivasan, Ranjani and Sanner, Scott and others, “Runtime assurance for safe learning-enabled aerospace systems,” 2020, aIAA Scitech.
- [161] Herbert, Sylvia and Chen, Mo and Han, Siyuan and Bansal, Somil and Fisac, Jaime and Tomlin, Claire, “Fastrack: A modular framework for fast and guaranteed safe motion planning,” 2017, proceedings of CDC.
- [162] Ames, Aaron D. and Molnar, Tamas and Orosz, Gabor and Sreenath, Koushil and Tabuada, Paulo, “Safety-critical control for autonomous systems,” 2021, annual Review of Control, Robotics, and Autonomous Systems.
- [163] Lederer, Johannes and Kohler, Jonas and Berkenkamp, Felix and Zeilinger, Melanie N., “A survey on safe reinforcement learning,” 2019, arXiv preprint.
- [164] Ulbrich, Simon and Reschka, Andreas and Rieken, Jonathan and Ernst, Sebastian and Bagschik, Gereon and Damm, Werner and Maurer, Markus, “Towards a functional system architecture for automated vehicles,” 2015, proceedings of ITSC.
- [165] Schwarting, Wilko and Alonso-Mora, Javier and Rus, Daniela, “Planning and decision-making for autonomous vehicles,” 2018, annual Review of Control, Robotics, and Autonomous Systems.
- [166] Badue, Claudine and Guidolini, Rafael and Carneiro, Raphael and Berriel, Rodrigo and Paixao, Thiago and Mutz, Filipe and de Paula Veronese, Luis and Oliveira-Santos, Thiago, “Self-driving cars: A survey,” 2021, expert Systems with Applications.

- [167] Paszke, Adam and Gross, Sam and Massa, Francisco and others, “Pytorch: An imperative style, high-performance deep learning library,” 2019, advances in Neural Information Processing Systems.
- [168] Sutton, Richard S. and Barto, Andrew G., “Reinforcement learning: An introduction,” 2018, mIT Press.
- [169] Kulkarni, Saurabh and Sahoo, Sarmila and Mohanty, Sasmita, “Industrial process monitoring and fault diagnosis: A survey,” 2022, computers and Chemical Engineering.
- [170] Yin, Shen and Ding, Steven and Xie, Xiaoxia and Luo, Haifeng, “A review on basic data-driven approaches for industrial process monitoring,” 2014, iEEE Transactions on Industrial Electronics.
- [171] Chiang, Leo H. and Russell, Edwin L. and Braatz, Richard D., “Fault detection and diagnosis in industrial systems,” 2001, springer.
- [172] Clifton, Lei and Clifton, David A. and Pimentel, Marco A. F. and Watkinson, Peter J. and Tarassenko, Lionel, “Gaussian processes for personalized e-health monitoring,” 2012, iEEE Transactions on Biomedical Engineering.
- [173] Goldstein, Benjamin A. and Navar, Ann Marie and Pencina, Michael J. and Ioannidis, John P. A., “Opportunities and challenges in developing risk prediction models with electronic health records data,” 2017, journal of the American Medical Informatics Association.
- [174] Johnson, Alistair E. W. and Pollard, Tom J. and Shen, Lu and Lehman, Li-wei H. and Feng, Mengling and Ghassemi, Mohammad and Moody, Benjamin and Szolovits, Peter and Celi, Leo Anthony and Mark, Roger G., “Mimic-iii, a freely accessible critical care database,” 2016, scientific Data.
- [175] Esteva, Andre and Robicquet, Alexandre and Ramsundar, Bharath and Kuleshov, Volodymyr and DePristo, Mark and Chou, Katie and Cui, Claire and Corrado, Greg and Thrun, Sebastian and Dean, Jeff, “A guide to deep learning in healthcare,” 2019, nature Medicine.
- [176] Lee, Edward A., “The past, present and future of cyber-physical systems: A focus on models,” 2015, sensors.
- [177] Wolf, Mark T. and others, “Safety and trustworthiness in physical ai systems,” 2018, aAAI Spring Symposium.

- [178] Gernstedt, Erik and Nilsson, Ola and Smith, Reid, “Physical ai systems: Architectures, risks, and safety layers,” 2024, arXiv preprint.
- [179] Astrom, Karl Johan and Murray, Richard M., “Feedback systems: An introduction for scientists and engineers,” 2008, princeton University Press.
- [180] Zhou, Kemin and Doyle, John C. and Glover, Keith, “Robust and optimal control,” 1996, prentice Hall.
- [181] Bertsekas, Dimitri P., “Dynamic programming and optimal control,” 2005, athena Scientific.
- [182] Vovk, Vladimir and Gammernan, Alex and Shafer, Glenn, “Conformal predictors and confidence machines,” 2005, algorithmic Learning in a Random World.
- [183] Shafer, Glenn and Vovk, Vladimir, “A tutorial on conformal prediction,” 2008, journal of Machine Learning Research.
- [184] Barber, Rina Foygel and Candes, Emmanuel J. and Ramdas, Aaditya and Tibshirani, Ryan J., “Conformal prediction beyond exchangeability,” 2023, annals of Statistics.
- [185] Kiyani, Shayan and Pappas, George J. and Roth, Aaron and Hassani, Hamed, “Decision theoretic foundations for conformal prediction: Optimal uncertainty quantification for risk-averse agents,” in *Proceedings of the 42nd International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, vol. 267, 2025, pp. 30 943–30 965.
- [186] Kwon, Minjae and Ingebrand, Tyler and Topcu, Ufuk and Feng, Lu, “Runtime safety through adaptive shielding: From hidden parameter inference to provable guarantees,” 2025, arXiv:2506.11033.
- [187] Alshiekh, Mohammad and Bloem, Roderick and Ehlers, Ruediger and Konighofer, Bettina and Niekum, Scott and Topcu, Ufuk, “Safe reinforcement learning via shielding,” 2018, proceedings of AAAI.
- [188] Konighofer, Bettina and Alshiekh, Mohammad and Bloem, Roderick, “Shielding techniques for safe reinforcement learning,” 2020, formal Methods in System Design.
- [189] Bastani, Osbert, “Safe reinforcement learning: A control-theoretic perspective,” 2021, foundations and Trends in Machine Learning.

- [190] Bertsimas, Dimitris and Brown, David B. and Caramanis, Constantine, “Theory and applications of robust optimization,” 2011, *SIAM Review*.
- [191] Birge, John R. and Louveaux, Francois, “Introduction to stochastic programming,” 2011, *Springer*.
- [192] Borisov, Vadim and Leemann, Tobias and Seßler, Kathrin and Haug, Johannes and Pawelczyk, Martin and Kasneci, Gjergji, “Deep neural networks and tabular data: A survey,” 2022, *IEEE Transactions on Neural Networks and Learning Systems*.
- [193] Chen, Min and Lu, Shengnan and Wang, Zhen and Wang, Xin, “A review of battery energy storage system safety and reliability,” 2020, *Journal of Energy Storage*.
- [194] Dibaji, Seyed Mohammad and Pirani, Mohsen and Flamholz, David and Annaswamy, Anuradha and Johansson, Karl Henrik and Chakraborty, Aranya, “A systems and control perspective of CPS security,” 2019, *Annual Reviews in Control*.
- [195] Lamport, Leslie and Shostak, Robert and Pease, Marshall, “The byzantine generals problem,” 1982, *ACM Transactions on Programming Languages and Systems*.
- [196] Lundberg, Scott M. and Lee, Su-In, “A unified approach to interpreting model predictions,” 2017, *Advances in Neural Information Processing Systems*.
- [197] Pineau, Joelle and Vincent-Lamarre, Philippe and Sinha, Koustuv and Larivière, Vincent and Beygelzimer, Alina and d’Alché-Buc, Florence and Fox, Emily and Larochelle, Hugo, “Improving reproducibility in machine learning research,” 2021, *Journal of Machine Learning Research*.
- [198] Rosewater, David and Ferraro, Paul and Byrne, Raymond and Santoso, Surya, “Risk-averse battery dispatch under forecast error,” 2020, *IEEE Transactions on Smart Grid*.
- [199] Leveson, Nancy, “Engineering a safer world: Systems thinking applied to safety,” 2011, *MIT Press*.
- [200] Koopman, Philip and Wagner, Michael, “Challenges in autonomous vehicle testing and validation,” 2016, *SAE International Journal of Transportation Safety*.
- [201] Dawid, A. Philip, “The well-calibrated bayesian,” 1982, *Journal of the American Statistical Association*.
- [202] Platt, John, “Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods,” 1999, *Advances in Large Margin Classifiers*.

- [203] Guo, Chuan and Pleiss, Geoff and Sun, Yu and Weinberger, Kilian Q., "On calibration of modern neural networks," 2017, proceedings of ICML.
- [204] Kuleshov, Volodymyr and Fenner, Nathan and Ermon, Stefano, "Accurate uncertainties for deep learning using calibrated regression," 2018, proceedings of ICML.
- [205] Vaicenavicius, Joris and Widmann, David and Andersson, Carl and Lindsten, Fredrik and Roll, Jacob and Sch"on, Thomas B., "Evaluating model calibration in classification," 2019, proceedings of AISTATS.
- [206] Nixon, Jeremy and Dusenberry, Michael W. and Zhang, Linchuan and Jerfel, Ghas-sen and Tran, Dustin, "Measuring calibration in deep learning," 2019, proceedings of CVPR Workshops.
- [207] Lee, Edward A., "Cyber physical systems: Design challenges," 2008, proceedings of ISORC.
- [208] Kopetz, Hermann, "Real-time systems: Design principles for distributed embedded applications," 2011, springer.
- [209] Alur, Rajeev, "Principles of cyber-physical systems," 2015, mIT Press.
- [210] Blanchini, Franco, "Set invariance in control," 1999, automatica.
- [211] Aubin, Jean-Pierre, "Viability theory," 1991, birkh"ausser.
- [212] Mitchell, Ian M. and Bayen, Alexandre M. and Tomlin, Claire J., "A time-dependent hamilton-jacobi formulation of reachable sets for continuous dynamic games," 2005, iEEE Transactions on Automatic Control.
- [213] Baier, Christel and Katoen, Joost-Pieter, "Principles of model checking," 2008, mIT Press.
- [214] Clarke, Edmund M. and Grumberg, Orna and Peled, Doron, "Model checking," 1999, mIT Press.
- [215] Geiger, Andreas and Lenz, Philip and Urtasun, Raquel, "Are we ready for autonomous driving? the kitti vision benchmark suite," 2012, proceedings of CVPR.
- [216] Caesar, Holger and Bankiti, Varun and Lang, Alex H. and Vora, Sourabh and Liong, Venice Erin and Xu, Qiang and Krishnan, Anush and Pan, Yu and Baldan, Giancarlo and Beijbom, Oscar, "nuscenec: A multimodal dataset for autonomous driving," 2020, proceedings of CVPR.

- [217] Ettinger, Scott and Cheng, Shuyang and Caine, Benjamin and Liu, Chenxi and Zhao, Hang and Pradhan, Saurabh and Chai, Yuning and Sapp, Benjamin and Qi, Charles R. and Zhou, Yin and others, “Large scale interactive motion forecasting for autonomous driving: The waymo open motion dataset,” 2021, proceedings of ICCV.
- [218] Caesar, Holger and Kabzan, Juraj and Tan, Kok Seang and Fong, Whye Kit and Wolff, Eric and Lang, Alex and Fletcher, Luke and Beijbom, Oscar and Omari, Sammy, “nuplan: A closed-loop ml-based planning benchmark for autonomous vehicles,” 2021, arXiv:2106.11810.
- [219] Karnchanachari, Napat and others, “Towards learning-based planning: The nuplan benchmark for real-world autonomous driving,” 2024, arXiv:2403.04133.
- [220] Peng, Mingxing and Yao, Ruoyu and Guo, Xusen and Ma, Jun, “nuplan-r: A closed-loop planning benchmark for autonomous driving via reactive multi-agent simulation,” 2025, arXiv:2511.10403.
- [221] Wilson, Benjamin and Qi, Wenda and Agarwal, Tushar and Lambert, John and Singh, Jitendra and Khandelwal, Siddhesh and Pan, Bowen and Wang, Cong and D’Sa, Sharath and Lee, Seung and others, “Argoverse 2: Next generation datasets for self-driving perception and forecasting,” 2021, arXiv preprint.
- [222] Downs, James J. and Vogel, Ernest F., “A plant-wide industrial process control problem,” 1993, computers and Chemical Engineering.
- [223] Rieth, Cory A. and Amsel, Ben D. and Tran, Rene and Cook, Miles B., “Revisiting the tennessee eastman process simulation benchmark for industrial machine learning,” 2017, proceedings of the AIChE Annual Meeting.
- [224] Goldberger, Ary L. and Amaral, Luis A. N. and Glass, Leon and Hausdorff, Jeffrey M. and Ivanov, Plamen Ch. and Mark, Roger G. and Mietus, Joseph E. and Moody, George B. and Peng, C.-K. and Stanley, H. Eugene, “Physiobank, physiotookit, and physionet: Components of a new research resource for complex physiologic signals,” 2000, circulation.
- [225] Huang, Yong and Yang, Zhongqi and Rahmani, Amir M., “Mimic-sepsis: A curated benchmark for modeling and learning from sepsis trajectories in the icu,” 2025, iEEE BHI / OpenReview.
- [226] Saxena, Abhinav and Goebel, Kai, “Turbofan engine degradation simulation data set,” 2008, nASA Ames Prognostics Data Repository.

- [227] Ramasso, Emmanuel and Saxena, Abhinav, "Performance benchmarking and analysis of prognostic methods for c-mapss datasets," 2014, international Journal of Prognostics and Health Management.
- [228] Sch" afer, Matthias and Strohmeier, Martin and Olive, Xavier and Smith, Matthew and Lenders, Vincent and Martinovic, Ivan, "Opensky report 2020: Analysing in-flight emergency events using big data," 2020, openSky Network Report.
- [229] Strohmeier, Martin and Sch" afer, Matthias and Lenders, Vincent and Martinovic, Ivan, "On perception and reality in wireless air traffic communication security," 2014, proceedings of IEEE CNS.
- [230] Peng, Roger D., "Reproducible research in computational science," 2011, science.
- [231] Sandve, Geir Kjetil and Nekrutenko, Anton and Taylor, James and Hovig, Eivind, "Ten simple rules for reproducible computational research," 2013, pLoS Computational Biology.
- [232] Wilson, Greg and Aruliah, D. A. and Brown, C. Titus and Chue Hong, Neil P. and Davis, Matt and Guy, Richard T. and Haddock, Steven H. D. and Huff, Kathryn D. and Mitchell, Ian M. and Plumbley, Mark D. and others, "Best practices for scientific computing," 2014, pLoS Biology.
- [233] Amershi, Saleema and Begel, Andrew and Bird, Christian and DeLine, Robert and Gall, Harald and Kamar, Ece and Nagappan, Nachiappan and Nushi, Besmira and Zimmermann, Thomas, "Software engineering for machine learning: A case study," 2019, proceedings of ICSE-SEIP.
- [234] Doshi-Velez, Finale and Kim, Been, "Towards a rigorous science of interpretable machine learning," 2017, arXiv preprint.
- [235] Rudin, Cynthia, "Stop explaining black box machine learning models for high stakes decisions and use interpretable models instead," 2019, nature Machine Intelligence.
- [236] Lipton, Zachary C., "The mythos of model interpretability," 2018, queue.
- [237] Russell, Stuart, "Human compatible: Artificial intelligence and the problem of control," 2019, viking.
- [238] Koopman, Philip and Wagner, Michael, "Autonomous vehicle safety: An interdisciplinary challenge," 2017, iEEE Intelligent Transportation Systems Magazine.

- [239] Favarò, Francesca M. and Eurich, Samantha O. and Nader, Nadine and Tripp, Matthew, “Examining accident reports involving autonomous vehicles in california,” 2017, pLoS ONE.
- [240] Kalra, Nidhi and Paddock, Susan M., “Driving to safety: How many miles of driving would it take to demonstrate autonomous vehicle reliability?” 2016, transportation Research Part A.
- [241] Cummings, Mary L., “Managing the complexity of autonomous vehicle safety,” 2017, journal of Transportation Safety and Security.
- [242] Dolan, John M. and others, “The field benchmark problem and evaluation standard for mobile robots,” 1994, aAAI Mobile Robot Workshop.
- [243] Thrun, Sebastian and Montemerlo, Michael and Dahlkamp, Hendrik and Stavens, David and Aron, Alex and Diebel, James and Fong, Philip and Gale, John and Halpenny, Morgan and Hoffmann, Gabe and others, “Stanley: The robot that won the darpa grand challenge,” 2006, journal of Field Robotics.
- [244] Montemerlo, Michael and Becker, Jan and Bhat, Surya and Dahlkamp, Hendrik and Dolgov, Dmitri and Ettinger, Scott and Haehnel, Dirk and Hilden, Tim and Hoffmann, Gabe and Huhnke, Burkhard and others, “Junior: The stanford entry in the urban challenge,” 2008, journal of Field Robotics.
- [245] Broggi, Alberto and Medici, Paolo and Cardarelli, Emanuele and Cerri, Pietro, “The vislab intercontinental autonomous challenge: An extensive road test of autonomous driving technologies,” 2013, iEEE Intelligent Transportation Systems Magazine.
- [246] LaValle, Steven M., “Planning algorithms,” 2006, cambridge University Press.
- [247] Kochenderfer, Mykel J., “Decision making under uncertainty: Theory and application,” 2015, mIT Press.
- [248] Rasmussen, Carl Edward and Williams, Christopher K. I., “Gaussian processes for machine learning,” 2006, mIT Press.
- [249] Murphy, Kevin P., “Machine learning: A probabilistic perspective,” 2012, mIT Press.
- [250] Jordan, Michael I. and Mitchell, Tom M., “Machine learning: Trends, perspectives, and prospects,” 2015, science.

- [251] Hendrycks, Dan and Gimpel, Kevin, “A baseline for detecting misclassified and out-of-distribution examples in neural networks,” 2017, international Conference on Learning Representations.
- [252] Hespanha, Joao P. and Naghshtabrizi, Payam and Xu, Yonggang, “A survey of recent results in networked control systems,” 2007, proceedings of the IEEE.